

Deterministic PinPoints:

**Representative and Repeatable Simulation Region
Selection with PinPlay and Sniper**

Harish Patil (*Intel Corporation*)
Trevor E. Carlson (*Ghent University*)

Legal Disclaimer & Optimization Notice

INFORMATION IN THIS DOCUMENT IS PROVIDED "AS IS". NO LICENSE, EXPRESS OR IMPLIED, BY ESTOPPEL OR OTHERWISE, TO ANY INTELLECTUAL PROPERTY RIGHTS IS GRANTED BY THIS DOCUMENT. INTEL ASSUMES NO LIABILITY WHATSOEVER AND INTEL DISCLAIMS ANY EXPRESS OR IMPLIED WARRANTY, RELATING TO THIS INFORMATION INCLUDING LIABILITY OR WARRANTIES RELATING TO FITNESS FOR A PARTICULAR PURPOSE, MERCHANTABILITY, OR INFRINGEMENT OF ANY PATENT, COPYRIGHT OR OTHER INTELLECTUAL PROPERTY RIGHT.

Software and workloads used in performance tests may have been optimized for performance only on Intel microprocessors. Performance tests, such as SYSmark and MobileMark, are measured using specific computer systems, components, software, operations and functions. Any change to any of those factors may cause the results to vary. You should consult other information and performance tests to assist you in fully evaluating your contemplated purchases, including the performance of that product when combined with other products.

Copyright © , Intel Corporation. All rights reserved. Intel, the Intel logo, Xeon, Core, VTune, and Cilk are trademarks of Intel Corporation in the U.S. and other countries.

Optimization Notice

Intel's compilers may or may not optimize to the same degree for non-Intel microprocessors for optimizations that are not unique to Intel microprocessors. These optimizations include SSE2, SSE3, and SSSE3 instruction sets and other optimizations. Intel does not guarantee the availability, functionality, or effectiveness of any optimization on microprocessors not manufactured by Intel. Microprocessor-dependent optimizations in this product are intended for use with Intel microprocessors. Certain optimizations not specific to Intel microarchitecture are reserved for Intel microprocessors. Please refer to the applicable product User and Reference Guides for more information regarding the specific instruction sets covered by this notice.

Notice revision #20110804

Executive Summary

PinPoints methodology effectively automates the tedious task of finding and check-pointing regions of programs for Pin-based simulation.*

** Creates check-points that are representative and repeatable.*

This Tutorial

IS about:

Tools and scripts for finding representative regions (***PinPoints***) of large programs, **check-pointing** them, and **feeding** them to Pin-based simulators.

IS NOT:

A detailed **description** of **Pin**, **PinPlay**, or **Sniper**

Outline

1. Introduction
2. An Overview of Pin, SimPoint, and PinPlay
3. PinPlay kit
4. Sniper:
 - Overview
 - PinPlay Integration
 - SPEC2006 Study
 - PinPoints Integration
5. Handling Parallel Programs

Pin-based Simulation: Two Usage Models

1. Pintool **is** the simulator



2. Pintool **feeds** the simulator



Simulation Challenges: Processor & Application Complexity

1. Complex Processors/ Slower Simulation


- Detailed, cycle-accurate, simulation is slow
- Typically few thousand instructions per second

2. Complex Applications/ Porting Difficulty

- OS dependency; large memory/disk requirement
- Porting/Running on simulators impractical


Instruction Counts : Some Real Applications

Real Applications Are Long-running
Instructions (billions)



Problem: Whole-Program Simulation is Slow

Simulation Time in YEARS
@ 10,000 Instructions/Second



Solution: Select Regions to Simulate

Select One Region

- At the beginning (no skip)
- After 1 billion instructions
- After skipping a random number of instructions

Fast-forward Simulation

Select Multiple Regions


Fast-forward Simulation

- Manually by looking at performance data – too tedious
- Randomly anywhere
- Randomly from uniform regions
- By program phase analysis (SimPoint:UCSD)
- Fine-grain sampling (SMARTS: CMU)

Do the selected regions represent whole-program behavior?

How Representative are Simulation Regions?


Compare Whole-program Statistics with Statistics Predicted with Simulation of Regions



Caveats:

1. Can never “prove” region selection is “best possible”
2. Whole-program statistics computation should be “fast” (cannot use detailed simulation which could take years!)

CPI: *Average Error* SPEC2000(IA32) Whole Program vs. Selected Points



PinPoints = Pin (Intel) + SimPoint (UCSD)*

❑ **What are PinPoints** : Representative Regions of Programs

- Automatically chosen
- Validated (multiple times: latest with Sniper from UGHent)

❑ **PinPoints Methodology**:

- Use : Pin, SimPoint, a Pin-based Simulator (e.g. Sniper), multiple scripts
- Find PinPoints, checkpoint them, validate them

** Acknowledgement: Timothy Sherwood, Erez Perelman, Greg Hamerly and Brad Calder*

Simulation Challenges: Our Solutions

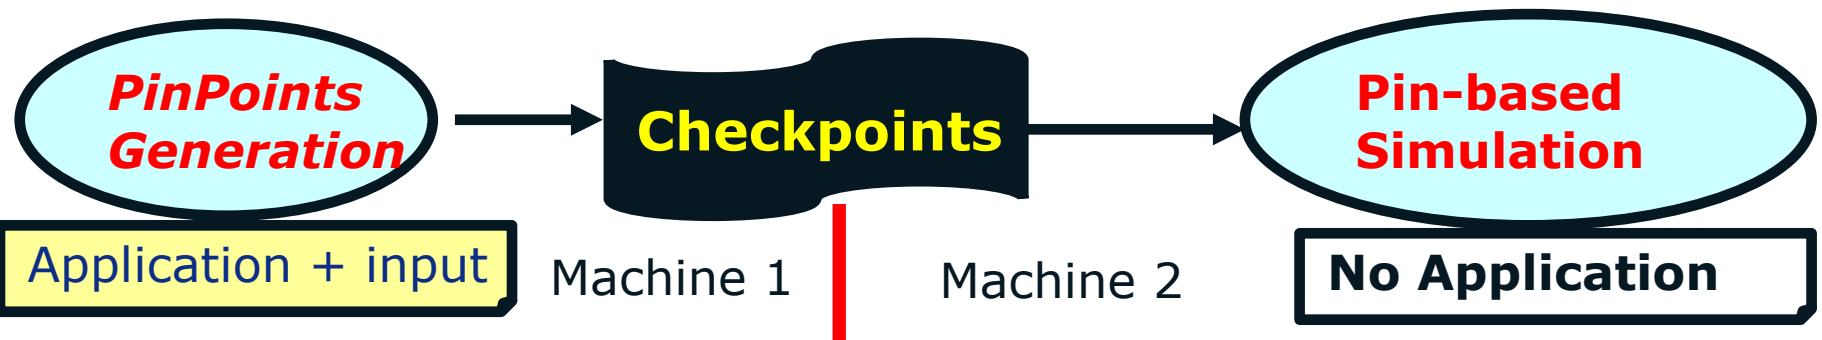
1. Detailed, cycle-accurate, simulation is slow

Solution: Find/simulate only representative regions

- PinPoints cover $\ll 1\%$ of whole-program execution
→ **vastly reduced simulation time**

2. Complex Applications/ Porting Difficulty

- Create check-points transfer to simulation



PinPoints-based Simulation + Prediction

1. Simulate with check-points for PinPoints



- Checkpoints capture region execution (details later)
- No fast-forwarding needed

2. Predict : CPI, MPI,... Stats-per-instruction

- Use statistics from PinPoints simulation, total instruction count, PinPoints weights...

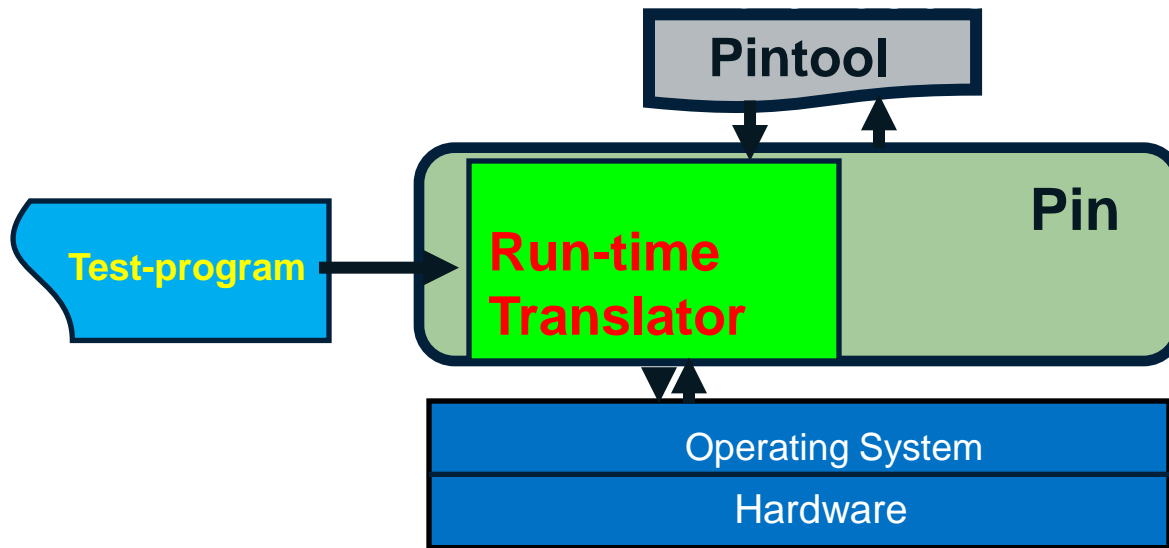
Pin, SimPoint, PinPlay : Overview



Pin: A Tool for Writing Program Analysis Tools

```
sub    $0xff, %edx
movl   0x8(%ebp), %eax
jle    <L1>
```

```
counter++; print(IP)
sub    $0xff, %edx
counter++; print(EA)
movl   0x8(%ebp), %eax
counter++; print(br taken)
jle    <L1>
```



```
$ pin -t pintool -- test-program
```

Normal output
+ *Analysis*
output

Pin: A Dynamic Instrumentation Framework from Intel

<http://www.pintool.org>

SimPoint from UCSD

Goals

- The goals of this research are:
 - To create an automatic system that is capable of intelligently **characterizing time-varying** program behavior
 - To provide both analytic and software tools to help with program **phase identification**
 - To demonstrate the utility of these tools for finding places to simulate (SimPoints)
 - **Without full program detailed simulation**




ASPLOS: Sherwood et. al.

UCSD SimPoint Work: Key Idea

Simulation Points

- Should represent the program behavior *independent of hardware*
- Identified using program-dependent metric *relative execution counts of basic blocks*

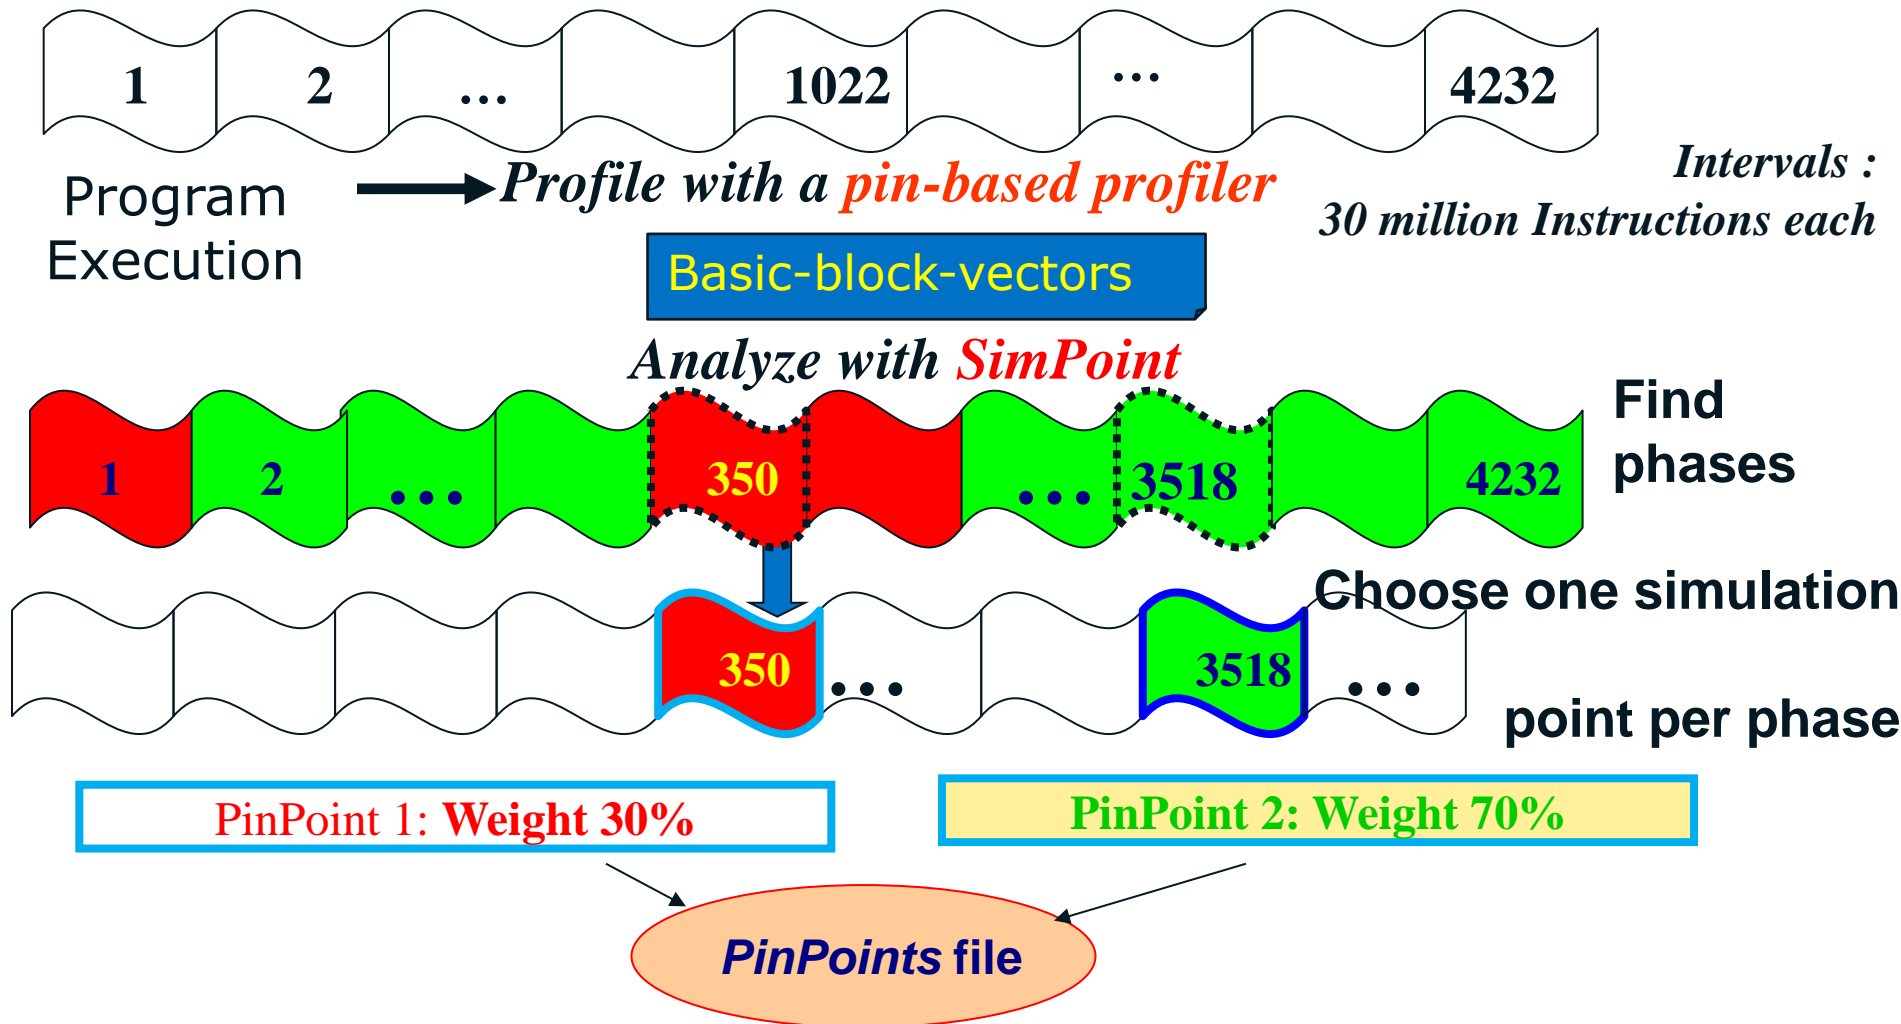
Large Scale Behavior (gzip)



UCSD SimPoint Work: 3 Steps

1. Profile → **Basic Block Vectors**
2. Analyze Profiles → **Simulation Points**
 - **SimPoints** : (skip length, weight)
3. Verify → Compare Selective Statistics (ipc, miss rates) → **Error rates**
 - Compare full run statistics vs SimPoints statistics

PinPoints = Pin + SimPoint



Two Phases => Two PinPoints

PinPoints: Estimating Total Execution Time

Total Execution Time = Total Cycles / Frequency

- We know the simulated Frequency; need to know Total Cycles for *full* run of the binary on the Simulator

Total Cycles Simulated = (Weighted CPI) * (Total Instructions)

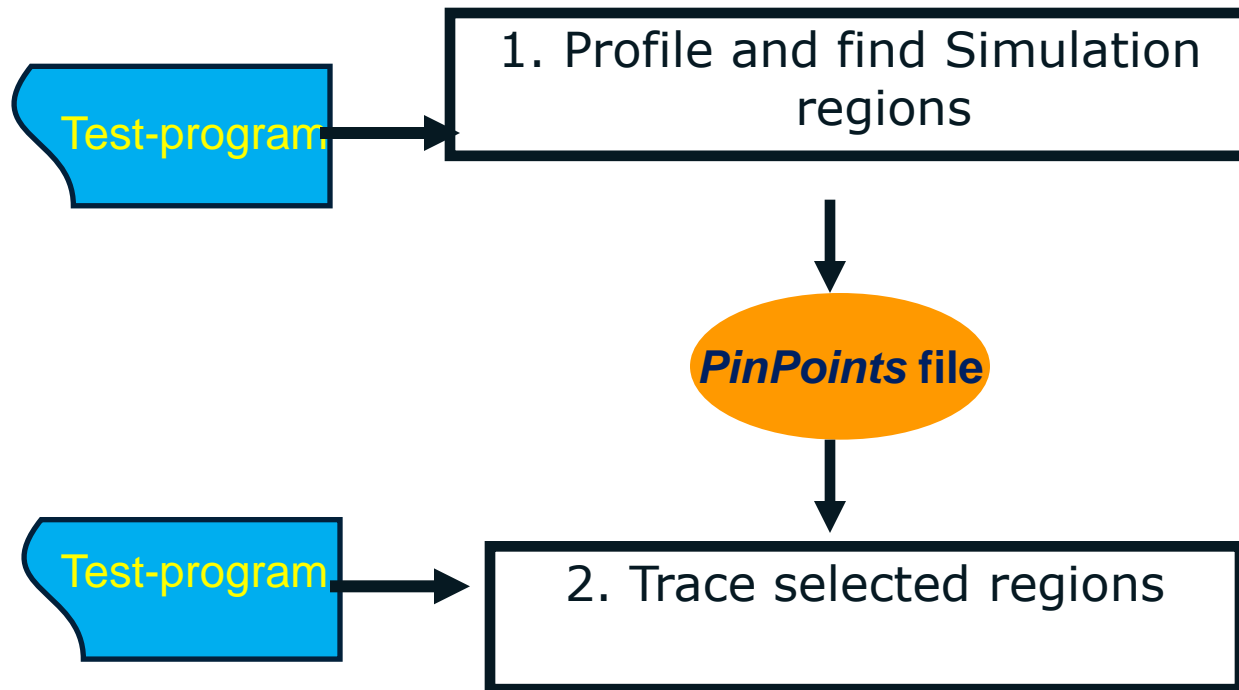
- PinPoints provides the Total number of instructions in the PinPoints file.

Weighted CPI can be determined through simulation of PinPoints regions and weighting of results:

$$\text{Weighted CPI} = \sum \text{Weight}_i * \text{CPI}_i$$

- CAUTION: Use the formula only for statistics normalized by instructions : CPI computation OK; IPC computation is NOT OK

PinPoints : The Repeatability Challenge

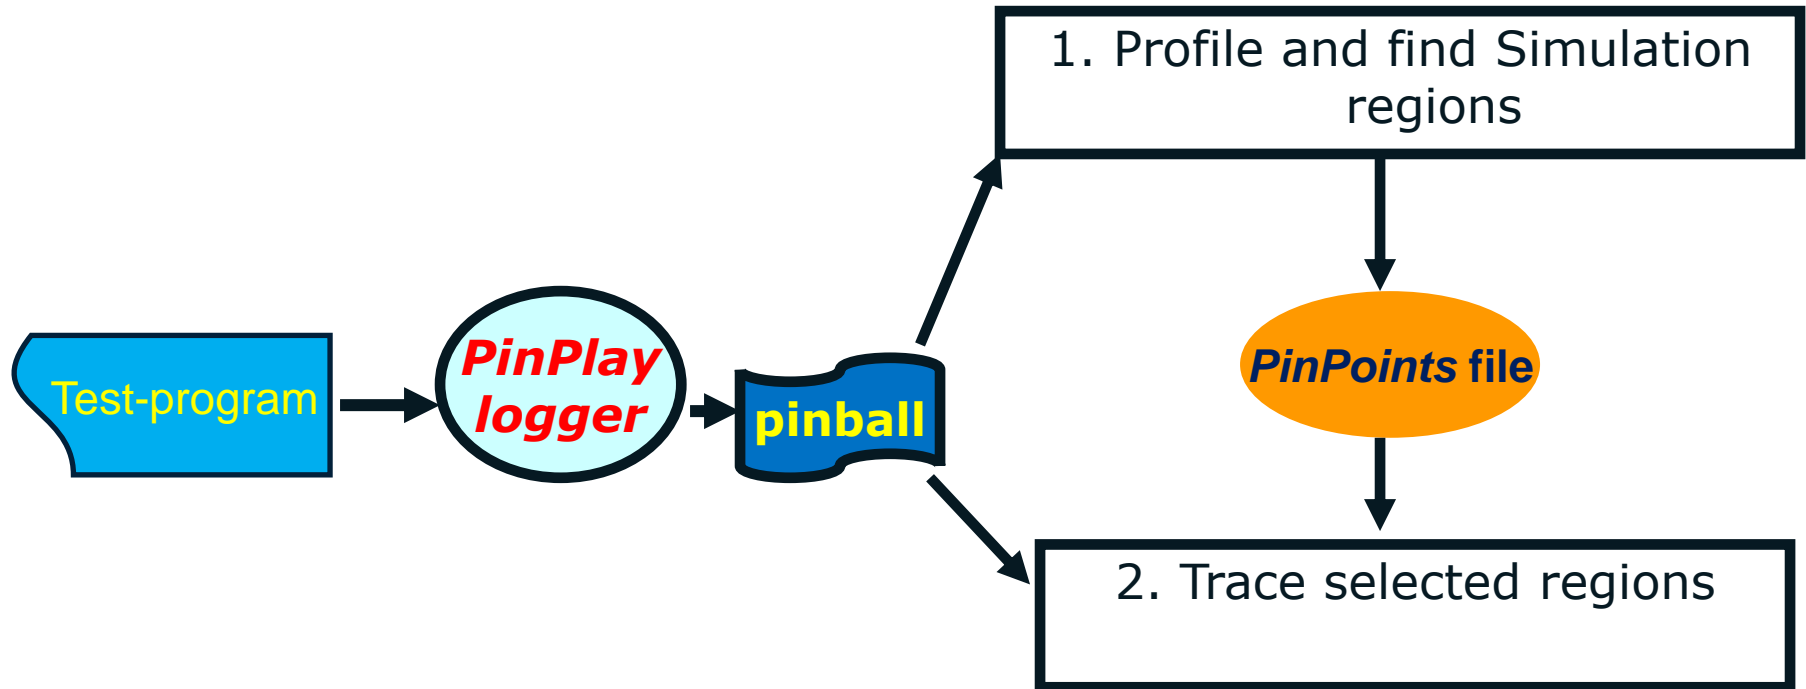


Problem: Two runs are not exactly same → PinPoints missed

Found this for 25/54 SPEC2006 runs!

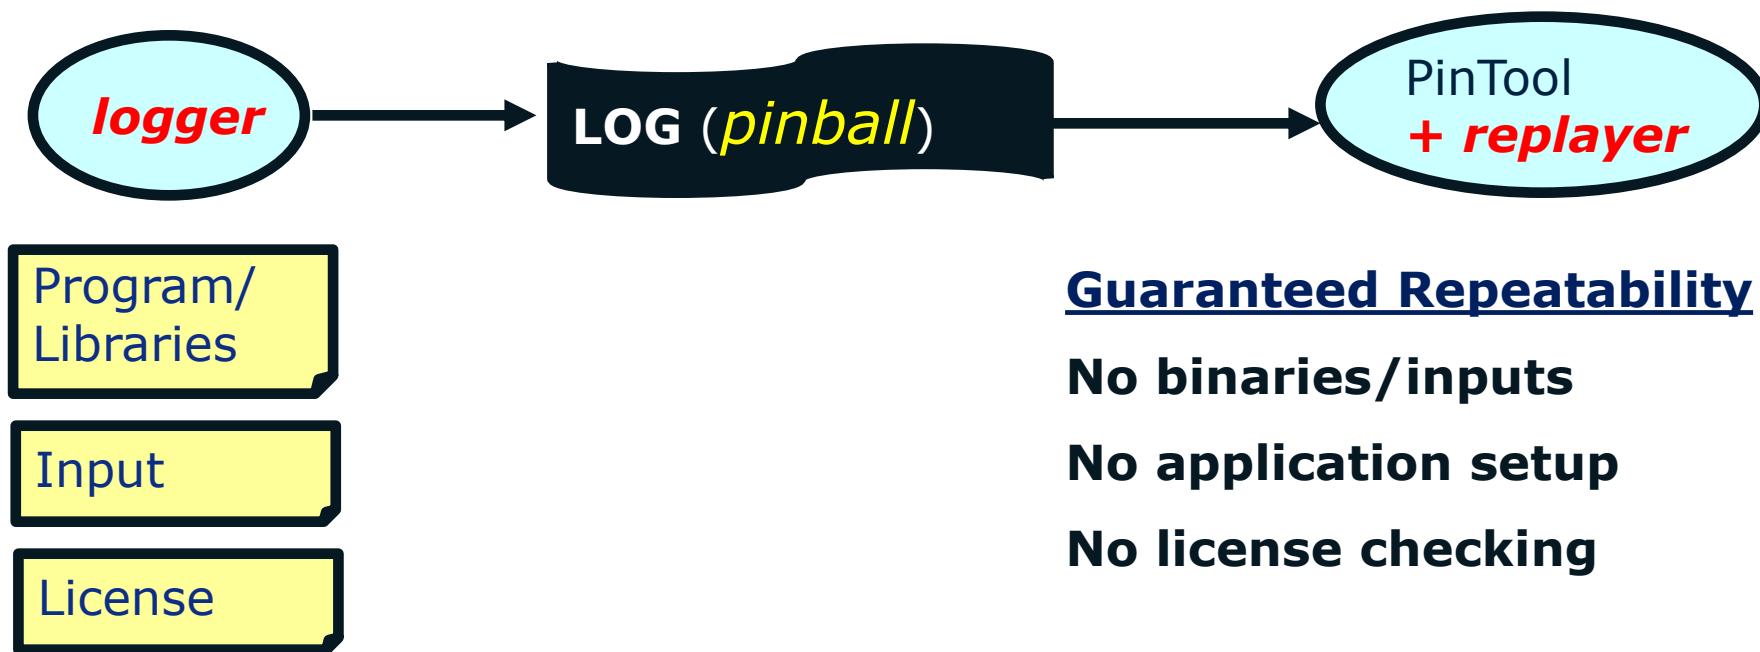
[*"PinPoints out of order" "PinPoint End seen before Start"*]

Enters PinPlay To Provide Repeatability



Two runs are same → PinPoints guaranteed to be reached

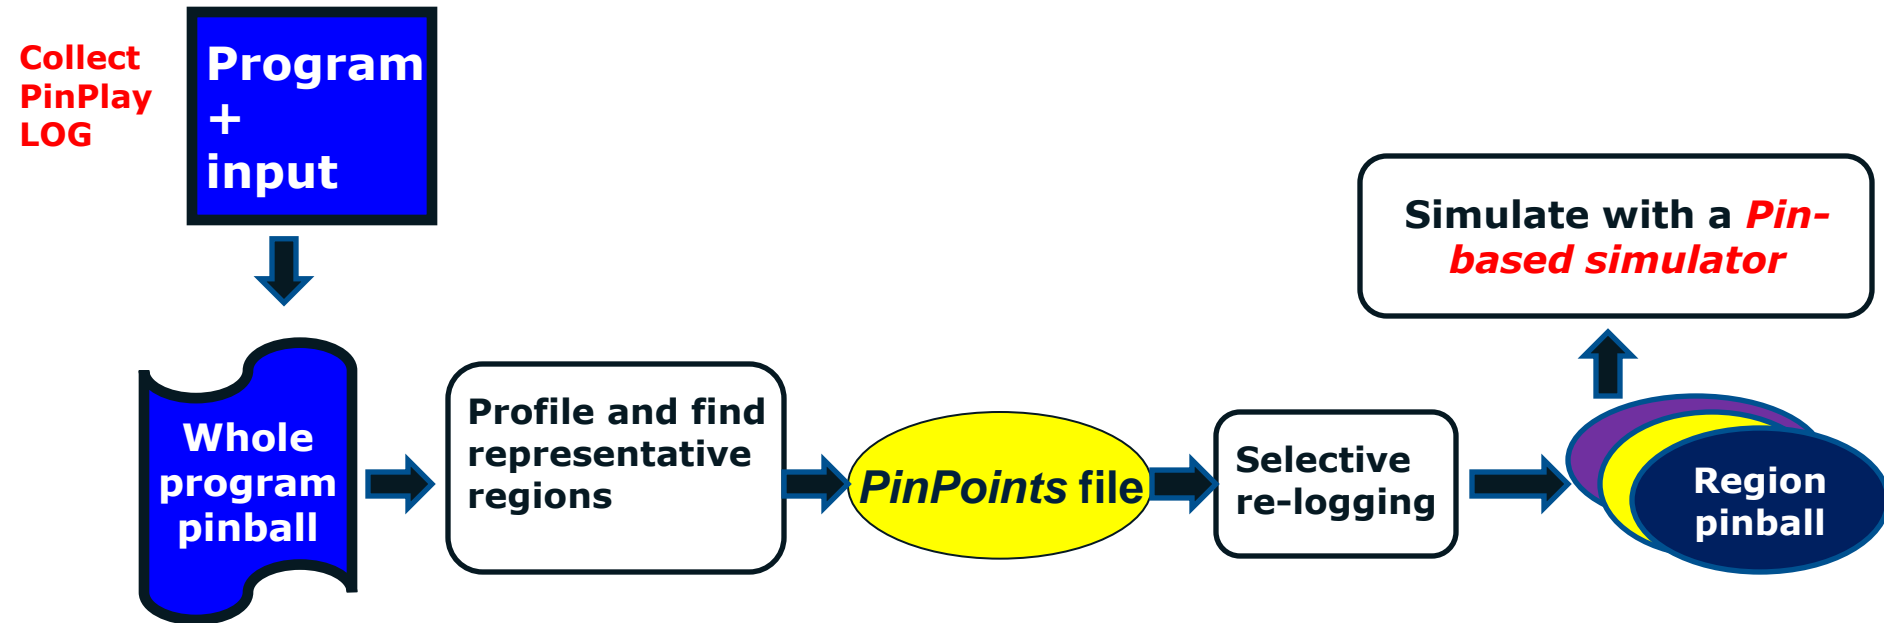
PinPlay*: Workload Capture and Deterministic Replay Framework



Record Once Replay + Analyze Multiple Times Anywhere!

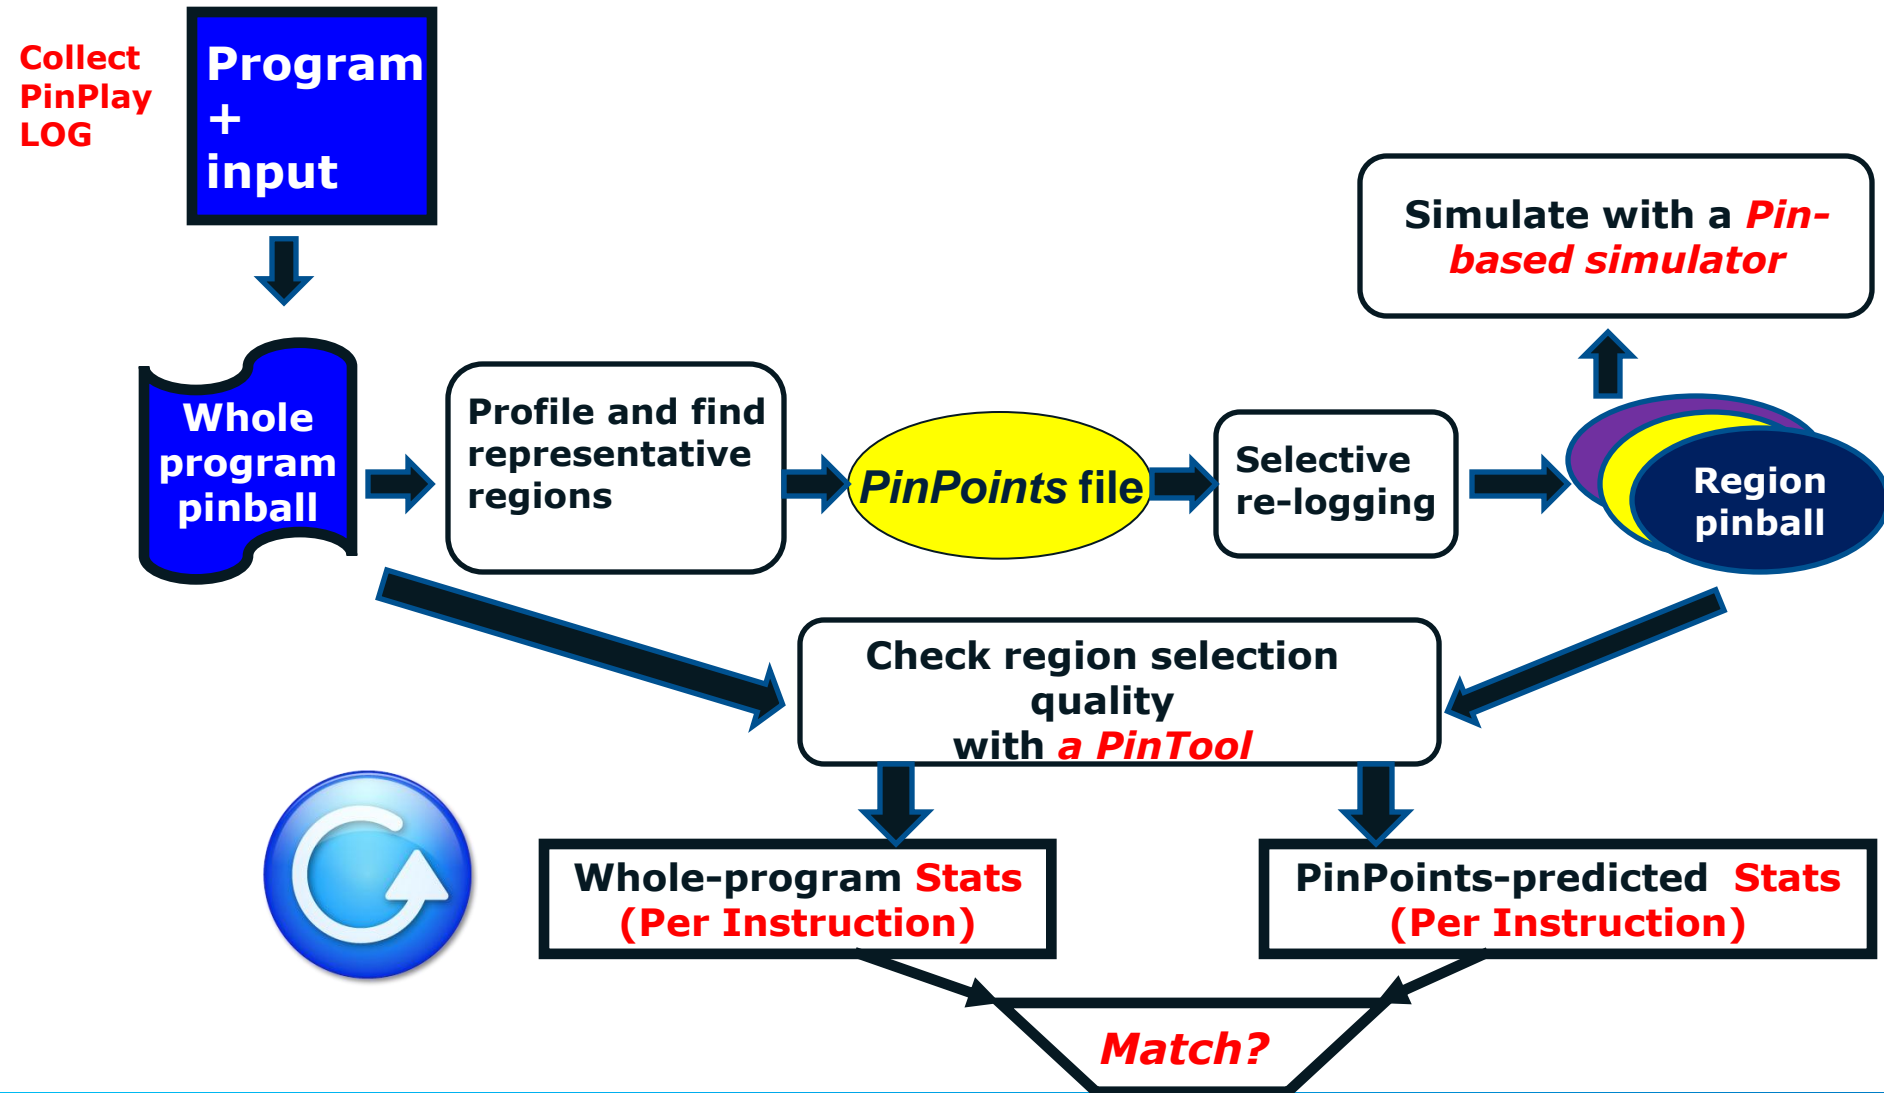
* *Developers: Cristiano Pereira, James Cownie, Harish Patil*

PinPlay + PinPoints: Basic Flow

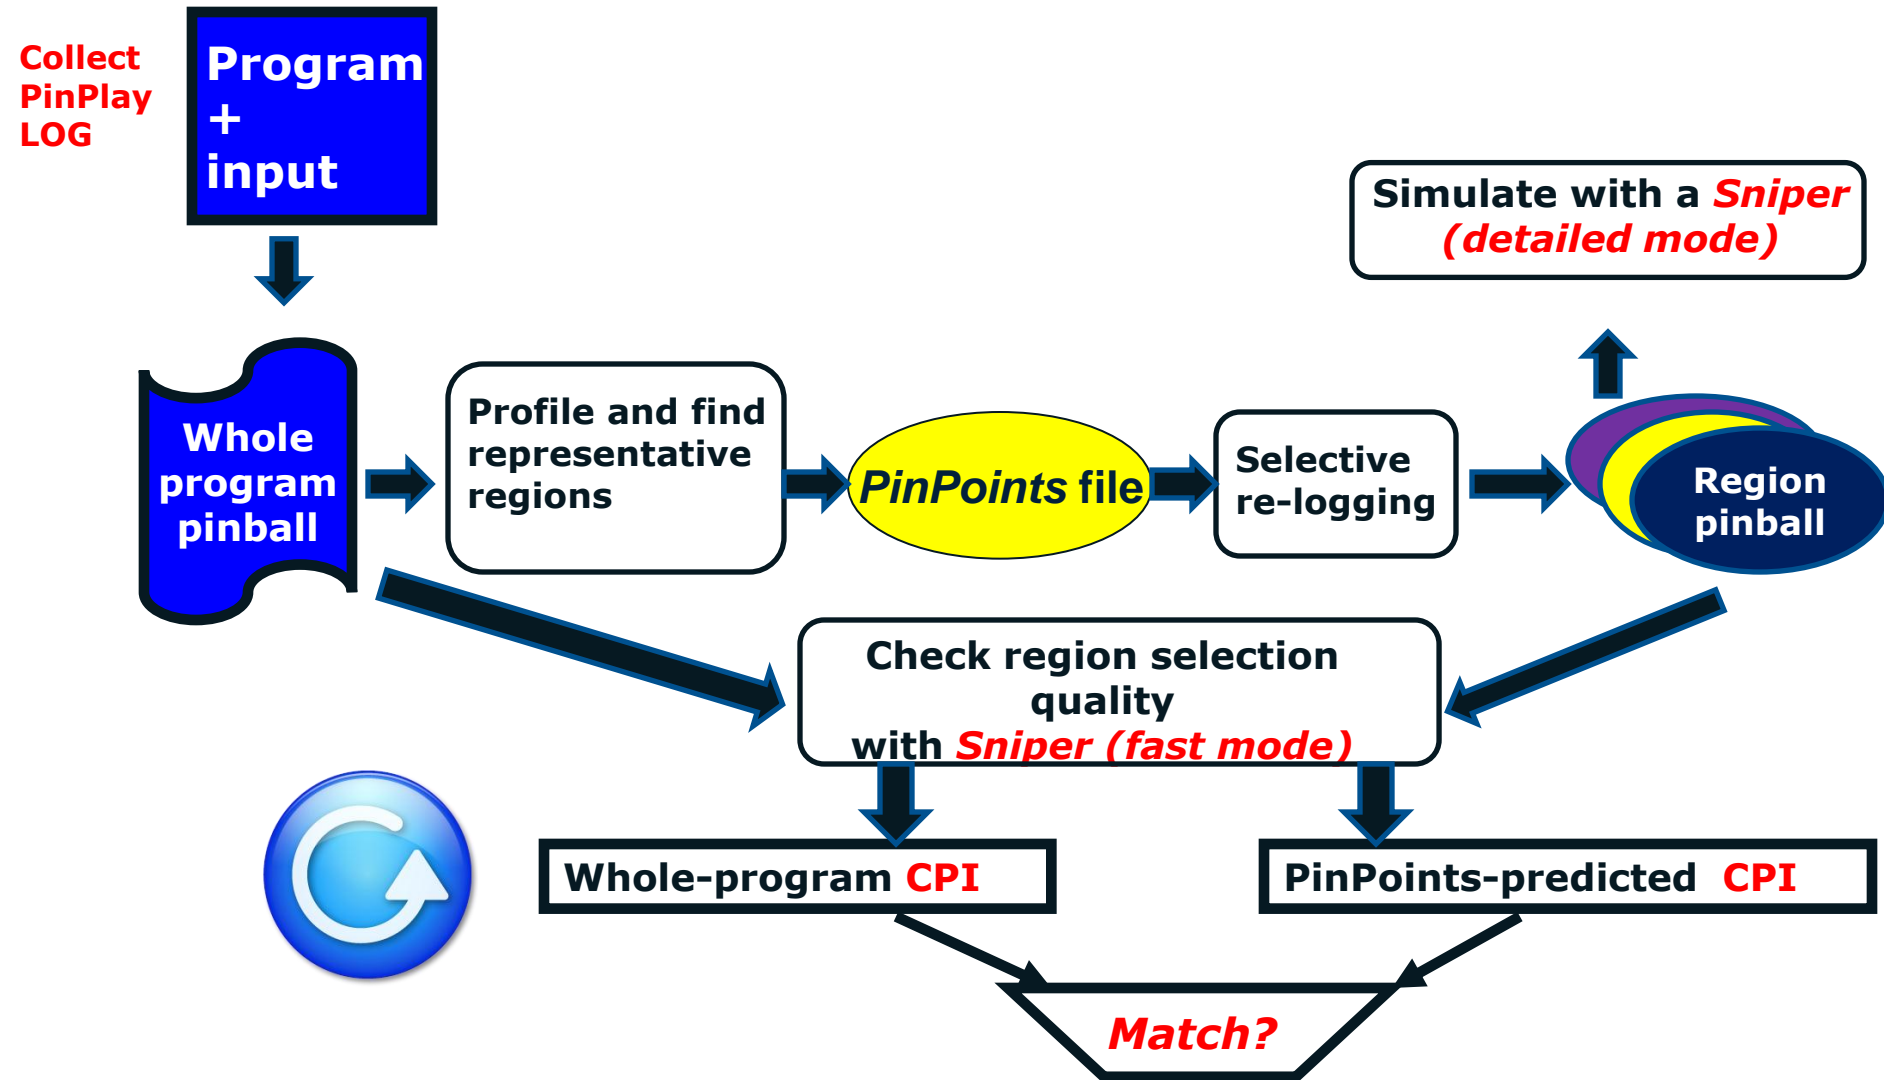


- PinPoints are representative (validation/tuning possible)
- PinPoints cover $\ll 1\%$ of whole-program execution
→ **vastly reduced simulation time**

PinPoints validation: “Functional Correlation”



PinPlay + PinPoints + Sniper



The PinPlay kit



Download from <http://www.pinplay.org>

pinplay-1.0-pin-2.12-55942-gcc.4.4.7-linux
<PinPlay version> <base Pin kit used>

PinPlay Kit = Pin kit + *extras/pinplay*

```
pinplay-1.0-pin-2.12-55942-gcc.4.4.7-linux/extras/pinplay/  
|-- PinPoints  
|  |-- bin  
|  |-- scripts  
|-- bin  
|  |-- ia32  
|  |-- intel64  
|-- examples  
|  |-- tests  
|-- include  
|-- include -- pinplay.H  
|-- lib  
|  |-- ia32 -- libpinplay.a  
|  |-- intel64  
|-- lib-ext  
|  |-- ia32  
|  |-- intel64  
|-- README.PinPoints  
|-- bin  
|-- LICENSE.simpoint  
|-- simpoint  
|-- Makefile  
|-- bimodal.H  
|-- pinplay-branch-predictor.cpp  
|-- pinplay-driver.cpp
```

Installation : Linux-64

[**Prerequisite:** gcc/g++ available for both 32 and 64 bit]

```
%setenv PIN_KIT <root of pinplay kit>
%cd $PIN_KIT/extras/pinplay/examples
% make
```

Builds and tests **pinplay-driver.so** and **pinplay-branch-predictor.so**

```
extras/pinplay/bin/
|-- ia32
|   |-- nullapp
|   |-- pinplay-branch-predictor.so
|   |-- pinplay-driver.so
|-- intel64
|   |-- nullapp
|   |-- pinplay-branch-predictor.so
|   |-- pinplay-driver.so
```

```
extras/pinplay/examples/
|-- Makefile
|-- bimodal.H
|-- foo.bimodal.ia32.out
|-- foo.bimodal.intel64.out
|-- hello32
|-- hello64
|-- pinball
|   |-- foo.0.dyn_text.bz2
|   |-- foo.0.race.bz2
|   |-- foo.0.reg.bz2
|   |-- foo.0.result
|   |-- foo.0.result_play
|   |-- foo.0.sel.bz2
|   |-- foo.0.sync_text.bz2
|   |-- foo.address
|   |-- foo.log.txt
|   |-- foo.procinfol.xml
|   |-- foo.replay.txt
|   |-- foo.text.bz2
```


Enabling a Pintool for PinPlay

```
#include "pinplay.H"
```

```
PINPLAY_ENGINE pinplay_engine;
```

```
Knob<BOOL>KnobReplayer(KNOB_MODE_WRITEONCE, KNOB_FAMILY,  
                      KNOB_REPLAY_NAME, "0", "Replay a pinball");  
Knob<BOOL>KnobLogger(KNOB_MODE_WRITEONCE, KNOB_FAMILY,  
                    KNOB_LOG_NAME, "0", "Create a pinball");
```

```
pinplay_engine.Activate(argc, argv, KnobLogger, KnobReplayer);
```

Link in *libpinplay.a*, *libzlib.a*,
libbz2.a

Restrictions:

1. PinTool shouldn't change application control flow
2. Image API not available during replay

Example: pinplay-branch-predictor.cpp

```
#define KNOB_LOG_NAME "log"
#define KNOB_REPLAY_NAME "replay"
#define KNOB_FAMILY "pintool:pinplay-driver"

PINPLAY_ENGINE pinplay_engine;

KNOB_COMMENT pinplay_driver_knob_family(KNOB_FAMILY, "PinPlay Driver Knobs");

KNOB<BOOL>KnobReplayer(KNOB_MODE_WRITEONCE, KNOB_FAMILY,
                      KNOB_REPLAY_NAME, "0", "Replay a pinball");
KNOB<BOOL>KnobLogger(KNOB_MODE_WRITEONCE, KNOB_FAMILY,
                    KNOB_LOG_NAME, "0", "Create a pinball");

int main(int argc, char *argv[])
{
    if( PIN_Init(argc,argv) )
    {
        return Usage();
    }

    outfile = new ofstream(KnobStatFileName.Value().c_str());
    bimodal.Activate(KnobPhases, outfile);

    pinplay_engine.Activate(argc, argv, KnobLogger, KnobReplayer);

    PIN_AddThreadStartFunction(threadCreated, reinterpret_cast<void *>(0));

    PIN_StartProgram();
}
```

PinPlay-enabled PinTools : 3 Modes

1. Regular Analysis mode

```
$ pin -t pintool -- test-program
```

Normal output
+ *Analysis*
output

2. Logging Mode

```
$ pin -t pintool -log -log:basename pinball/foo -- test-program
```

pinball

3. Replay Mode

```
$ pin -t pintool -replay -replay:basename pinball/foo -- nullapp
```

Example: pinplay-branch-predictor.so

```
% $PIN_KIT/pin -t  
$PIN_KIT/extras/pinplay/bin/intel64/pinplay-  
branch-predictor.so -- hello
```

Creates "bimodal.out"

```
%pin -t pinplay-branch-predictor.so -log -  
log:basename pinball/foo - hello
```

Creates "biomodal.out" and "pinball/foo*"

```
%pin -xyzzz -reserve_memory pinball/foo.address -  
replay -replay:basename pinball/foo --  
$PIN_KIT/extras/pinplay/intel64/bin/nullapp
```

Creates "bimodal.out"

Using PinPoints/scripts* :Logging (uses pinplay-driver.so):

```
:%$PIN_KIT/extras/pinplay/PinPoints/scripts/logger.py  
--pinplayhome $PIN_KIT --mode st --logfile pinball/foo  
hello
```

```
Creates: pinball/foo_28293.*  
         foo_<pid>.*
```

```
Usage: logger.py [options] --logfile FILE application app_arguments  
Version: 1.48
```

```
-h, --help          show this help message and exit
```

```
--mode=MODE          MODE specifies the type of program to be logged. No  
                     default. Must be defined in either the tracing  
                     configuration file or this option.  
                     st - single-threaded  
                     mt - multi-threaded  
                     mpi - MPI single-threaded  
                     mpi_mt - MPI multi-threaded
```

* *Developed by Mack Stallcup*

Using PinPoints/scripts: Replaying (uses pinplay-driver.so)

```
:%$PIN_KIT/extras/pinplay/PinPoints/scripts/replayer.py  
--pinplayhome $PIN_KIT -replay_file pinball/foo_28293
```

```
Usage: replayer.py [options] --replay_file=FILE
```

```
-h, --help          show this help message and exit
```

Example: PinPoints on specrand (SPEC2006)

```
% $PIN_KIT/extras/pinplay/PinPoints/scripts/pinpoints.py -h
```

```
Usage: pinpoints.py phase [options]
```

```
--cfg FILE, --config_file FILE
                        Give one, or more, file(s) containing the application
                        tracing parameters. Must use '--cfg' for each file.
```

```
% cat specrand.test.cfg
[Parameters]
program_name:    specrand
input_name:      test
command:         "./base.exe 1255432124 234923 > rand.234923.out"
epilog_length:   0
maxk:            5
mode:            st
num_proc:        1
prolog_length:   0
slice_size:      100000
warmup_length:   300000
```

Warmup

Prolog

Simulation

Epilog

Specrand PinPoints: Basic Flow

-l, --log

Generate whole program pinballs for the application.

**Program
+
input**



**Profile and find
representative
regions**



***PinPoints* file**

**Selective
re-logging**



**Region
pinball**

-p, --region_pinball

Re-log whole program pinballs using representative regions from Simpoint to generate region pinballs.

-b, --basic_block_vector

Generate basic block vectors for whole program pinballs.

-s, --simpoint

Run Simpoint using whole program pinball basic block vectors.

```
% $PIN_KIT/extras/pinplay/PinPoints/scripts/pinpoints.py --cfg specrand.test.cfg  
--pinplayhome=$PIN_KIT -l -b -s -p
```


Specrand PinPoints: Output

```
Script args:      --cfg specrand.test.cfg --pinplayhome=/nfs/mmdc/disks/pinplay/proj/PinPlayExternal/Wh
atIf/released/pinplay-1.0-pin-2.12-55942-gcc.4.4.7-linux -l -b -s -p
Tracing mode:    st
Program name:    specrand
Input name:      test
Command:         "./base.exe 1255432124 234923 > rand.234923.out"
Maxk:           20
Cutoff:         1.0
Warmup length:   3,001,500
Prolog length:   0
Slice size (region): 1,000,000
Epilog length:   0
Pinplayhome:    /nfs/mmdc/disks/pinplay/proj/PinPlayExternal/WhatIf/released/pinplay-1.0-pin-2.12-55
942-gcc.4.4.7-linux
Whole program directory: whole_program.test
Data/lit/pp directory:  specrand.1.test
Trace file name format: specrand.1.test_t0rX_warmup3001500_prolog0_region1000000_epilog0
Number cores/system: 8
```

Warm-up: Extra 1500 instructions added
for overlap avoidance due to basic-block
level instruction counting

Trace file name format (for region pinballs):

specrand.1.test_**t0rX_warmup3001500_prolog0_region1000000_epilog0**

Warmup

Prolog

Simulation

Epilog

Specrand PinPoints: Results

```
999.specrand/  
|-- specrand.1.test_31095.Data  
|-- specrand.1.test_31095.pp  
`-- whole_program.test
```

```
whole_program.test/  
|-- specrand.1.test_31095.0.dyn_text.bz2  
|-- specrand.1.test_31095.0.reg.bz2  
|-- specrand.1.test_31095.0.result  
|-- specrand.1.test_31095.0.result_play  
|-- specrand.1.test_31095.0.sel.bz2  
|-- specrand.1.test_31095.address  
|-- specrand.1.test_31095.log.txt  
|-- specrand.1.test_31095.procinfol.xml  
|-- specrand.1.test_31095.replay.txt  
`-- specrand.1.test_31095.text.bz2
```

```
specrand.1.test_31095.Data  
|-- create_region_file.out  
|-- simpoint.out  
|-- specrand.1.test_31095.T.0.bb  
|-- specrand.1.test_31095.pinpoints.csv  
|-- t.bb  
|-- t.labels  
|-- t.simpoints  
`-- t.weights
```

```
% du -h base.exe whole_program.test  
492K    base.exe  
192K    whole_program.test
```

```
% grep inscount whole_program.test/*  
whole_program.test/specrand.1.test_31095.0.result:inscount: 617602978  
whole_program.test/specrand.1.test_31095.0.result_play:inscount: 617602978
```

- Dynamic instruction count : 617 million
Whole-program pinball size : 192K (base.exe 492K)

specrand.1.test_31095.pinpoints.csv

```
# Regions based on '/nfs/mmdc/disks/pinplay/proj/PinPlayExternal/WhatIf/released
/pinplay-1.0-pin-2.12-55942-gcc.4.4.7-linux/extras/pinplay/PinPoints/scripts/cre
ate_region_file.pl -seq_region_ids -tid 0 -region_file t.simpoints -weight_file
t.weights t.bb':
comment,thread-id,region-id,simulation-region-start-icount,simulation-region-end
-icount,region-weight
# Region = 1 Slice = 467 Icount = 467003527 Length = 1000005 Weight = 0.0016
cluster 0 from slice 467,0,1,467003527,468003532,0.001618
# Region = 2 Slice = 93 Icount = 93000803 Length = 1000004 Weight = 0.2006
cluster 1 from slice 93,0,2,93000803,94000807,0.200647
# Region = 3 Slice = 577 Icount = 577004084 Length = 1000007 Weight = 0.2427
cluster 2 from slice 577,0,3,577004084,578004091,0.242718
# Region = 4 Slice = 186 Icount = 186001529 Length = 1000001 Weight = 0.3447
cluster 3 from slice 186,0,4,186001529,187001530,0.344660
# Region = 5 Slice = 51 Icount = 51000481 Length = 1000006 Weight = 0.2104
cluster 4 from slice 51,0,5,51000481,52000487,0.210356
# Total instructions in 5 regions = 5000023.
# Total instructions in workload = 617602969.
# Total slices in workload = 618.
# Overall dynamic coverage of workload by these regions = 1.0000 (including lost
instrs).
```

5 PinPoints → 5 Region Pinballs

```
% tree specrand.1.test_31095.pp | grep address
|-- specrand.1.test_31095_t0r1_warmup3001500_prolog0_region1000005_epilog0_001_0-00161.0.address
|-- specrand.1.test_31095_t0r2_warmup3001500_prolog0_region1000004_epilog0_002_0-20064.0.address
|-- specrand.1.test_31095_t0r3_warmup3001500_prolog0_region1000007_epilog0_003_0-24271.0.address
|-- specrand.1.test_31095_t0r4_warmup3001500_prolog0_region1000001_epilog0_004_0-34466.0.address
|-- specrand.1.test_31095_t0r5_warmup3001500_prolog0_region1000006_epilog0_005_0-21035.0.address
```

Warmup length:	3,001,500	Warm-up: Extra 1500 instructions
Prolog length:	0	
Slice size (region):	1,000,000	
Epilog length:	0	

Trace file name format (for region pinballs):

specrand.1.test_**t0rX**_**warmup3001500**_**prolog0**_region**1000000**_epilog**0**

Warmup

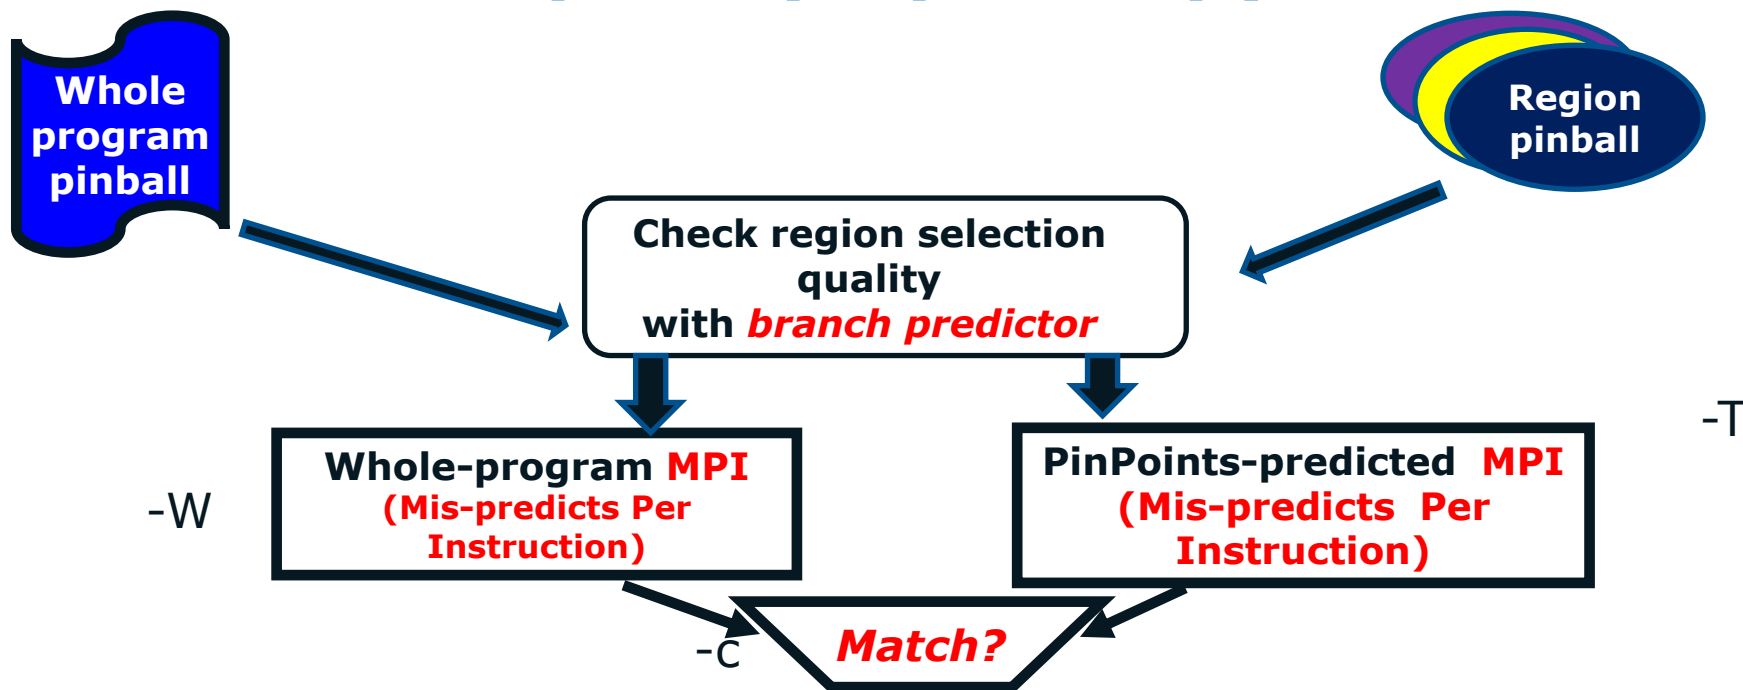
Prolog

Simulation

Epilog

```
% grep inscount specrand.1.test_31095_t0r*.0.result
specrand.1.test_31095_t0r1_warmup3001500_prolog0_region1000005_epilog0_001_0-00161.0.result:inscount: 4001505
specrand.1.test_31095_t0r2_warmup3001500_prolog0_region1000004_epilog0_002_0-20064.0.result:inscount: 4001490
specrand.1.test_31095_t0r3_warmup3001500_prolog0_region1000007_epilog0_003_0-24271.0.result:inscount: 4001500
specrand.1.test_31095_t0r4_warmup3001500_prolog0_region1000001_epilog0_004_0-34466.0.result:inscount: 4001499
specrand.1.test_31095_t0r5_warmup3001500_prolog0_region1000006_epilog0_005_0-21035.0.result:inscount: 4001504
```

Functional Correlation with brpred_pinpoints.py

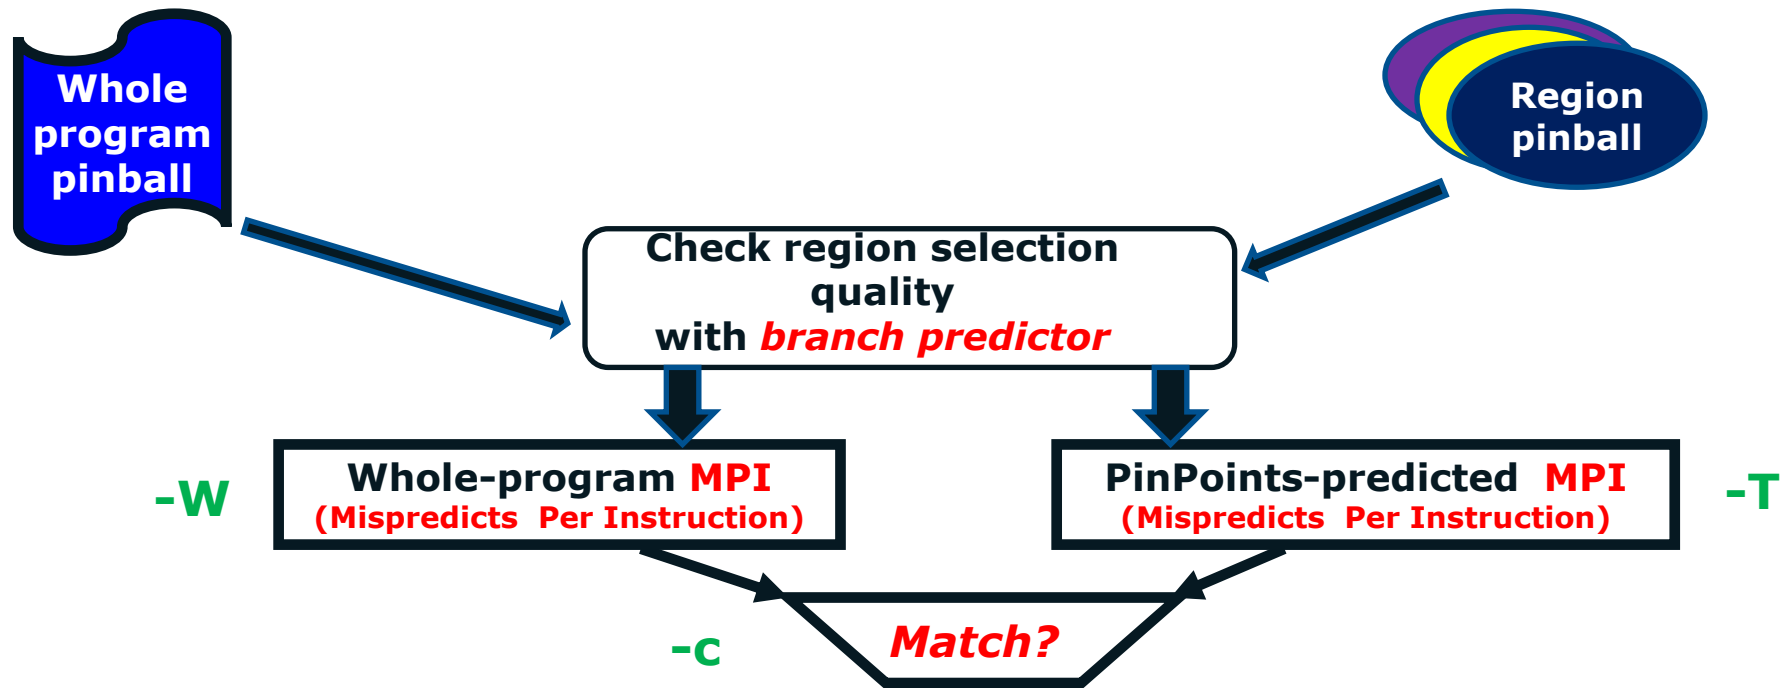


```
% $PIN_KIT/extras/pinplay/PinPoints/scripts/brpred_pinpoints.py -h
Usage: brpred_pinpoints.py phase [options]
```

```
-T, --region_sim
-W, --whole_sim
-c, --calc_func
```

Run the simulator on the whole program pinballs.
Run the simulator on the region pinballs.
Calculate the functional correlation, using the metric of interest, for a set of representative regions. Must have already generated simulator data, either using phases '--whole_sim' and '--region_sim' or the appropriate options for your simulator, before running this phase.

Specrand PinPoints : *Functional Correlation* with pinplay-branch-predictor.so



```
% $PIN_KIT/extras/pinplay/PinPoints/scripts/brpred_pinpoints.py --cfg specrand.test.cfg --pinplayhome=$PIN_KIT -W -T -C
```

```
*** Calculating functional correlation *** February 18, 2013 20:55:17
```

```
PID: 31095
```

```
Predicted MPI: 6.952073
```

```
Measured MPI: 6.969056
```

```
Functional correlation: 0.997563 (p/m)
```


Tuning PinPoints Selection: bzip2: input "chicken"

Initial results (maxK=5): Error 31%

Re-generated PinPoints with maxK=20

```
% $PIN_KIT/extras/pinplay/PinPoints/scripts/pinpoints.py --pinplayhome=$PIN_KIT -b -s  
-p --cfg bzip2.chicken.cfg
```

```
maxk: - 20
```



Tuned results (maxK=20) : Error 2.7%

```
% $PIN_KIT/extras/pinplay/PinPoints/scripts/brpred_metric.py --pinplayhome=$PIN_KIT  
--actual=whole_program.chicken/ --predicted=bzip2.1.chicken_9999.pp --calc_func  
Predicted metric: 6.8173  
Actual metric: 7.0338  
Functional correlation: 0.969226 (p/a)
```

Logging a Specific Region

Use “PinPlay Controller” knobs:

```
% $PIN_KIT/pin -t $PIN_KIT/extras/pinplay/bin/intel64/pinplay-driver.so -help -- /bin/true
```

Start of region: `-log:skip [default]`
Number of instructions to skip from beginning

`-log:start_address`
Address and count to trigger a start (e.g. 0x400000, main, memcpy:2, /lib/tls/libc.so.6+0x1563a:1)

End of region:

`-log:stop_address`
Address and count to trigger a stop (e.g. 0x400000, main, memcpy:2, /lib/tls/libc.so.6+0x1563a:1)

`-log:length [default]`
Number of instructions to execute before stopping

Uniform sampling

```
-log:uniform_count [default ]  
    Number of uniform samples to trigger.  
-log:uniform_length [default ]  
    Number of instructions to capture periodically  
-log:uniform_period [default ]  
    Number of instructions to skip periodically  
-log:uniform_skip [default ]  
    Number of skip before uniform sampling starts.
```


Example: Logging a Specific Region

Skip 5000 instructions, log 1000, and then exit:

```
% $PIN_KIT/extras/pinplay/PinPoints/scripts/logger.py --pinplayhome=$PIN_KIT --mode st --logfile  
pinball/foo --log_options="-log:skip 5000 -log:length 1000 -log:early_out" -- /bin/ls
```

```
% grep inscount pinball/*  
pinball/foo_589.0.result:inscount: 996
```

pinball/*.log.txt:

```
[0] -----  
[0] region-start  
[0] + RECEIVED AND PROCESSED START event: ip:0x2aaaaaabd290 mcount: 0 icount: 0 time:Tue Feb 19 13:10:13 2013  
[0] Region# : 0  
[0] Region has 1 existing thread(s)  
[0] InitRegion called by thread: 0  
[0] First thread in the region  
[0] region-end  
[0] + RECEIVED STOP event: mcount: 345 icount: 986 time:Tue Feb 19 13:10:13 2013  
[0] getting ready to terminate due to : log:early_out 345  
[0] TerminateTrace: 348 0x2aaaaaaadef5  
[0] EndRegion called by: 0 in_region: 1  
[0] Exited_thread_count: 0  
[0] Last thread, writting global files  
[0] + FINISHED PROCESSING STOP event: mcount: 348 icount: 996  
[0] Exiting due to log:early_out ic: 996 mc: 348
```

Sniper Overview





ExaScience Lab
Intel Labs Europe



EXASCALE COMPUTING

THE SNIPER MULTI-CORE SIMULATOR

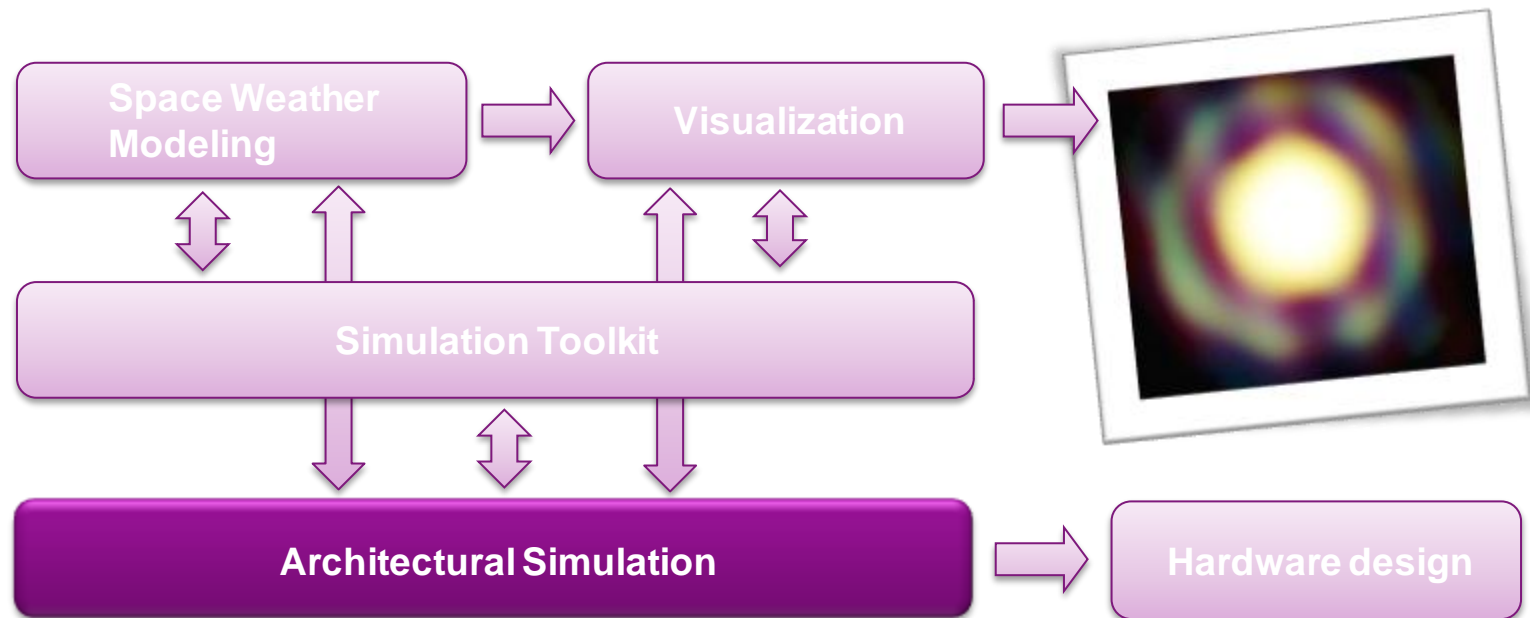
TREVOR E. CARLSON, WIM HEIRMAN, IBRAHIM HUR
KENZO VAN CRAEYNEST AND LIEVEN EECKHOUT



[HTTP://WWW.SNIPERSIM.ORG](http://www.snipersim.org)
SATURDAY, FEBRUARY 23, 2013
HPCA 2013, SHENZHEN, CHINA

INTEL EXASCIENCE LAB

- Collaboration between Intel, imec and 5 Flemish universities
- Study Space Weather as an HPC workload



SIMULATION


- Design tomorrow's processor using today's hardware
- Simulation
 - Obtain performance characteristics for new architectures
 - Architectural exploration
 - Early software optimization

DEMANDS ON SIMULATION ARE INCREASING

- Increasing core counts
 - Hardware today is seeing large core-counts
 - 2011: 10-core Intel Xeon Westmere-EX
 - 2012: Intel MIC Knights Corner (60+ cores)
 - Linear increase in simulator workload
 - Single-threaded simulator sees a rising gap
 - workload: increasing target cores
 - available processing power: near-constant single-thread performance of host machine
 - Need to use all cores of the host machine
 - Parallel simulation

DEMANDS ON SIMULATION ARE INCREASING

- Increasing cache size
 - Need a large working set to fully exercise a large cache
 - Scaled-down app behavior
 - Long-running sim



UPCOMING CHALLENGES

- Future systems will be diverse
 - Varying processor speeds
 - Varying failure rates for different components
 - Homogeneous applications become heterogeneous
- Software and hardware solutions are needed to solve these challenges
 - Handle heterogeneity (reactive load balancing)
 - Be fault tolerant
 - Improve power efficiency at the algorithmic level (extreme data locality)
- Hard to model accurately with analytical models

FAST AND ACCURATE SIMULATION IS NEEDED

- Simulation use cases
 - Architecture exploration
 - Pre-silicon software optimization
 - [Validation]
- Cycle-accurate simulation is too slow for exploring multi/many-core design space and software
- Key questions
 - Can we raise the level of abstraction?
 - What is the right level of abstraction?
 - When to use these abstraction models?

EXPERIMENT DESIGN IN ARCHITECTURE EXPLORATION/EVALUATION


- Optimizing the probability of success (i.e., finding the best architecture/parameters):
 - Coverage: how many architecture configurations can I run
 - Confidence: # benchmarks, re-runs for variable applications
 - Accuracy: simulation model detail vs. runtime
- How many scenarios can I run?
 - N = total number of simulation scenarios
 - d = days until paper deadline
 - t = average time per simulation
 - B = number of benchmarks
 - A = number of architectures

$$N = \frac{d}{t \cdot B \cdot A}$$


minimize t to maximize N

FAST OR ACCURATE SIMULATION?

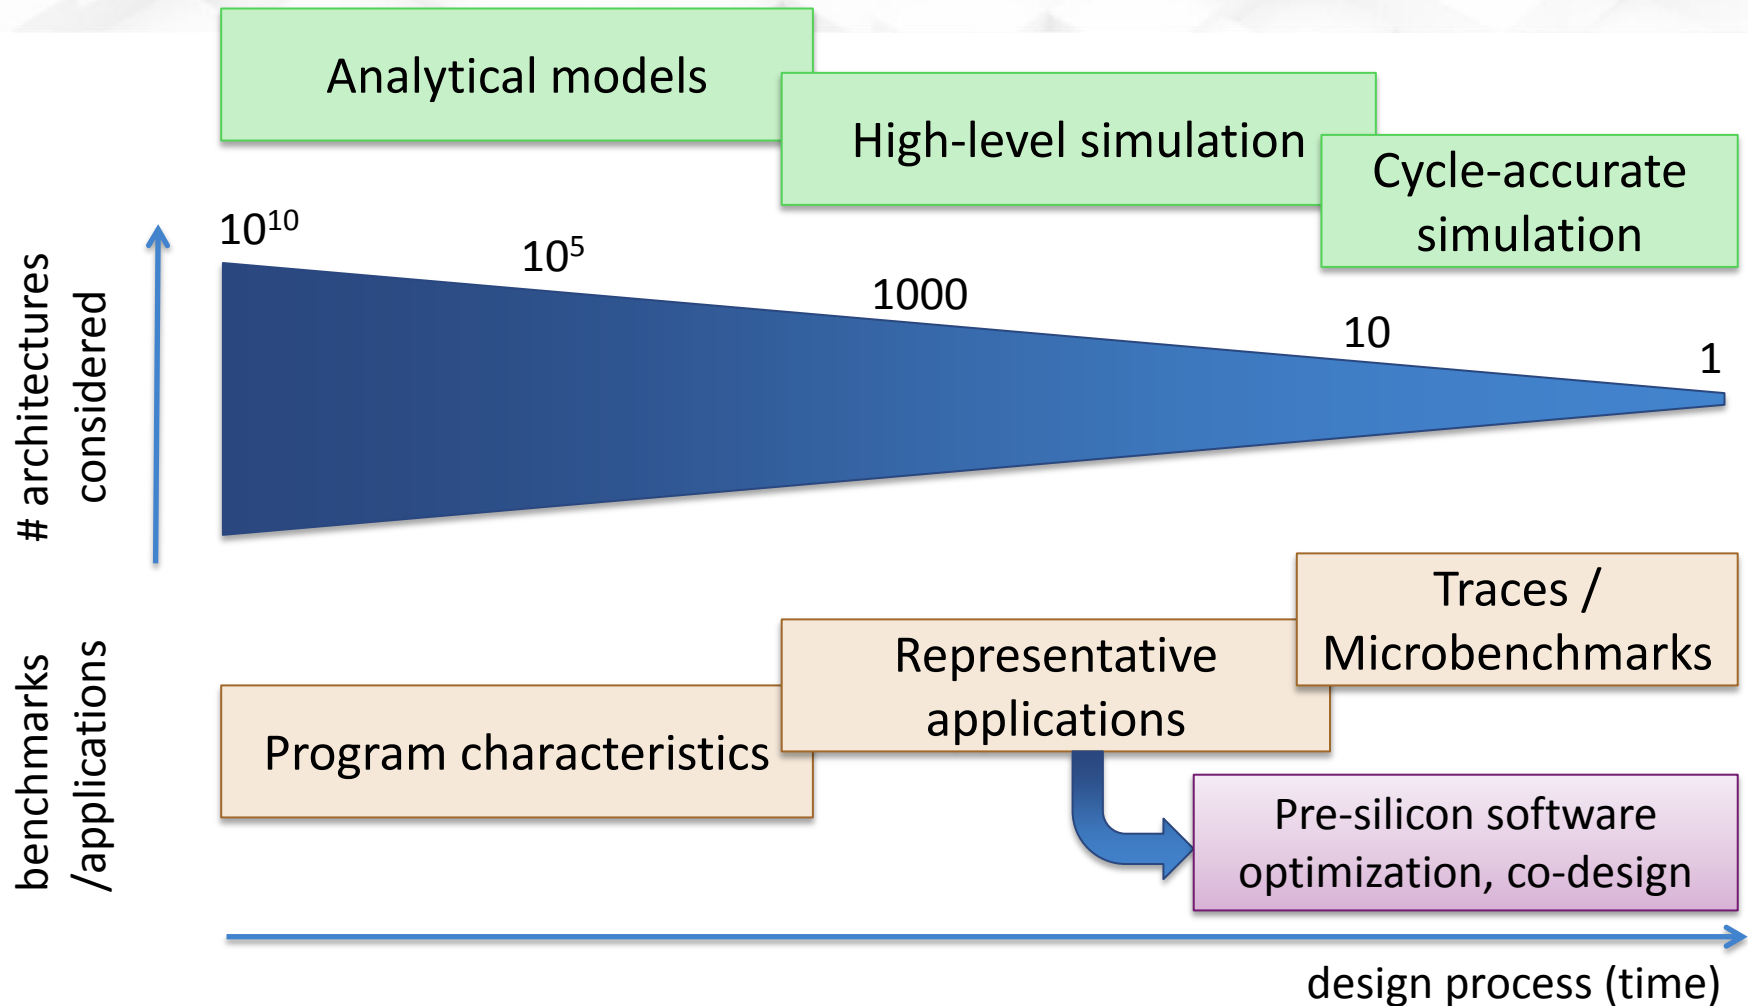
Cycle-accurate simulator



Higher-abstraction level simulator




THE ARCHITECTURE DESIGN WATERFALL



NEEDED DETAIL DEPENDS ON FOCUS

Component	Single-event time scale	Required sim time
RTL	single clock cycle	millions of cycles
OOO execution		
Core memory ops		
L1 cache access		
LLC access		
Off-socket	microseconds	seconds

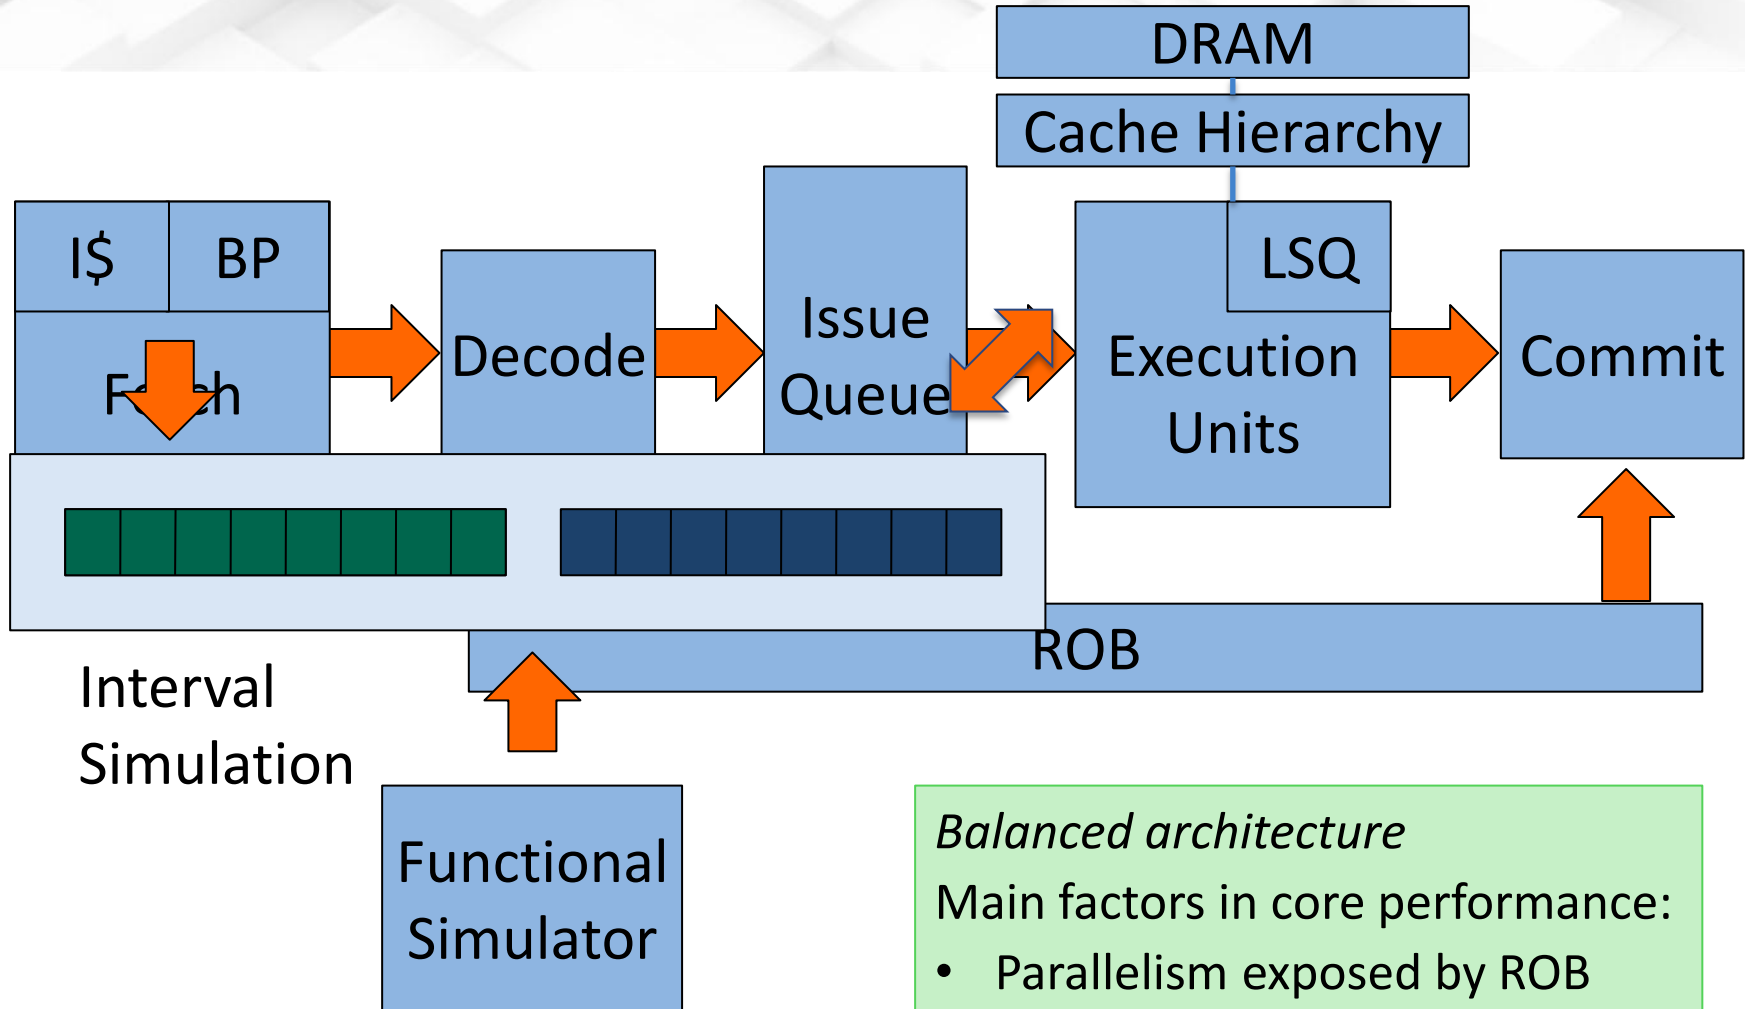


SNIPER: A FAST AND ACCURATE SIMULATOR

- Hybrid simulation approach
 - Analytical interval core model
 - Micro-architecture structure simulation
 - branch predictors, caches (incl. coherency), NoC, etc.
- Hardware-validated, Pin-based
- Models multi/many-cores running multi-threaded and multi-program workloads
- Parallel simulator scales with the number of simulated cores
- Available at <http://snipersim.org>



DETAILED MODEL VS. INTERVAL SIM




Balanced architecture

Main factors in core performance:


- Parallelism exposed by ROB
- Miss events

INTERVAL SIMULATION

Out-of-order core performance model with in-order simulation speed




INTERVAL SIMULATION FROM 30,000 FEET



Interval simulation considers instructions (in-order) at *dispatch*

- dispatch not possible
 - Instruction cache / TLB miss
 - Branch misprediction (not dispatching *useful* instructions)
 - Front-end refill after misprediction
 - ROB full: long-latency miss at head of ROB

INTERVAL SIMULATION FROM 30,000 FEET





Interval simulation considers instructions (in-order) at *dispatch*

- dispatch not possible
- dispatch possible: at rate governed by ROB
 - Little's law: progress rate = #elements / time spent in queue
 - Computed using ROB fill and critical path through ROB
 - Computed using dynamic instruction dependencies and latencies


CYCLE STACKS

- Where did my cycles go?
- CPI stack
 - Cycles per instruction
 - Broken up in components
- Normalize by either
 - Number of instructions (CPI stack)
 - Execution time (time stack)
- Different from miss rates:
cycle stacks directly quantify
the effect on performance



CYCLE STACKS AND SCALING BEHAVIOR

- Scaling to more cores, larger input set size
- How does execution time scale, and why?



VIZ: CYCLES STACKS IN TIME

Options

☐ Simple
 ☒ Normalized
 ☐ Detailed
 ☐ Absolute CPI

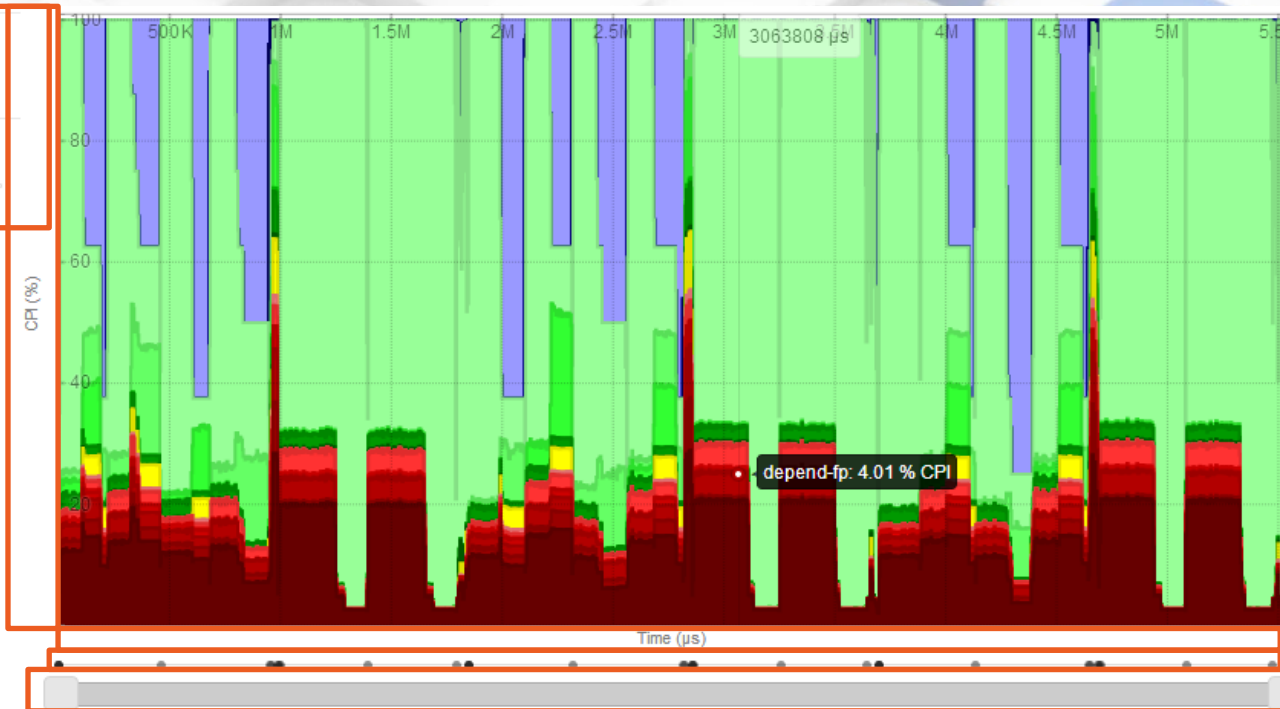
Smoothing

☒ Show IPC graph

Cycles (%)

Time

Markers




- ☒ imbalance-end
- ☒ imbalance-start
- ☒ sync-unscheduled
- ☒ sync-sleep
- ☒ sync-futex
- ☒ mem-dram
- ☒ mem-remote
- ☒ mem-l3
- ☒ mem-l2
- ☒ mem-l1d
- ☒ ifetch
- ☒ branch
- ☒ serial
- ☒ issue-port015
- ☒ issue-port5
- ☒ issue-port34
- ☒ issue-port2
- ☒ issue-port1
- ☒ issue-port0
- ☒ depend-branch
- ☒ depend-fp
- ☒ depend-int
- ☒ dispatch_width
- ☒ base

Legend


IPC Visualization

Slider


Smoothing



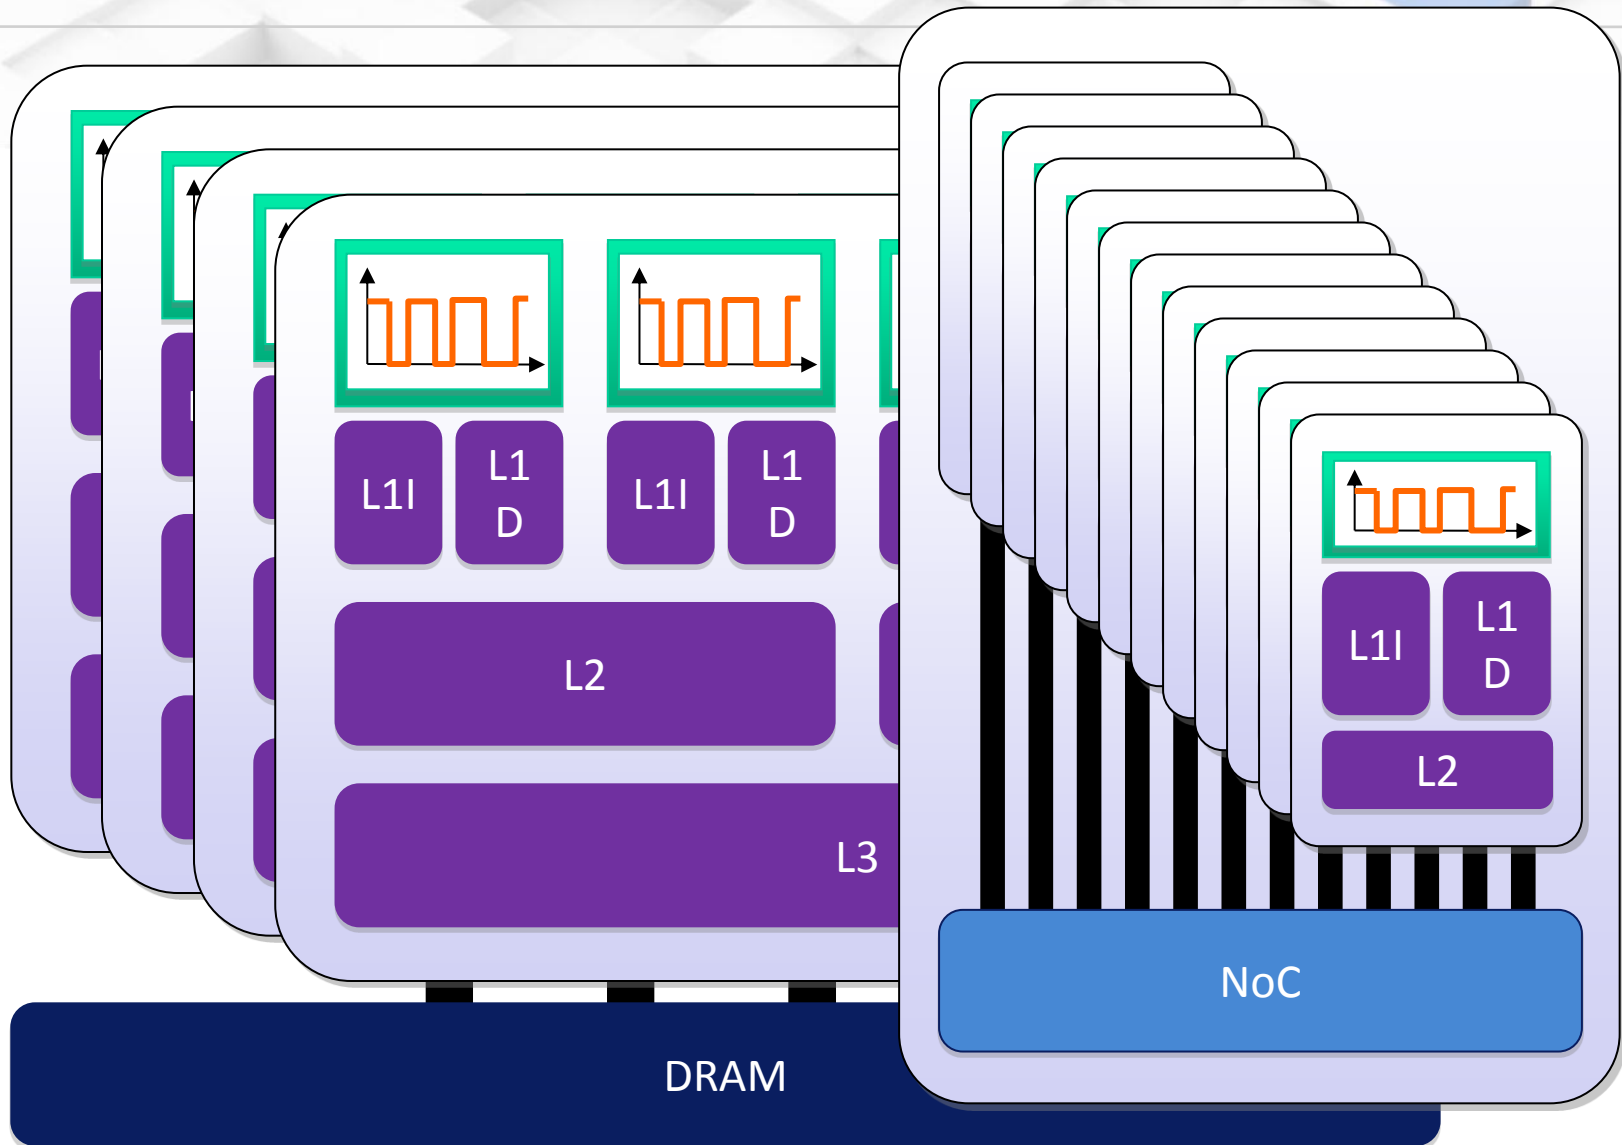
Viz: McPAT OUTPUT OVER TIME



3D VISUALIZATION: IPC vs. TIME vs. CORE



MANY ARCHITECTURE OPTIONS

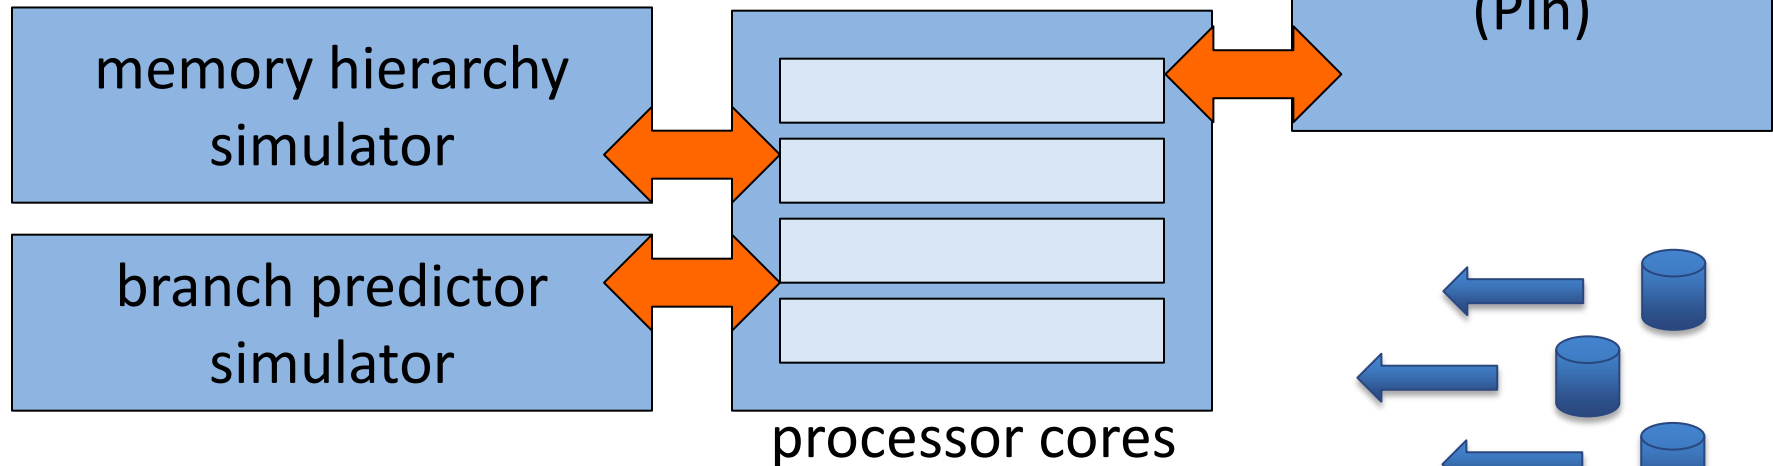


SIMULATION IN SNIPER

Execution-driven simulation


A single-process,
multithreaded
workload (v1.0)

functional
simulator
(Pin)



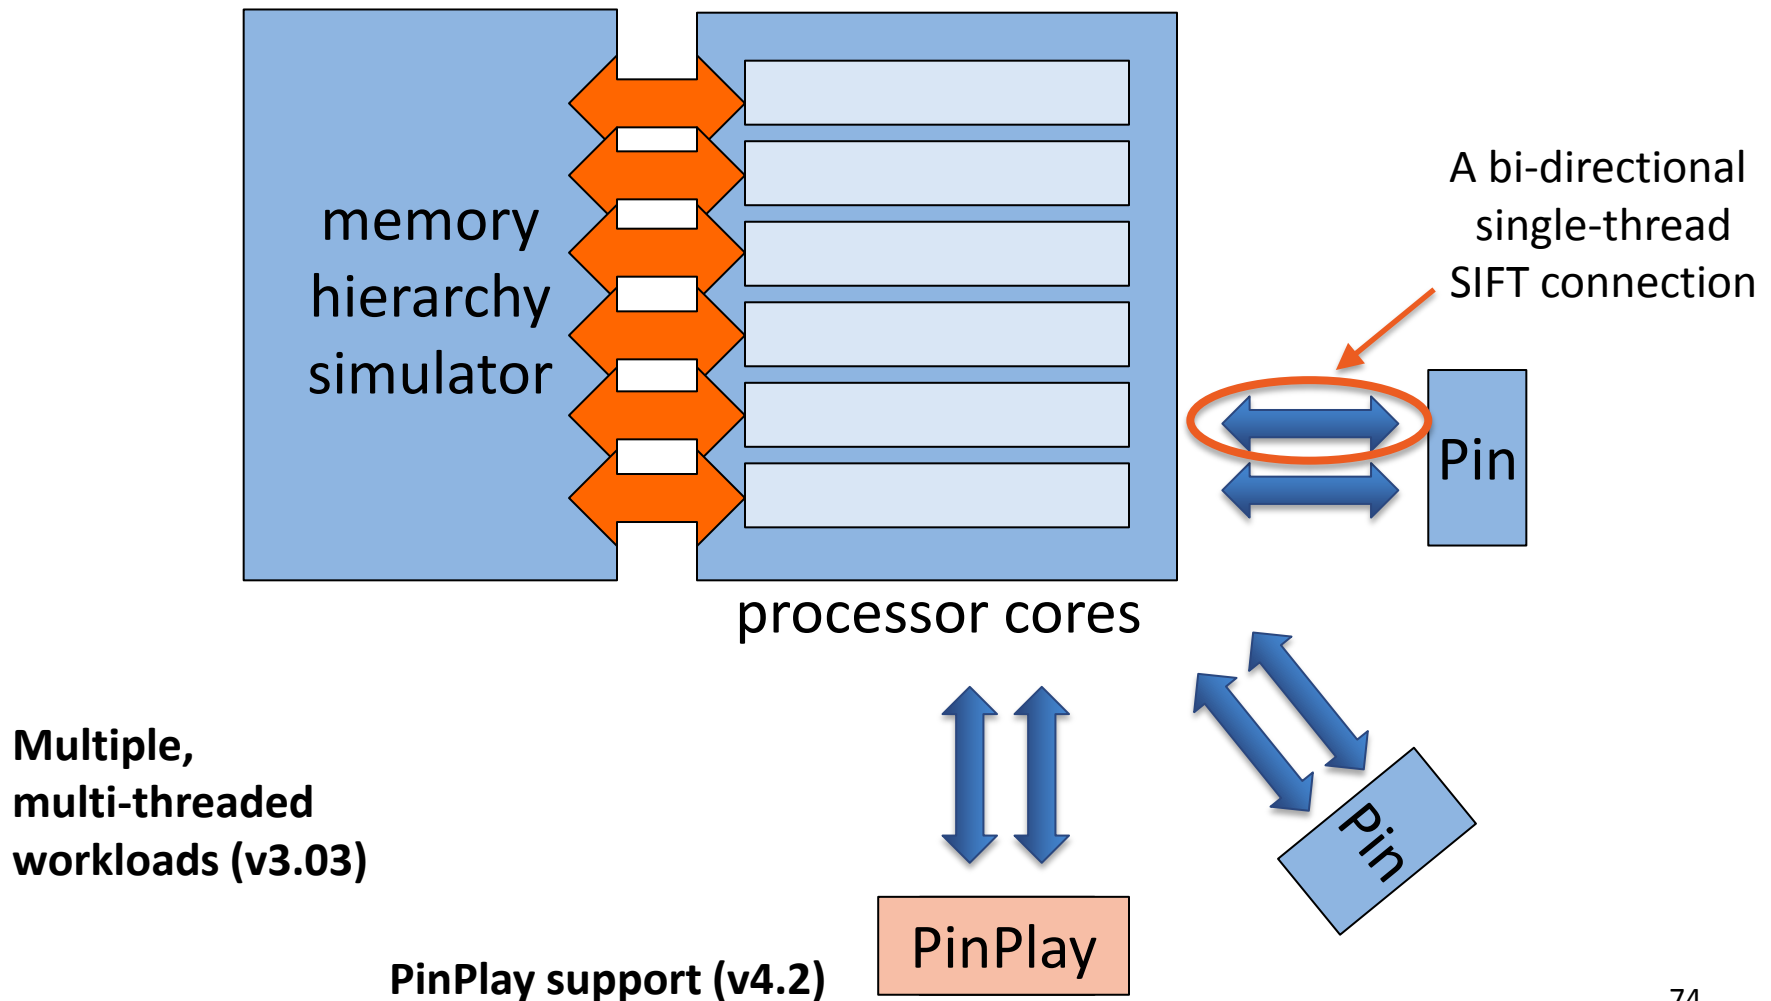
Trace-driven simulation

Multiple,
single-threaded
workloads (v2.0)




SIMULATION IN SNIPER WITH SIFT

Functional-directed simulation + timing-feedback



THREAD SCHEDULING AND MIGRATION



TOP SNIPER FEATURES

- Interval Model
- CPI Stacks and Interactive Visualization
- Parallel Multithreaded Simulator
- x86-64 and SSE2 support
- Validated against Core2, Nehalem
- Thread scheduling and migration
- Full DVFS support
- Shared and private caches
- Modern branch predictor
- Supports pthreads and OpenMP, TBB, OpenCL, MPI, ...
- SimAPI and Python interfaces to the simulator
- Many flavors of Linux supported (Redhat, Ubuntu, etc.)



SNIPER LIMITATIONS

- User-level
 - Perfect for HPC
 - Not the best match for workloads with significant OS involvement
- Functional-directed
 - No simulation / cache accesses along false paths
- High-abstraction core model
 - Not suited to model all effects of core-level changes
 - Perfect for memory subsystem or NoC work
- x86 only

SNIPER HISTORY

- November, 2011: SC'11 paper, first public release
- May 2012, version 3.0: Heterogeneous architectures
- November 2012, version 4.0: Thread scheduling and migration
- December 2012, version 4.1: Visualization (2D and 3D)
- February 2012, version 4.2: PinPlay support
- Today: 300+ downloads from 45 countries



Sniper PinPlay Integration



Sniper PinPlay Integration



- Integration of PinPlay with Sniper was a perfect match
 - Since v2.0, Sniper has supported multi-program execution
 - We support functional-directed simulation
 - Pin generates an instruction stream
 - Sniper consumes that stream, but doesn't allow Pin to get too far ahead
 - Critical synchronization instructions halt functional execution
 - The SIFT format stores dynamic instruction trace information
 - Trace sizes are $O(\text{instruction count})$
 - But Pinball data is executed directly; pinball sizes can be quite small $O(\text{binary-size} + \text{input-size})$
 - 100's MBs to GBs for a SIFT trace can result in just MBs for a pinball
- Goal to Integrate PinPlay into Sniper
 - Code to insert is straight-forward
 - But, we ran into a few problems along the way

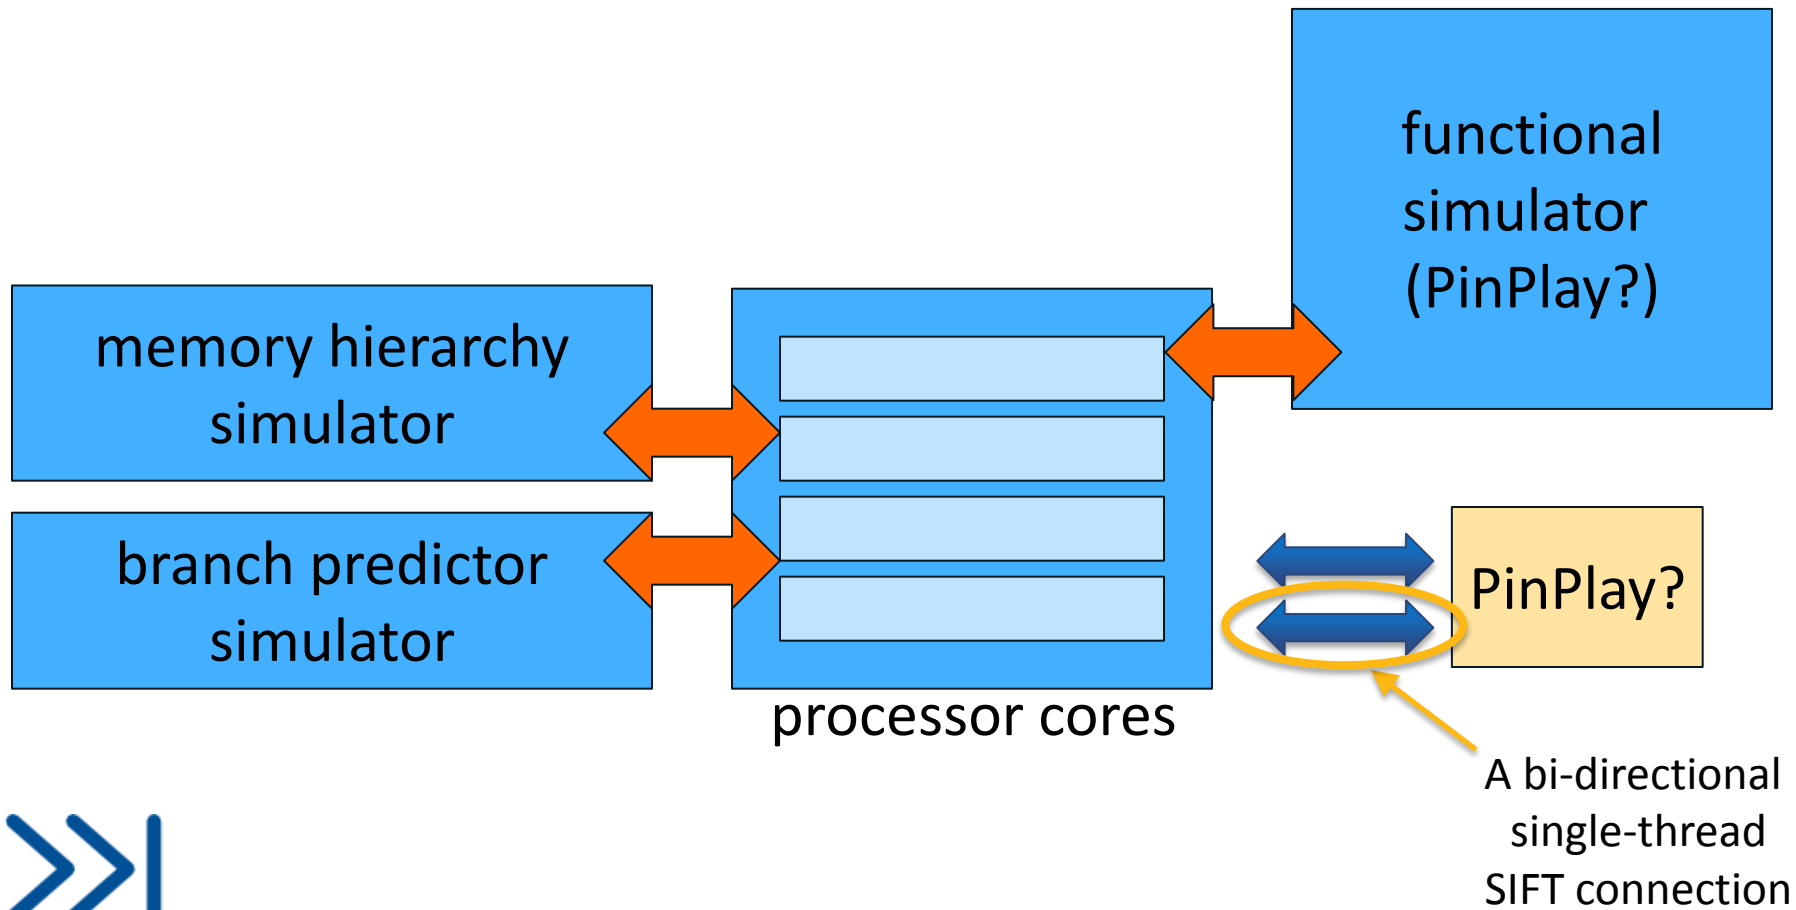
Sniper PinPlay Integration

- Would like to integrate PinPlay into Sniper
 - Code to insert is straight-forward
 - Initial goal was to integrate PinPlay into Sniper directly

```
#ifdef PINPLAY_SUPPORTED
# include "pinplay.H"
PINPLAY_ENGINE pinplay_engine;
#endif

bool pinplay_enabled = cfg->getBool("general/enable_pinplay");
#ifdef PINPLAY_SUPPORTED
    if (pinplay_enabled) {
        pinplay_engine.Activate(argc, argv, false /*logger*/, true /*replayer*/);
    }
#else
    if (pinplay_enabled) {
        LOG_PRINT_ERROR("PinPlay support not compiled in. Please use a compatible pin kit when compiling.");
    }
#endif
```

Initial Sniper PinPlay Integration



Sniper Integration Issue

- Sniper Pintool is slightly more complex than a typical Pintool
 - Per-basic-block callback with decoded instructions
 - Static information: list of micro-ops, loads, stores, memory accesses
 - Dynamic information: branch taken?, memory addresses, conditional execution status
 - Sniper also tries to control the state of running threads
 - Syscall interception and updates
 - gettimeofday() (returns simulated time)
 - get_nprocs() (returns number of simulated processors)
- Controlling the state of the thread is an issue with PinPlay

PinPlay Controls Execution

- Sniper would like to control execution
- PinPlay would like to control execution

```
$ ./run-sniper --pinballs=pin_kit/extras/pinplay/examples/pinball/foo
[SNIPER] Start
Running ['bash', '-c', ...]
[SNIPER] Enabling performance models
[SNIPER] Setting instrumentation mode to DETAILED
```

```
A:restore_mem.cpp:RestorePage:RestorePage:223:
```

```
Entry is not a page for thread 0 @icount: 5720 @mcount: 1635
```

```
#####
## STACK TRACE
#####
```

```
...
```

```
Pin 2.12
```

```
Copyright (c) 2003-2012, Intel Corporation. All rights reserved.
```

```
...
```

```
Pin app terminated abnormally due to signal 6.
```

```
[SNIPER] End
```

```
[SNIPER] Elapsed time: 1.53 seconds
```

PinPlay Controls Execution

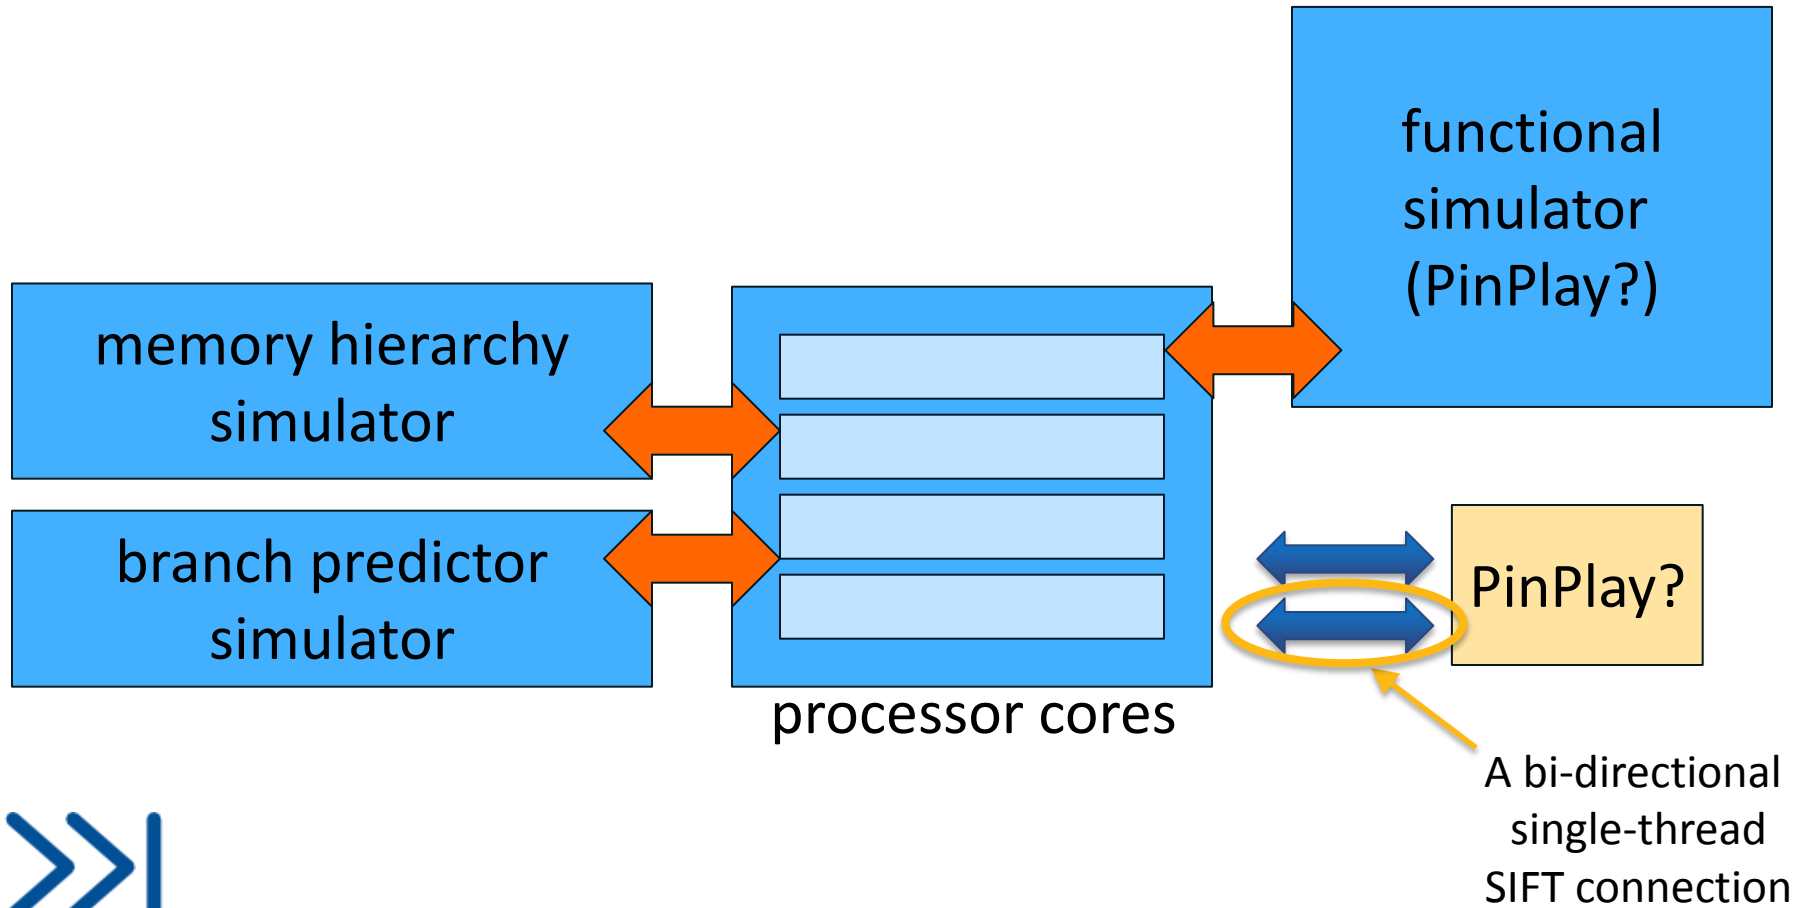
A:restore_mem.cpp:RestorePage:RestorePage:223:
Entry is not a page for thread 0 @icount: 5720 @mcount: 1635

- This is a PinPlay memory lookup error
- PinPlay needs to maintain control of the execution state of the application
 - The Pinball only contains memory, instruction data and syscalls that have been accessed during initial execution
 - Other data is not available, and therefore, changes in execution will cause PinPlay to return an error

Sniper Integration

- The end result took a different approach
 - Because of the deep integration with Pin, implementing PinPlay checks for these cases could be error prone
 - Instead, we implemented PinPlay support as a part of the SIFT trace generator (record-trace)
- Our trace recorder has a much lower impact on the execution traces
 - It supports command-line options to easily disable Linux syscall emulation
 - Behind the scenes, `run-sniper --pinballs=foo` automatically spawns a new record-trace process to generate a SIFT trace

Current Sniper-PinPlay Integration



Sniper PinPoints Evaluation



SPEC CPU2006 PinPoints Study

- We can use Sniper as a way to evaluate the accuracy of the Pinballs created using runtime as a metric
- Experiment with PinPoints inside of Sniper
 - Show results for a basic SimPoint configuration
 - Compare a matrix of results to improve outliers



SPEC CPU2006 PinPoints Study

- SPEC CPU2006
 - Reference inputs
 - 50 of the 55 benchmarks completed (5 full-simulations are still running)
- Environment
 - Compiled with GCC 4.3, -O2, SSE2 enabled
 - Sniper version 4.2, gainestown configuration, single core
 - PinPlay prerelease version



SPEC CPU2006 PinPoints Study

- Runtime
 - Whole-program pinball generation takes about a week per SPEC CPU2006 ref-input benchmark
 - PinPoints pinball generation takes an additional 1 to 5 days
 - Sniper Timing Simulation
 - 100M warmup + 30M detailed
 - 5 minutes
 - Full-reference inputs
 - Days to months
- Trace Sizes
 - Total of 3.8 GiB, from 2.7MiB (gameess) to 1.1GiB (milc)



SPEC CPU2006 CPI Error



30M detailed, 100M warmup, maxk=5

astar_1-ref-1	2.54	gcc_1-ref-1	0.30	gobmk-ref-1	1.12	perlbench_2-ref-1	2.24
astar-ref-1	8.11	gcc_2-ref-1	18.70	gromacs-ref-1	2.81	perlbench-ref-1	5.33
bzip2_1-ref-1	14.90	gcc_3-ref-1	9.37	h264ref_1-ref-1	1.03	povray-ref-1	2.25
bzip2_2-ref-1	0.57	gcc_4-ref-1	11.67	h264ref_2-ref-1	0.33	sjeng-ref-1	1.06
bzip2_3-ref-1	15.29	gcc_5-ref-1	1.66	h264ref-ref-1	2.36	soplex_1-ref-1	4.22
bzip2_4-ref-1	5.95	gcc_6-ref-1	3.14	hmmer_1-ref-1	0.04	soplex-ref-1	1.23
bzip2_5-ref-1	10.53	gcc_7-ref-1	16.26	hmmer-ref-1	0.08	sphinx3-ref-1	0.82
bzip2-ref-1	9.55	gcc_8-ref-1	9.23	lbm-ref-1	1.79	tonto-ref-1	0.20
calculix-ref-1	1.30	gcc-ref-1	9.79	mcf-ref-1	4.63	wrf-ref-1	14.27
dealII-ref-1	21.89	gobmk_1-ref-1	0.82	milc-ref-1	6.01	xalancbmk-ref-1	18.42
gamess_1-ref-1	0.56	gobmk_2-ref-1	1.09	namd-ref-1	31.79	zeusmp-ref-1	5.16
gamess_2-ref-1	2.15	gobmk_3-ref-1	0.76	omnetpp-ref-1	3.51		
gamess-ref-1	4.26	gobmk_4-ref-1	2.70	perlbench_1-ref-1	0.93	avg abs err	5.89

SPEC CPU2006 PinPoints Study

- Can we improve these timing error rates?
 - Compare a number of different maxK and warmup values

cpu2006-gcc_2-1, detailed = 30M insn		
	maxk 5	maxk 20
warmup 100M	18.70	6.99
warmup 300M	18.61	7.26



Sniper PinPoints Integration



Sniper PinPoints Integration

- The `pinpoints.py` script provides a basic infrastructure for creating PinPoints
- The `sniper_pinpoints.py` script enhances `pinpoints.py` to support automatic reference and PinPoint playback and comparison
- Provide an overview to extending the PinPoints scripts for your simulator



PinPoints Simulator Integration

- Register additional phases as options
- If those options have been selected, run the simulation

```
import pinpoint

class SniperPinPoints(pinpoints.PinPoints):
    def AddAdditionalOptions(self, parser):
        """Add additional options specific for Sniper"""

    def AddAdditionalPhaseOptions(self, parser, phase_group):
        """Add additional phase options specific for Sniper"""

    def RunAdditionalPhases(self, wp_log_dir, sim_replay_cmd, options):
        """Run additional phases specific for Sniper"""

        return 0

if __name__ == "__main__":
    sys.exit(SniperPinPoints().Run())
```


PinPoints Simulator Integration

- Register additional options and phases

```
import pinpoint, cmd_options

def sniper_root(parser):
    parser.add_option("--sniper_root", dest="sniper_root",
                      default=os.getenv("SNIPER_ROOT"), help="Sniper root")

def whole_sim(parser, group):
    method = cmd_options.GetMethod(parser, group)
    method("-W", "--whole_sim", dest="whole_sim", action="store_true",
           help="Run Sniper on the whole program pinballs.")

class SniperPinPoints(pinpoint.PinPoints):
    def AddAdditionalOptions(self, parser):
        sniper_root(parser)

    def AddAdditionalPhaseOptions(self, parser, phase_group):
        whole_sim(parser, phase_group)


    def RunAdditionalPhases(self, wp_log_dir, sim_replay_cmd, options):
        if options.whole_sim or options.default_phases:
            pass # Run the whole program simulation

    return 0
```

Handling Parallel Programs




Model 1: Parallel Capture : Parallel Replay For Multi-threaded Programs




Useful for parallel analysis/simulation
[Can focus on one thread with `-log:focus_thread`]

PinPlay's Determinism == Same Access Order for Conflicting Shared Memory Accesses




- Instructions from each thread replayed in program order
- RAW, WAR, WAW order for multiple threads is preserved

Relative Speed of Threads will Change During Replay



- Instructions from each thread replayed in program order
- RAW, WAR, WAW order for multiple threads is preserved

Model 2: **Parallel Capture : Isolated Replay** For Multi-process Programs



• multi-process → **multi-programmed**

Extending PinPoints to Parallel Programs [Work In Progress]

- **Multi-threaded Programs:**

- **Per thread pinball:** *-log:focus_thread tid*
Whole-program logging, PinPoints generation same as single-threaded program
- **Truly multi-threaded** ("co-operative") **pinball:**
 - Use basic block vector from thread 0 (or any specified thread)
 - Generate multi-threaded PinPoint pinballs
 - Caveat: Simulator cannot change thread order!
[Sniper runs into a problem trying to emulate futex()]

- **Multi-process Programs:** (per-process pinballs only)
 - Focus on one process pinball for PinPoints generation

Example: Multi-threaded PinPoints Generating Cooperative Pinballs

```
% pinpoint.py --pinplayhome=$PIN_KIT --cfg thread.coop.cfg --coop_pinball -l -b -s -p
```

```
% cat thread.coop.cfg
[Parameters]
program_name:    thread
input_name:      parallel
command:         "base.exe 3"
epilog_length:   0
maxk:            5
mode:            mt
num_proc:        1
prolog_length:   0
slice_size:      1000000
warmup_length:   3000000
```

```
% tree thread.1.parallel_14134.pp/ | grep t0r1 | grep race
|-- thread.1.parallel_14134_t0r1_warmup3001500_prolog0_region215307_epilog0_001_1-000000.0.race.bz2
|-- thread.1.parallel_14134_t0r1_warmup3001500_prolog0_region215307_epilog0_001_1-000000.1.race.bz2
|-- thread.1.parallel_14134_t0r1_warmup3001500_prolog0_region215307_epilog0_001_1-000000.2.race.bz2
|-- thread.1.parallel_14134_t0r1_warmup3001500_prolog0_region215307_epilog0_001_1-000000.3.race.bz2
```

- PinPoints pinballs:
 - truly multi-threaded (co-operative) for 1 (main) + 3 (created) threads
 - **Will not work with Sniper** currently

Example: Multi-threaded PinPoints Generating Per-thread Pinballs [thread 2]

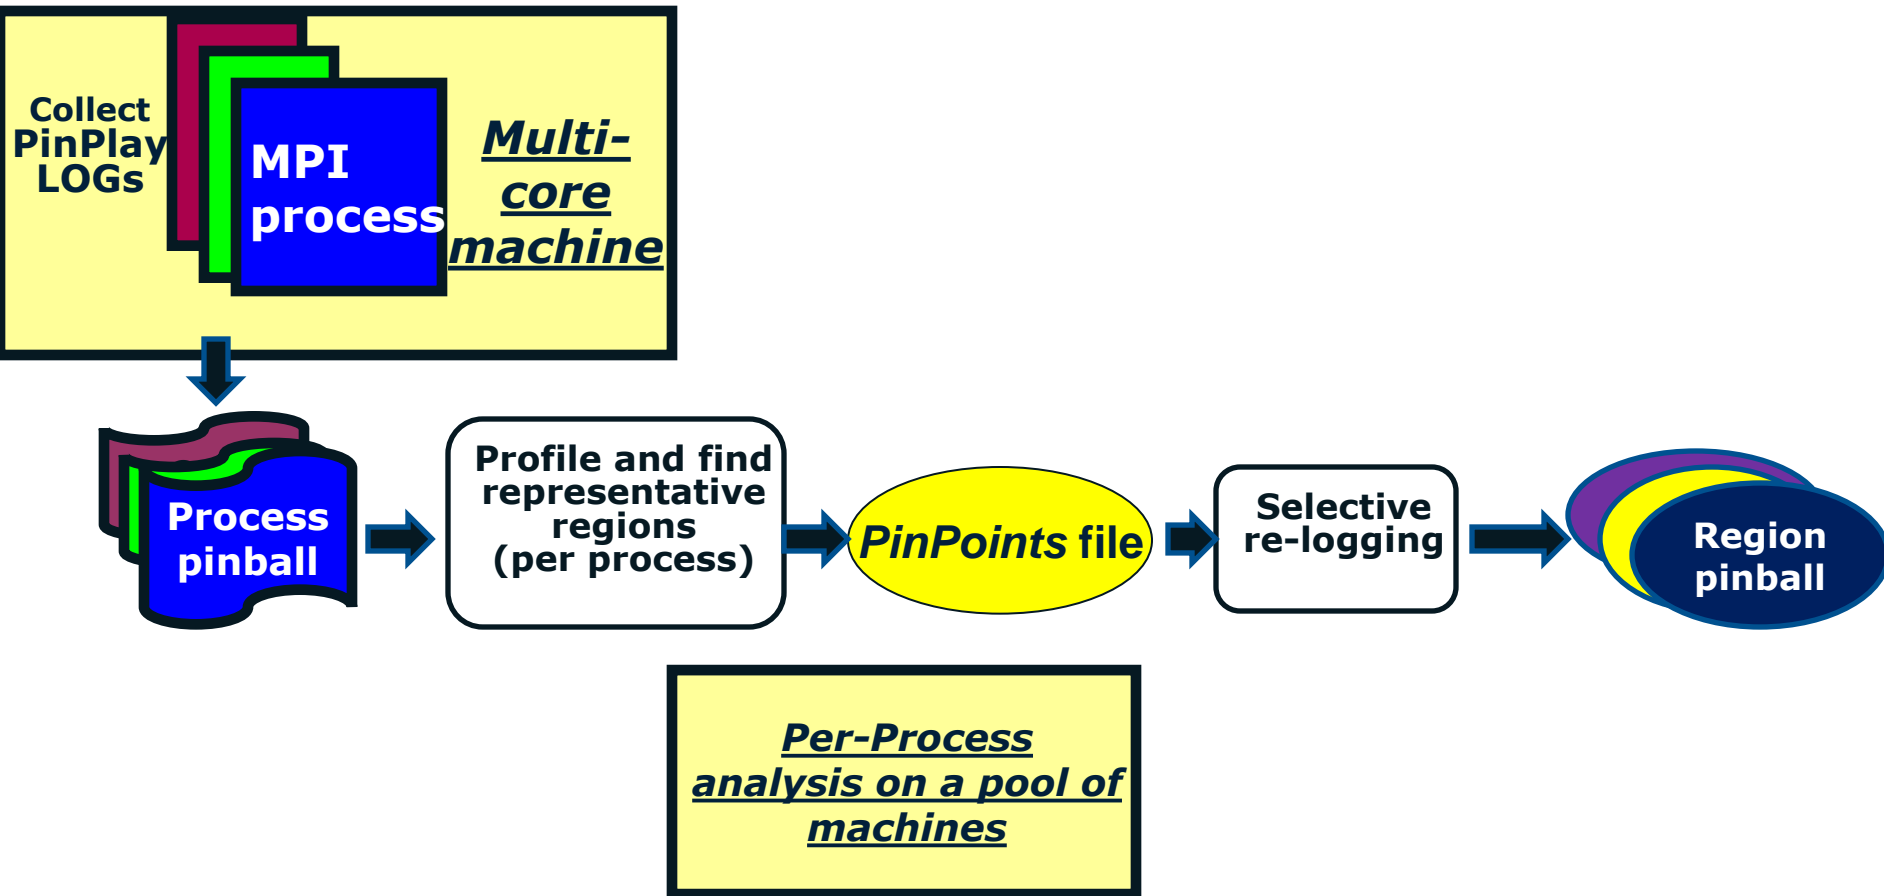
```
% pinpoints.py --pinplayhome=$PIN_KIT --cfg thread.perthread.cfg --focus_thread=2 -l -b -s -p
```

```
% cat thread.perthread.cfg ~
[Parameters]
program_name:      thread
input_name:        per-thread
command:           "base.exe 3"
epilog_length:     0
maxk:              5
mode:              mt
num_proc:          1
prolog_length:     0
slice_size:        1000000
warmup_length:     3000000
```

```
% tree thread.1.per-thread_22112.pp/ | grep t2r1 | grep race
|-- thread.1.per-thread_22112_t2r1_warmup0_prolog0_region1000002_epilog0_001_1-000000.2_race.bz2
```

- PinPoints pinballs:
 - Single-threaded
 - **Will work with Sniper** currently

MPI PinPoints Process Flow: Parallel Logging : Isolated Replay



Example: Intel MPI program with 2 processes (milc)

```
% $PIN_KIT/extras/pinplay/PinPoints/scripts/pinpoints.py --pinplayhome  
=$PIN_KIT --cfg milc.mpi.cfg -l -b -s -p | & tee out
```

```
% cat milc.mpi.cfg  
[Parameters]  
program_name:    milc  
input_name:      test  
command:         "base.exe test < /dev/null"  
epilog_length:   0  
maxk:            5  
mode:            mpi  
num_proc:        2  
prolog_length:   0  
slice_size:      1000000  
warmup_length:   3000000
```

```
% tree whole_program.test | grep address  
|-- milc.2.test_22555.address  
|-- milc.2.test_22556.address
```

```
milc/  
|-- milc.2.test_22555.Data  
|-- milc.2.test_22555.pp  
|-- milc.2.test_22556.Data  
|-- milc.2.test_22556.pp  
`-- whole_program.test
```

- num_proc:2 → "mpirun -n 2" was used
2 whole-program pinballs for 2 processes
PinPoints created for 2 processes per-process

Other Multi-process Models

1. A program creating multiple processes

- Parent forks children → Children fork grand-children
- Entire family tree communicate using inter-process shared memory

2. Multiple programs/processes interacting:

- Two independently invoked programs
- Communicate using inter-process shared memory

<..>/PinPoints/scripts/*.py do not currently support these models yet

Example: Program with fork()

```
% $PIN_KIT/pin -t $PIN_KIT/extras/pinplay/bin/intel64/pinplay-driver.so -log -log:mp_mode -log:basename pinball/foo -log:pid -- fork_app
APPLICATION: Before fork
APPLICATION: After fork in parent
APPLICATION: After fork in child
```

```
pinball/
|-- fk0.fork_app
```

```
|-- fk0.fork_app
|   |-- foo_18786_18793.0.race
|   |-- foo_18786_18793.0.reg
|   |-- foo_18786_18793.0.result
|   |-- foo_18786_18793.address
|   |-- foo_18786_18793.procinfol.xml
|-- foo_18786.0.race
|-- foo_18786.0.reg
|-- foo_18786.0.result
|-- foo_18786.address
|-- foo_18786.procinfol.xml
```

```
% $PIN_KIT/pin -xyzzzy -reserve_memory pinball/foo_18786.address -t $PIN_KIT/extras/pinplay/bin/intel64/pinplay-driver.so -replay -replay:basename pinball/foo_18786 -- $PIN_KIT/extras/pinplay/bin/intel64/nullapp
APPLICATION: Before fork
APPLICATION: After fork in parent
```

```
% $PIN_KIT/pin -xyzzzy -reserve_memory pinball/fk0.fork_app/foo_18786_18793.address -t $PIN_KIT/extras/pinplay/bin/intel64/pinplay-driver.so -replay -replay:basename pinball/fk0.fork_app/foo_18786_18793 -- $PIN_KIT/extras/pinplay/bin/intel64/nullapp
APPLICATION: After fork in child
```

One pinball per process

Recipe for Independent MP Logging

Generic model:

```
%pin-t pinplay-driver.so -log:mp_create_pool -- <any_valid_binary>
```

```
% pin -t pinplay-driver.so -log -log:basename pinball1/foo -log:mp_mode -log:mp_attach --  
<program1>
```

```
%pin -t pinplay-driver.so -log -log:basename pinball2/bar -log:mp_mode -log:mp_attach --  
<program2>
```

Client/Server model (only Client to be traced):

```
%pin -t pinplay-driver.so -log -log:basename pinball1/foo -log:mp_mode_lock_only --  
<server_program_to_be_ignored>
```

```
% pin -t pinplay-driver.so -log -log:basename pinball2/bar -log:mp_mode -log:mp_attach --  
<client_program_of_interest>
```

Summary

PinPoints (Pin + SimPoint + Sniper) methodology effectively automates the tedious task of finding and check-pointing regions of programs for Pin-based simulation.*

- *Creates check-points that are representative and repeatable*

<http://www.pinplay.org> <http://www.snipersim.org>

References

Pin: [Pin: Building Customized Program Analysis Tools with Dynamic Instrumentation](#); Chi-Keung Luk, Robert Cohn, Robert Muth, Harish Patil, Artur Klauser, Geoff Lowney, Steven Wallace, Vijay Janapa Reddi, and Kim Hazelwood. *Proceedings of the 2005 ACM SIGPLAN conference on Programming language design and implementation*.

PinPoints: [Pinpointing Representative Portions of Large Intel® Itanium® Programs with Dynamic Instrumentation](#); Patil, H., Cohn, R., Charney, M., Kapoor, R., Sun, A., and Karunanidhi, A. In *Proceedings of the 37th Annual IEEE/ACM international Symposium on Microarchitecture* (Portland, Oregon, December 04 - 08, 2004).

PinPlay: [PinPlay: A Framework for Deterministic Replay and Reproducible Analysis of Parallel Programs](#); Harish Patil, Cristiano Pereira, Mack Stallcup, Gregory Lueck, James Cownie. CGO 2010. **CGO 2010 Best Paper Award Winner!**

SimPoint : [Automatically Characterizing Large Scale Program Behavior](#); Timothy Sherwood, Erez Perelman, Greg Hamerly and Brad Calder. In *proceedings of the 10th International Conference on Architectural Support for Programming Languages and Operating Systems*, October 2002.

Sniper: [Sniper: Exploring the Level of Abstraction for Scalable and Accurate Parallel Multi-Core Simulation](#); Trevor E. Carlson; Wim Heirman; Lieven Eeckhout. In *proceedings of International Conference for High Performance Computing, Networking, Storage and Analysis (SC)*, 2011.

谢谢

xièxiè
Thank you!

