

PinPoints: A methodology for *Simulation Region Selection, Validation, and Simulator Staging* with Intel Pin and SDE

Harish Patil

Principal Engineer, Development Tools Software, Intel Corporation

With input from Alen Sabu (National University of Singapore)

November 7th, 2024

Last modified: 14th November 2024

<https://github.com/intel/pinplay-tools/tree/main/PinPoints>

Architecture Simulation : Key Questions

Where to simulate?

Regions of Interest (ROIs)
($\ll 0.1$ % of whole-program)
NOT whole-program

How to simulate?

Drive simulations for ROIs

Are the regions representative?

Region Selection Validation

PinPoints Methodology

Where to simulate?

SDE-profiler + SimPoint

Region
Selection

How to simulate?

1. Checkpoint (pinball) driven
2. Binary-driven
 - Pinballs → ELFies

Simulator
Staging

Are the regions representative?

ROIPerf:
CPU time-stamps with a Pin-probe tools

Region
Validation

What is a “Region”?



Region specification: *Icount* vs *Pccount*

- *ICount*: instruction count based

Region Start: *icount1*
Region End: *icount2* (or *length* after Start)

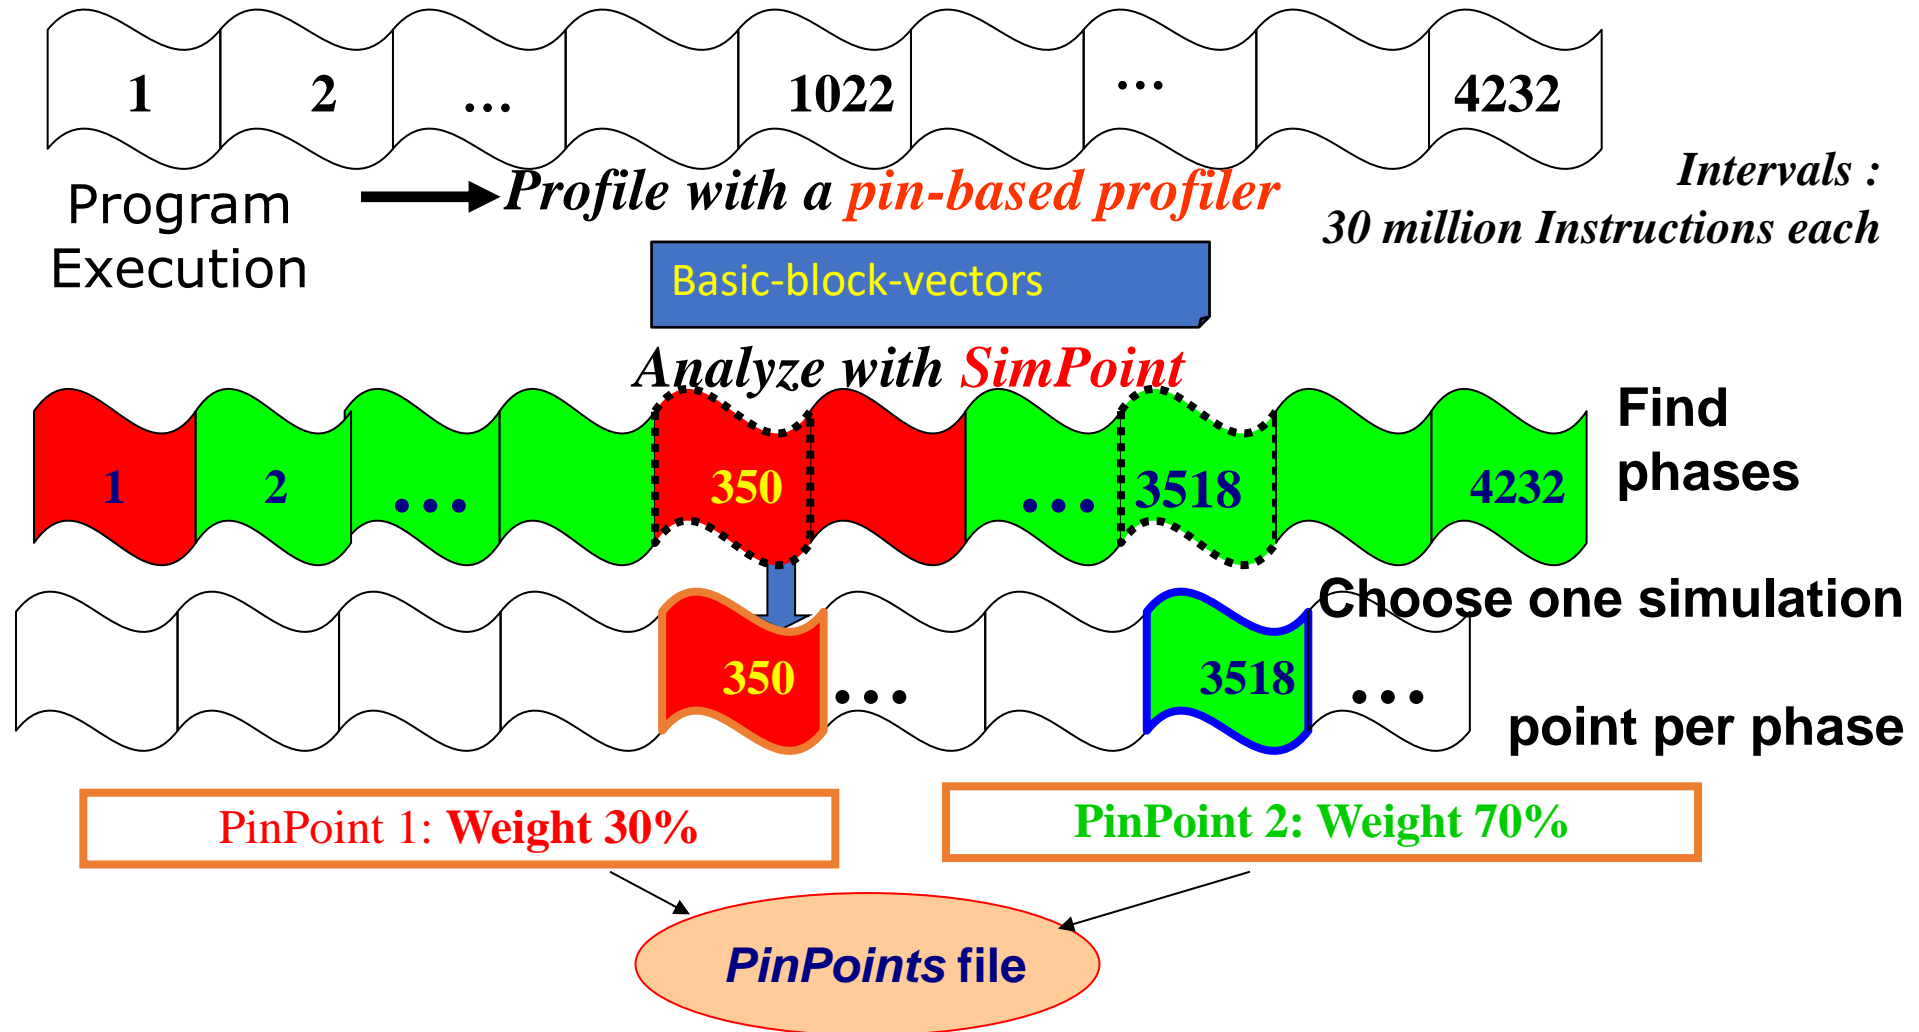
Drawbacks:

1. Different tools count instructions differently (e.g. REP-prefixed instructions as 1 or N)
2. (Multi-threaded programs) Instruction counts can change widely from run to run

- *PCcount*: program-counter (pc) based (select marker PCs with invariant counts across runs)

Region Start: *PC1+count1*
Region End: *PC2+count2* (relative to Start)

PinPoints = Pin + SimPoint (UC San Diego)



Two Phases => Two PinPoints

Projection Formulas

1. With weights (**ICount regions**) : ST only

$$\text{CPI}_{\text{predicted}} = \sum \text{Weight}_{\text{region}} * \text{CPI}_{\text{region}}$$

$$\text{Cycles}_{\text{predicted}} = \text{CPI}_{\text{predicted}} \times \text{Total_Instruction_Count (pathlength)}$$

2. With multipliers (**PCCount regions**) :ST and MT (no instruction counts used directly)

$$\text{Cycles}_{\text{predicted}} = \sum \text{Cycles}_{\text{region}} * \text{Multiplier}_{\text{region}}$$

where

$$\text{Multiplier}_j = \text{inscount}_j / \sum_{i=0..m} \text{inscount}_i$$

for m, the number of regions that are represented by the j^{th} representative region

Agenda

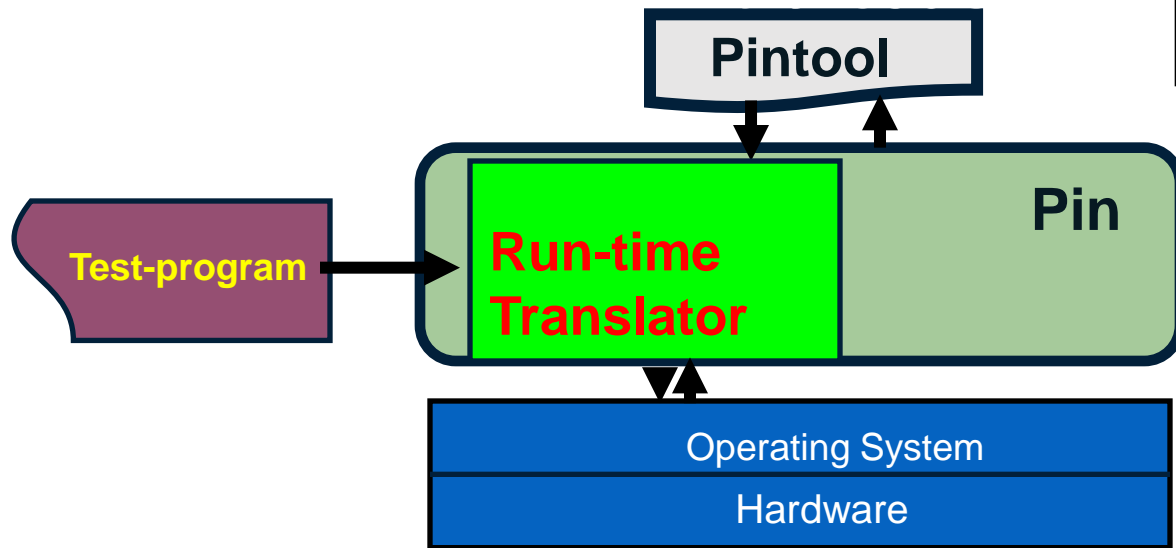
1. Tools Overview
2. Example/Demo
3. Discussion

Tools Overview

Pin: A Tool for Writing Program Analysis Tools

```
sub    $0xff, %edx
movl   0x8(%ebp), %eax
jle    <L1>
```

```
counter++; print(IP)
sub    $0xff, %edx
counter++; print(EA)
movl   0x8(%ebp), %eax
counter++; print(br taken)
jle    <L1>
```



```
$ pin -t pintool -- test-program
```

Normal output
+ *Analysis*
output

Pin: A Dynamic Instrumentation Framework from Intel
<http://www.pintool.org> (or search "Intel Pintool"

JIT Mode vs Probe Mode

- JIT Mode

- Pin creates a modified copy of the application on-the-fly
- Original code never executes

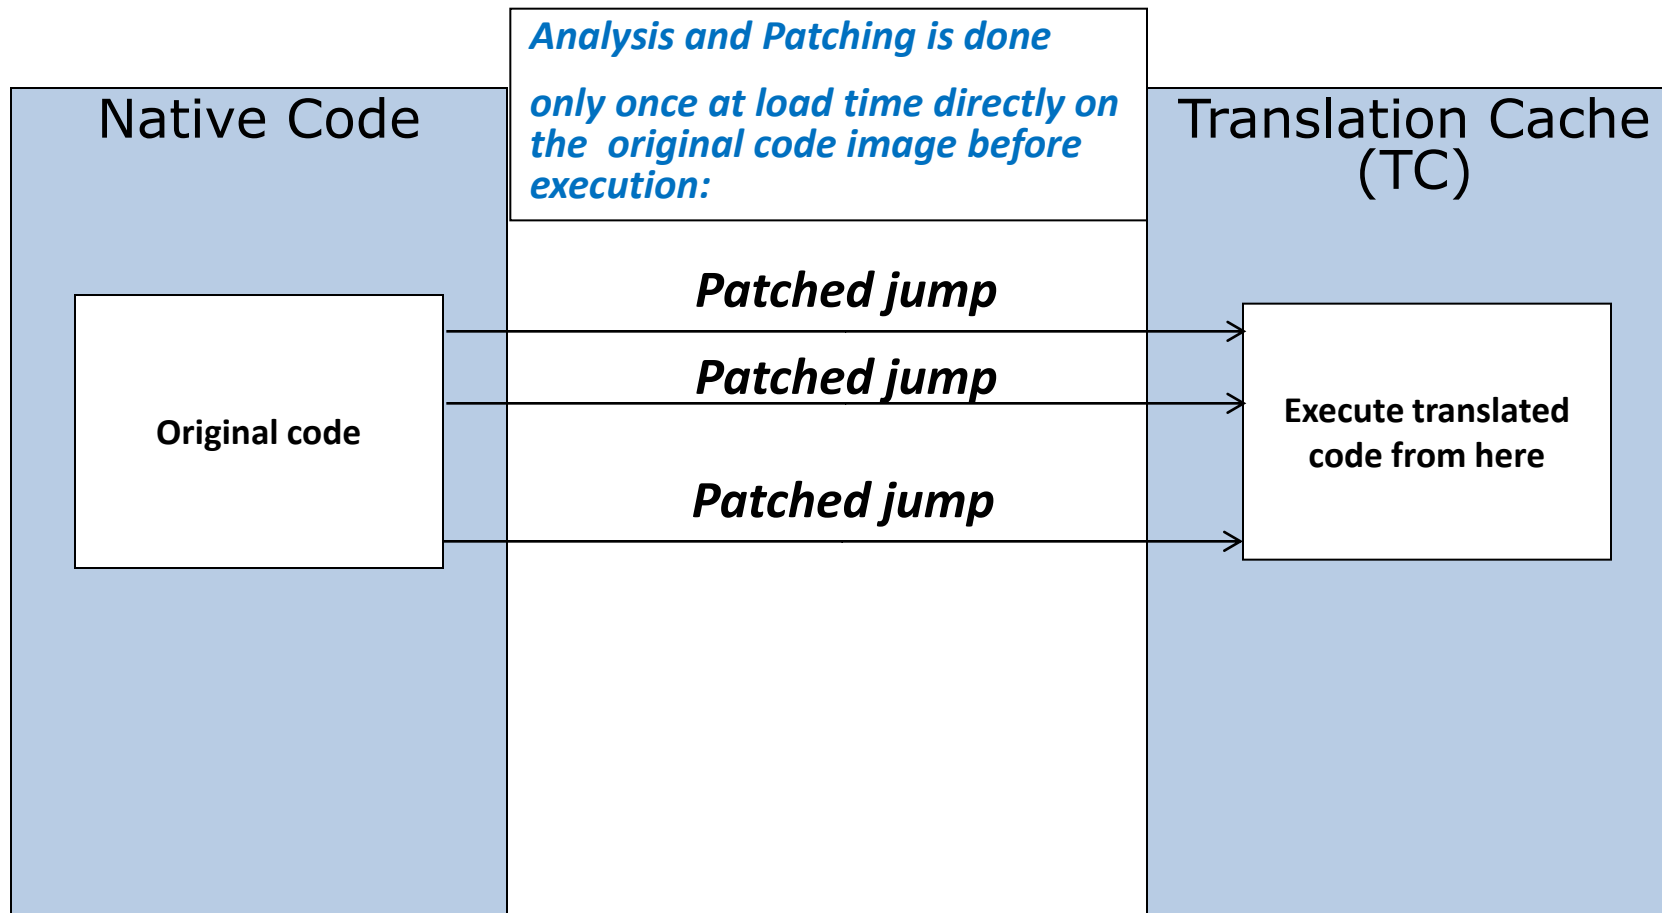
➤ More flexible, more common approach

- Probe Mode

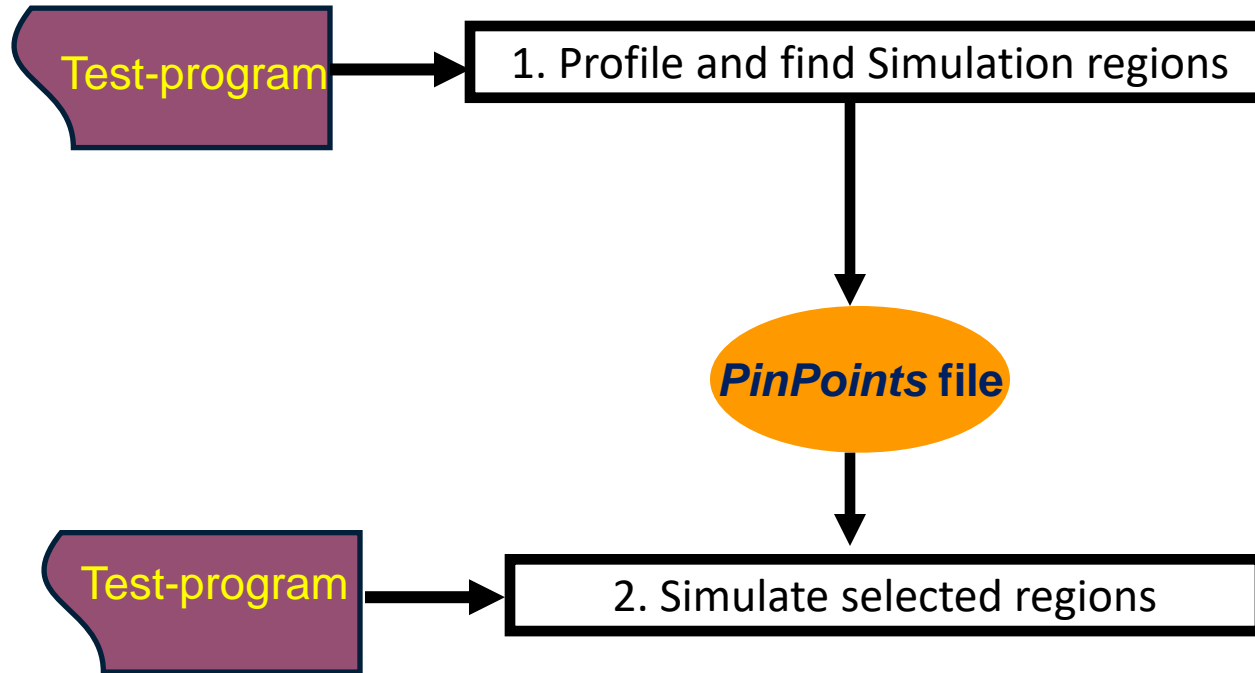
- Pin modifies the original application instructions
- Inserts jumps to instrumentation code (trampolines)

➤ Lower overhead (less flexible) approach

Pin Probe Mode – General Flow



PinPoints : The Repeatability Challenge

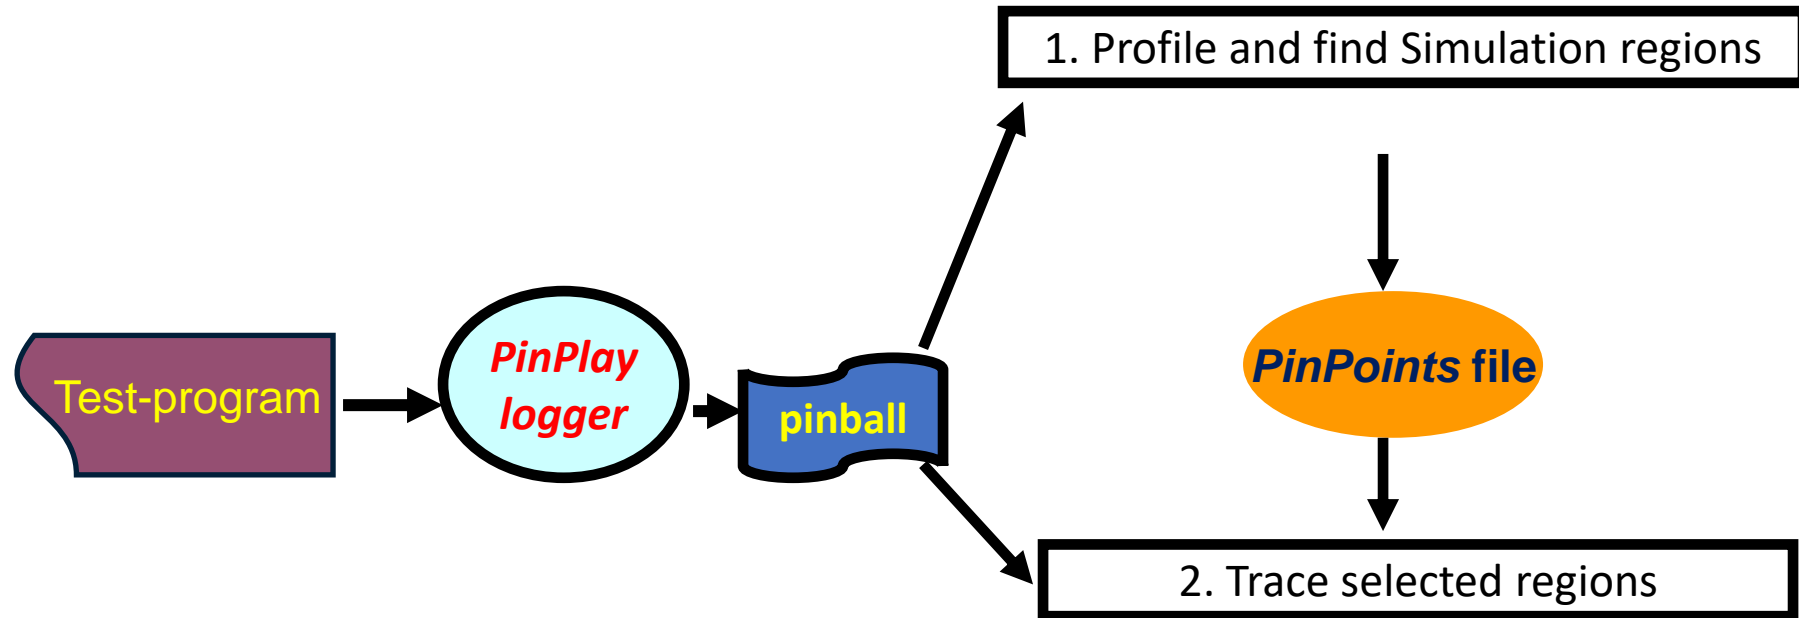


Problem: Two runs are not exactly same → PinPoints missed

Found this for 25/54 SPEC2006 runs!

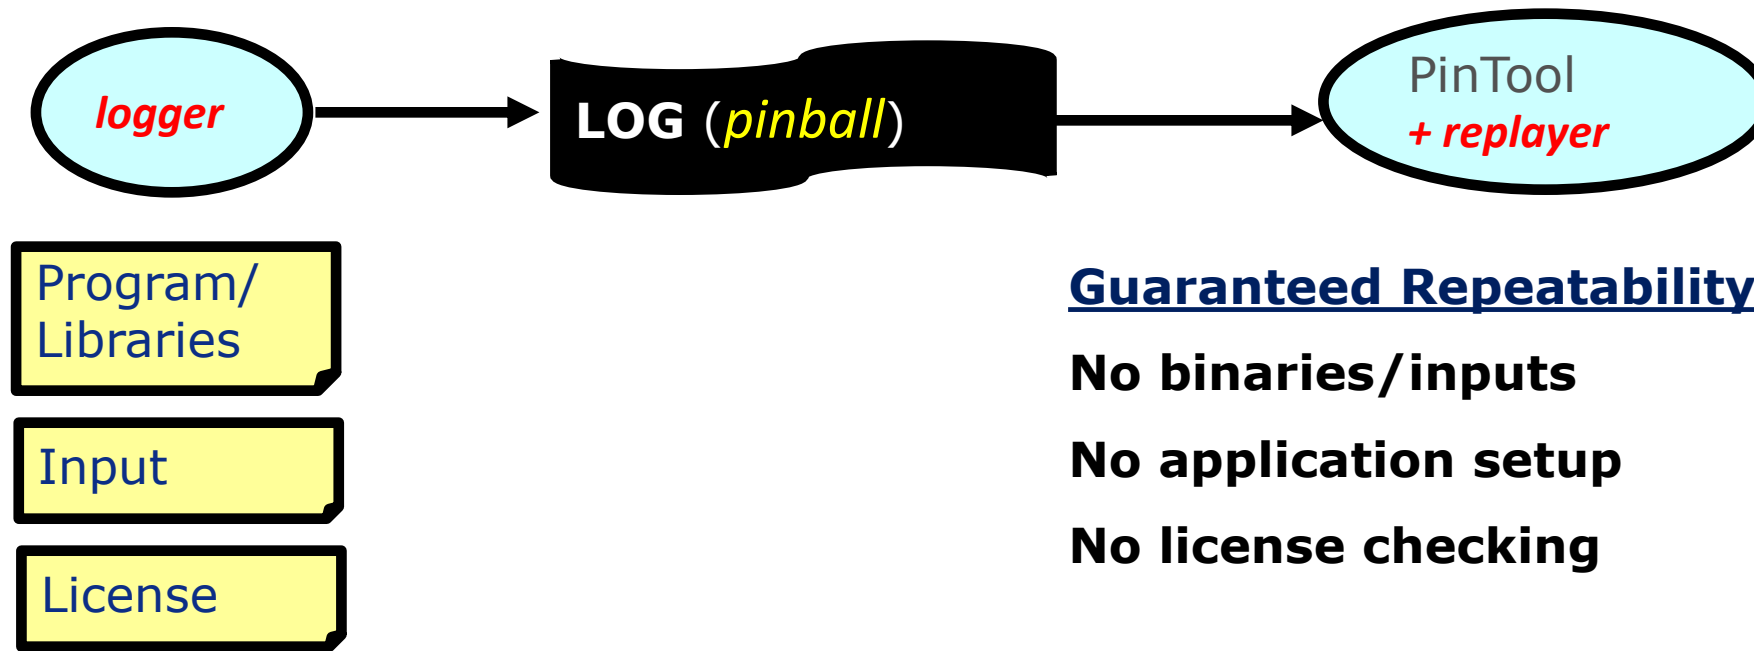
[*"PinPoints out of order" "PinPoint End seen before Start"*]

Enters PinPlay To Provide Repeatability



Two runs are same → PinPoints guaranteed to be reached

PinPlay*: Workload Capture and Deterministic Replay Framework



Record Once Replay + Analyze Multiple Times Anywhere!

PinPlay: Included in SDE

SDE : Pin + Emulation + PinPlay

- **SDE driver** runs with a default tool (sde-mix.so)
 - Emulation, log/replay, loop generation (DCFG json files), checkers.....
- **SDE Kit** also allows people to write tools
 - Regular Pin tools + use of additional SDE API
 - Can use emulation, log/replay, DCFG....

SDE : Intel Software Development Emulator
(search "Intel SDE")

Direct pinball generation with SDE

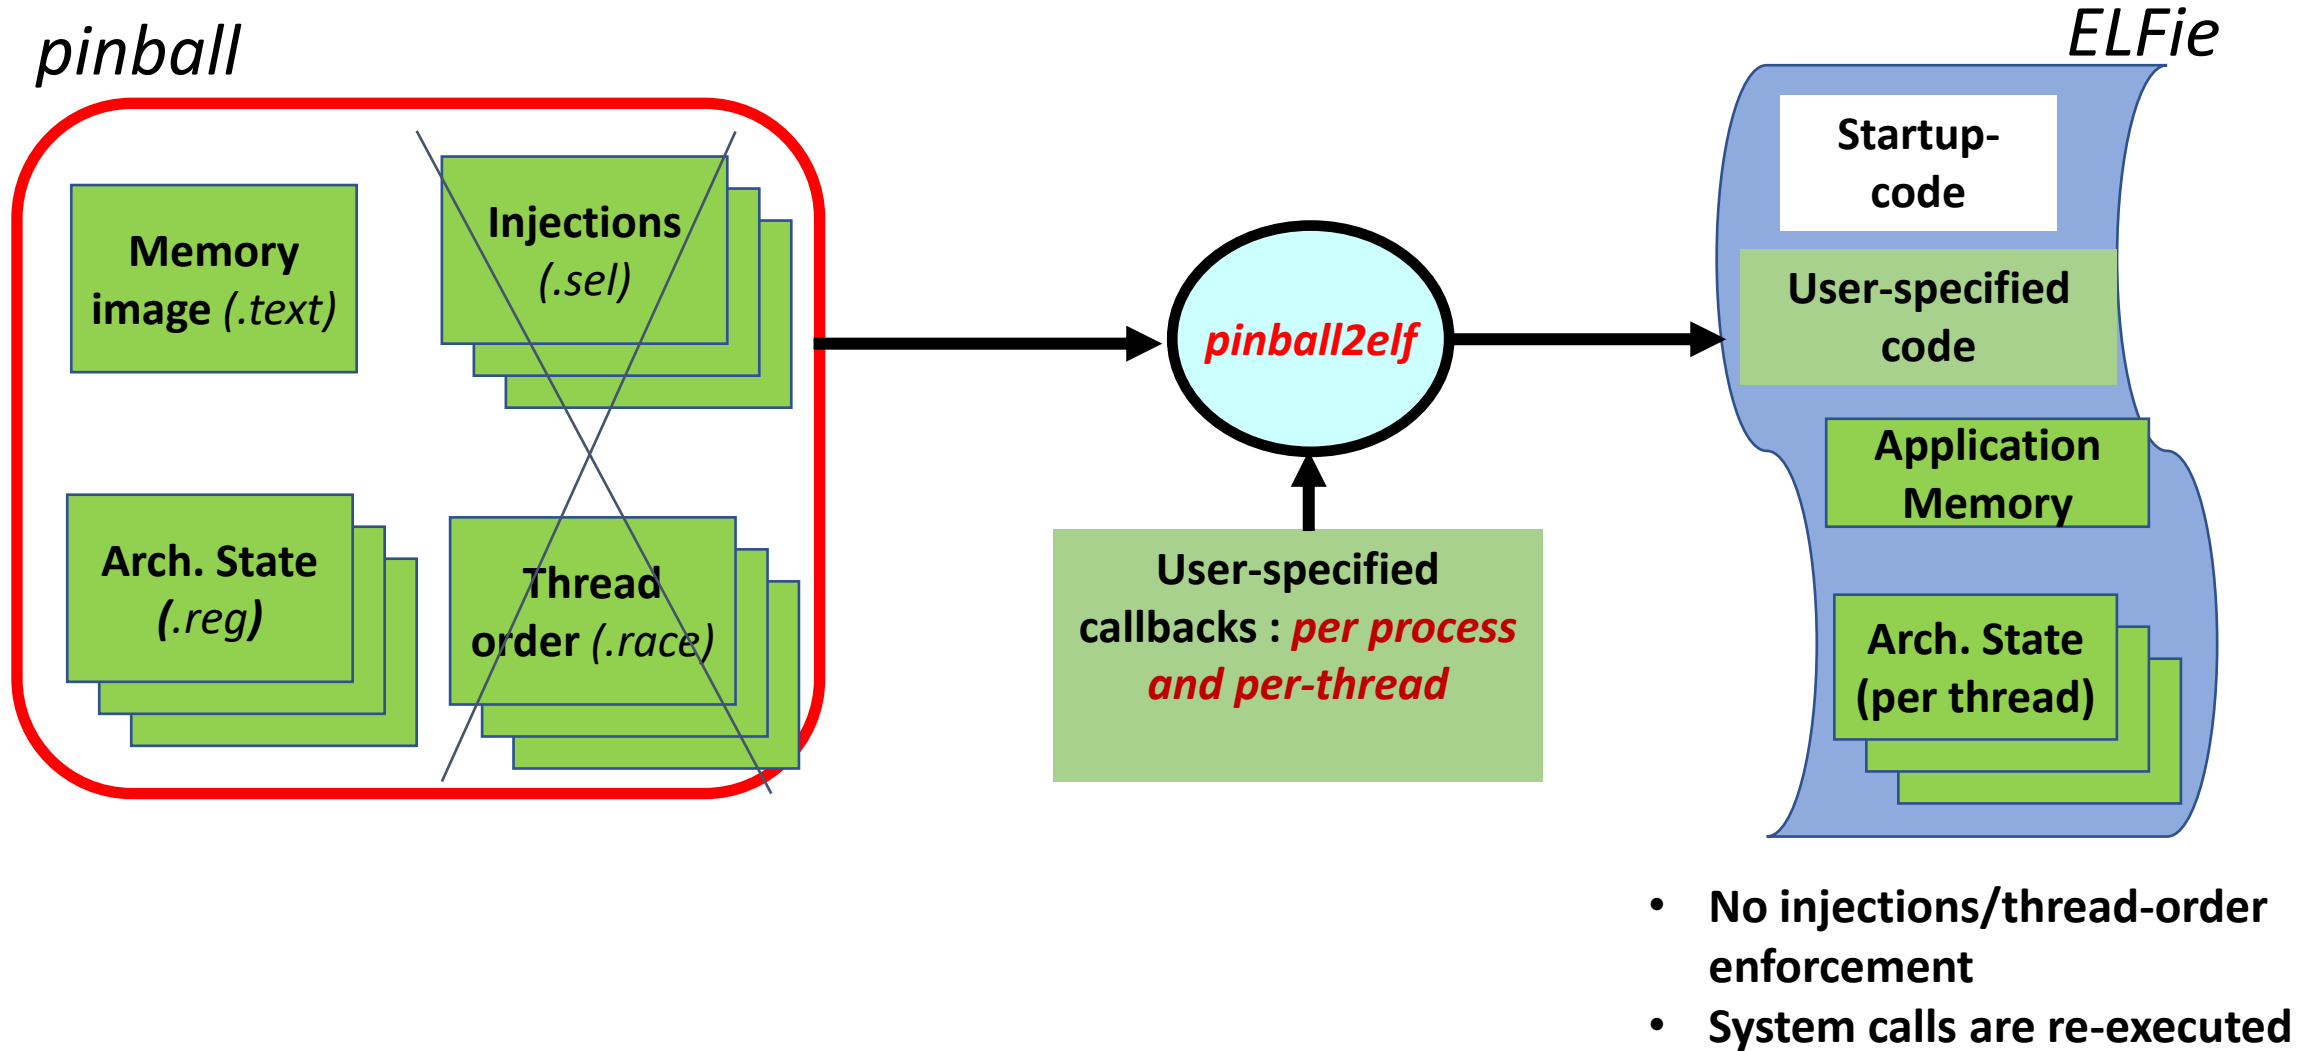
If you know your “Regions of Interest”, you may use SDE directly to generate pinballs for them

- SDE supports a “Controller” mechanism which has multiple ways of specifying ROI
- See [this article](#) or search for “Pintool regions”
- Example: (See run.sde.binary.log.sh in the PinPoints/Test directory)

```
% $SDE_ROOT/sde64 -controller_log -controller_olog region.binary -control  
start:address:dotproduct-st+0x12b0:count2692 -control stop:icount:300000 -log -log:basename  
pinball.region.binary/rpb -- dotproduct-st
```

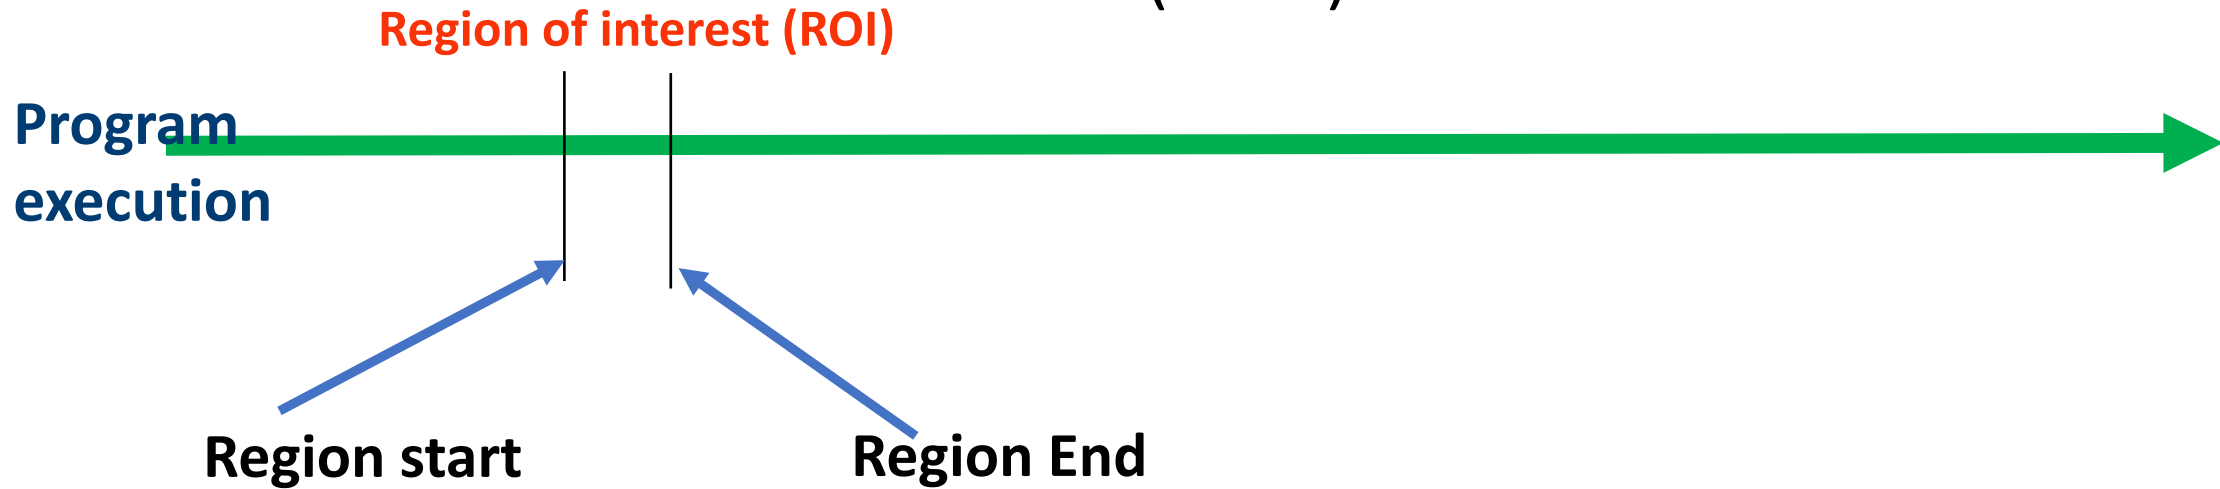
Will create a pinball region starting at the PC=“image+offset”=“dotproduct-st+0x12b0” for a length of 300,000 instructions (as counted by SDE)

pinball2elf: Pinball converter to ELF



Open-sourced on *GitHub* : <http://www.pinelfie.org>
(or search “pinball2elf”)

ROIperf: A Pin-probe tool for Targeted Performance monitoring of a region-of-interest (ROI)



Output: '*rdtsc*' and HW counter values at region boundaries

Hardware performance counter specification:

Comma separated pairs '*perftype:counter*'

ROIperf_LIST="0:0,0:1..."

Counter: 0:0 hw_cpu_cycles, 0:1: hw_instructions

based on /usr/include/linux/perf_event.h

perftype: 0 --> HW 1 --> SW

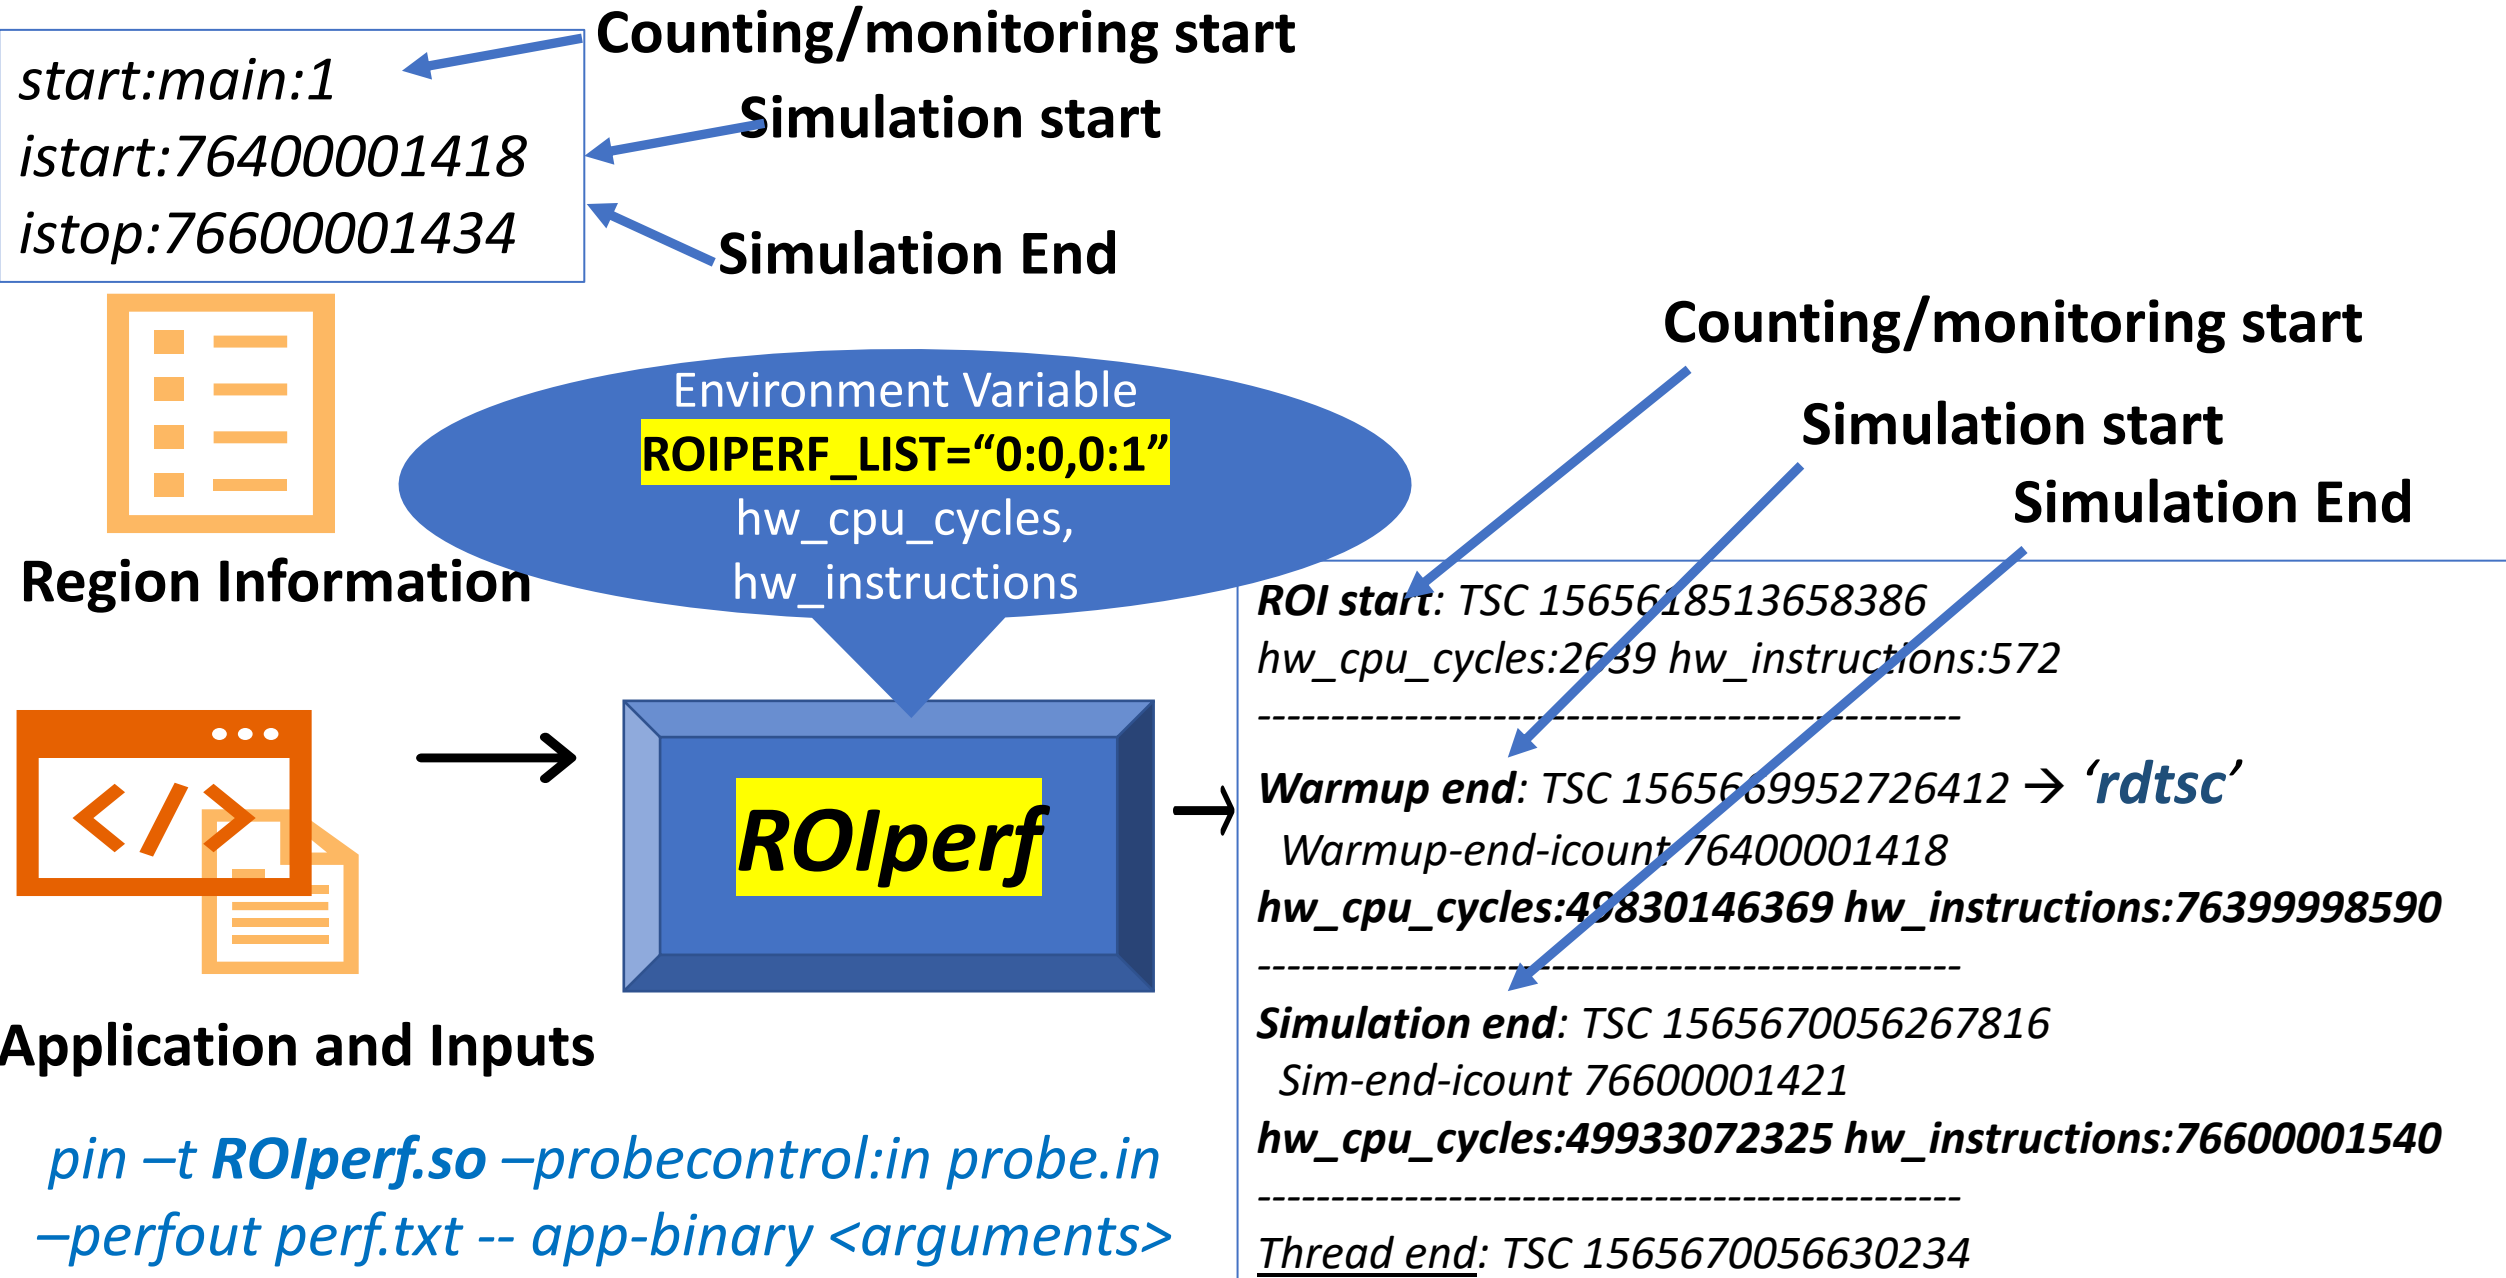
HW counter: 0 --> PERF_COUNT_HW_CPU_CYCLES

HW counter: 1 --> PERF_COUNT_HW_CPU_INSTRUCTIONS

SW counter: 0 --> PERF_COUNT_SW_CPU_CLOCK

... <see 'enum perf_hw_id' and 'enum perf_sw_id'>

Usage example



PinPoints Methodology

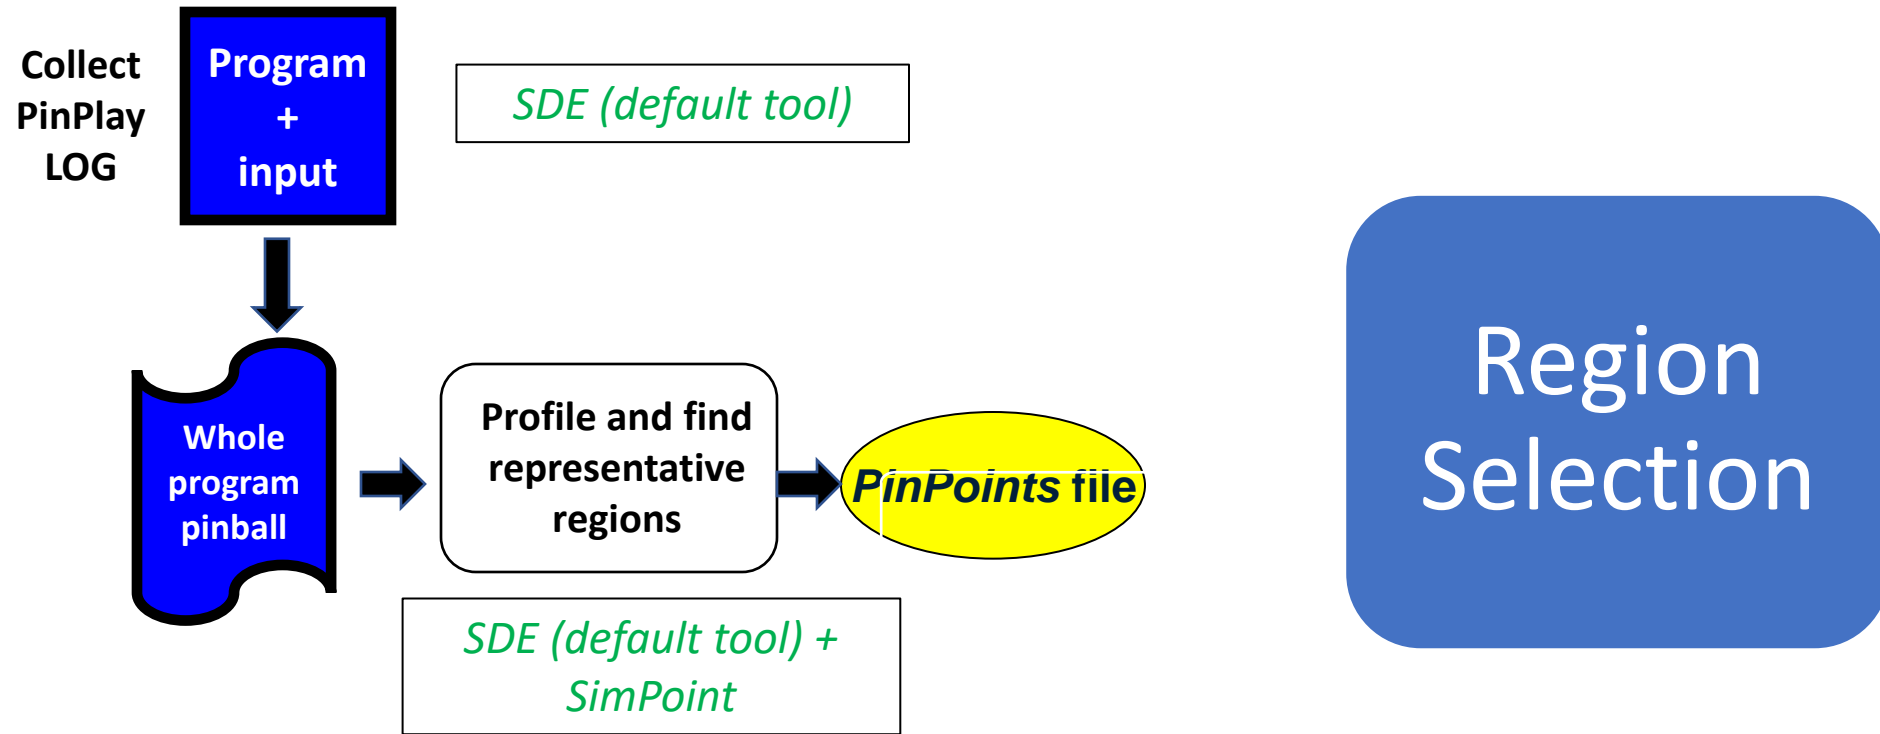
The diagram illustrates the PinPoints Methodology. It features a large, light blue arrow pointing to the right, which serves as a background for three blue rounded rectangular boxes. The boxes are arranged horizontally and contain the text 'Region Selection', 'Region Validation', and 'Simulator Staging' in white. The title 'PinPoints Methodology' is positioned at the top center of the image.

Region
Selection

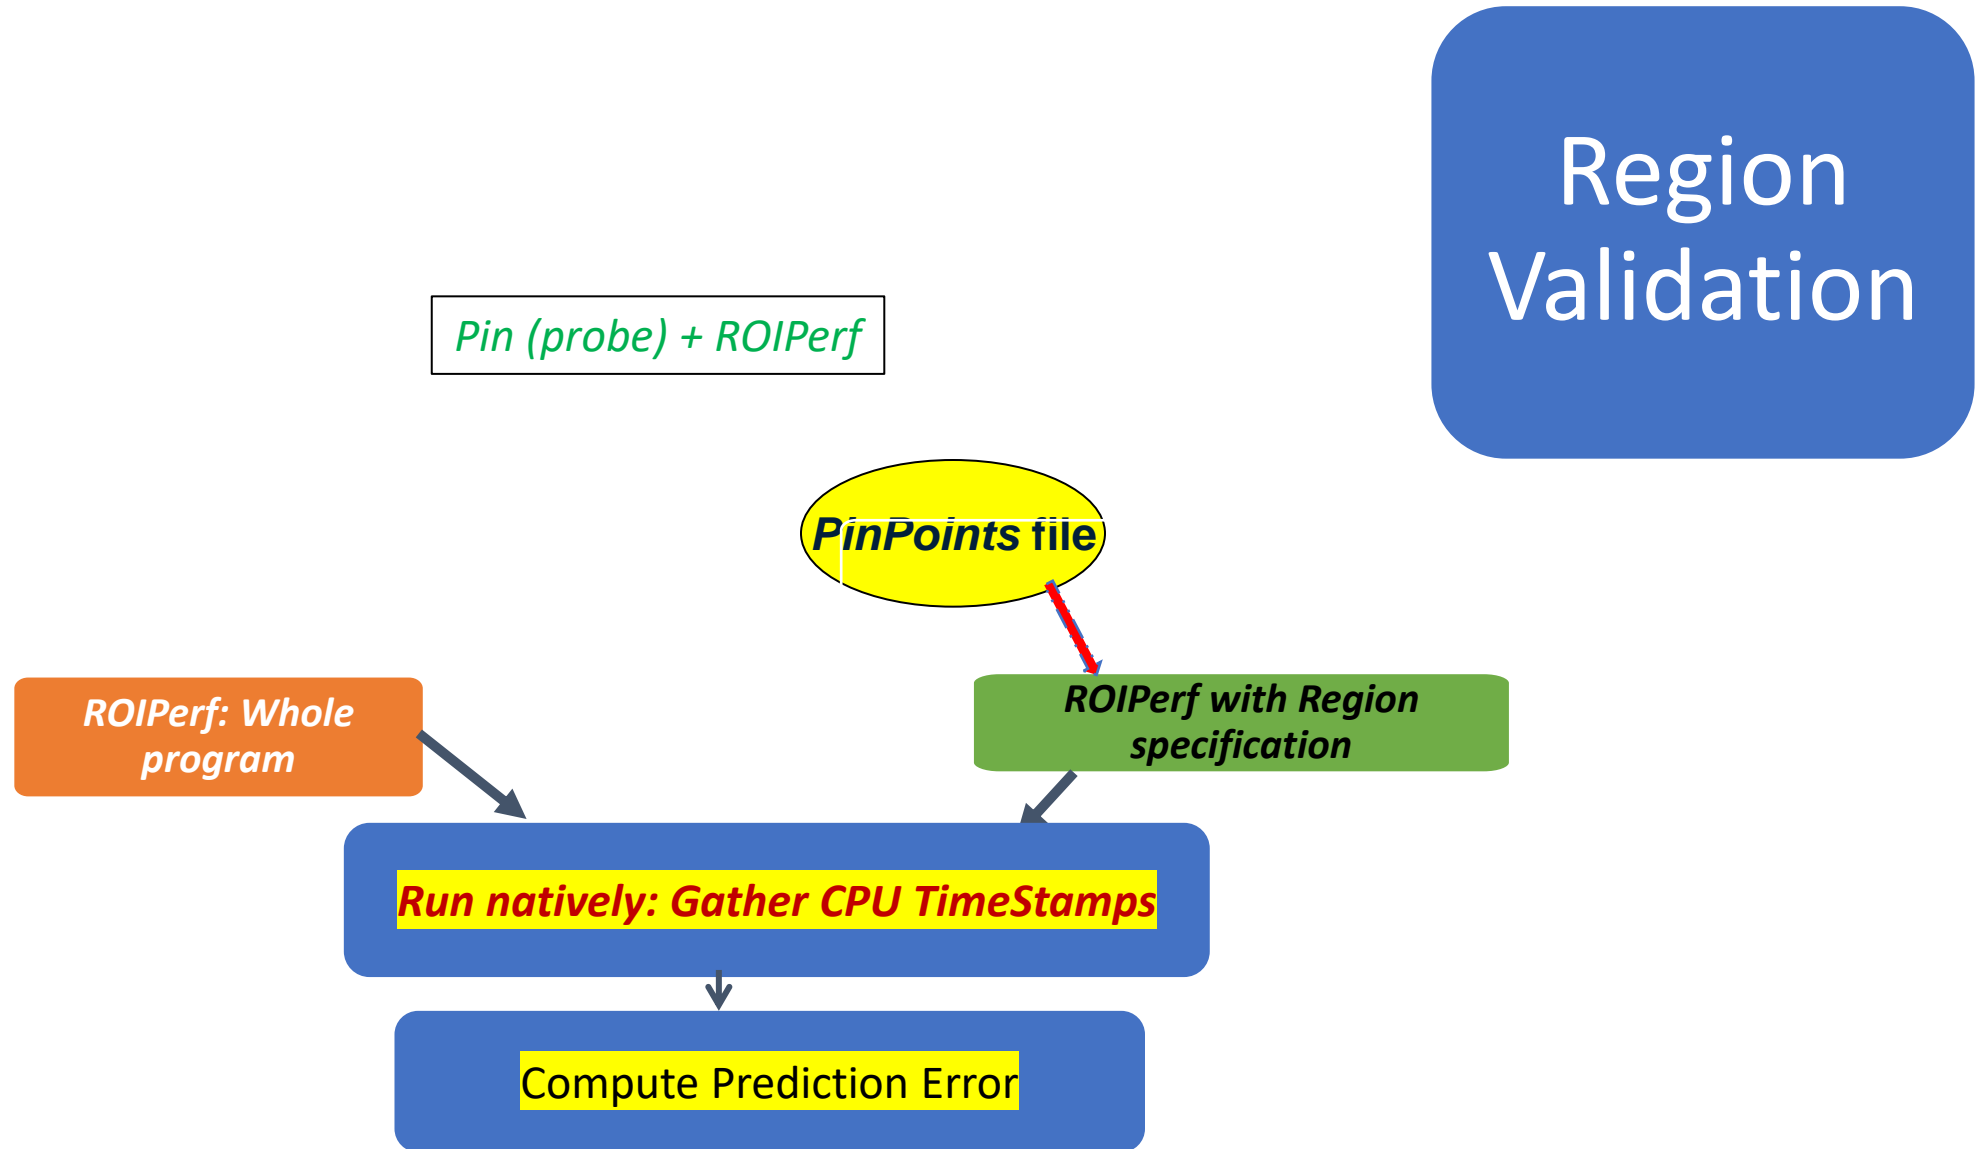
Region
Validation

Simulator
Staging

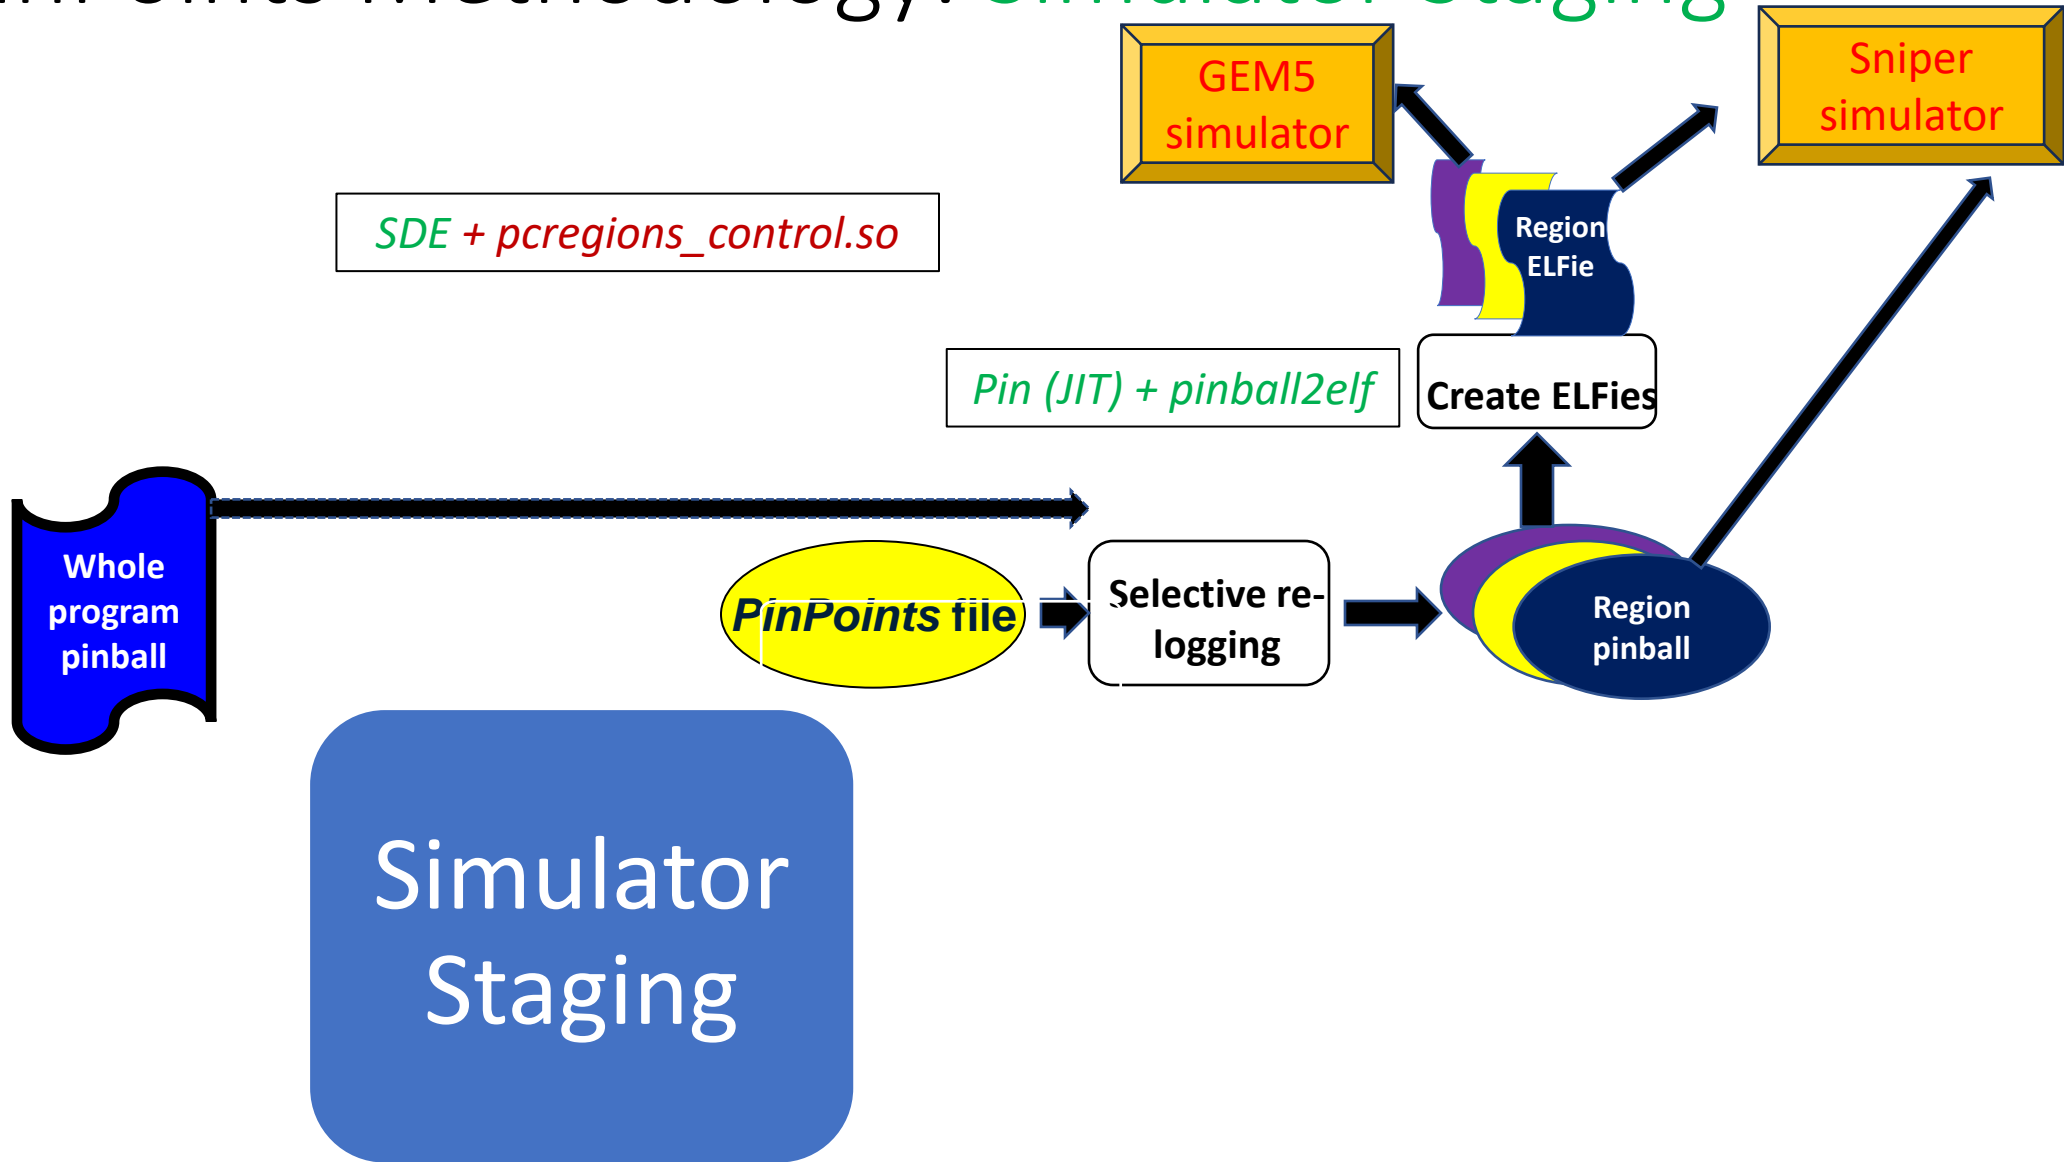
PinPoints Methodology: Simulation Region Selection



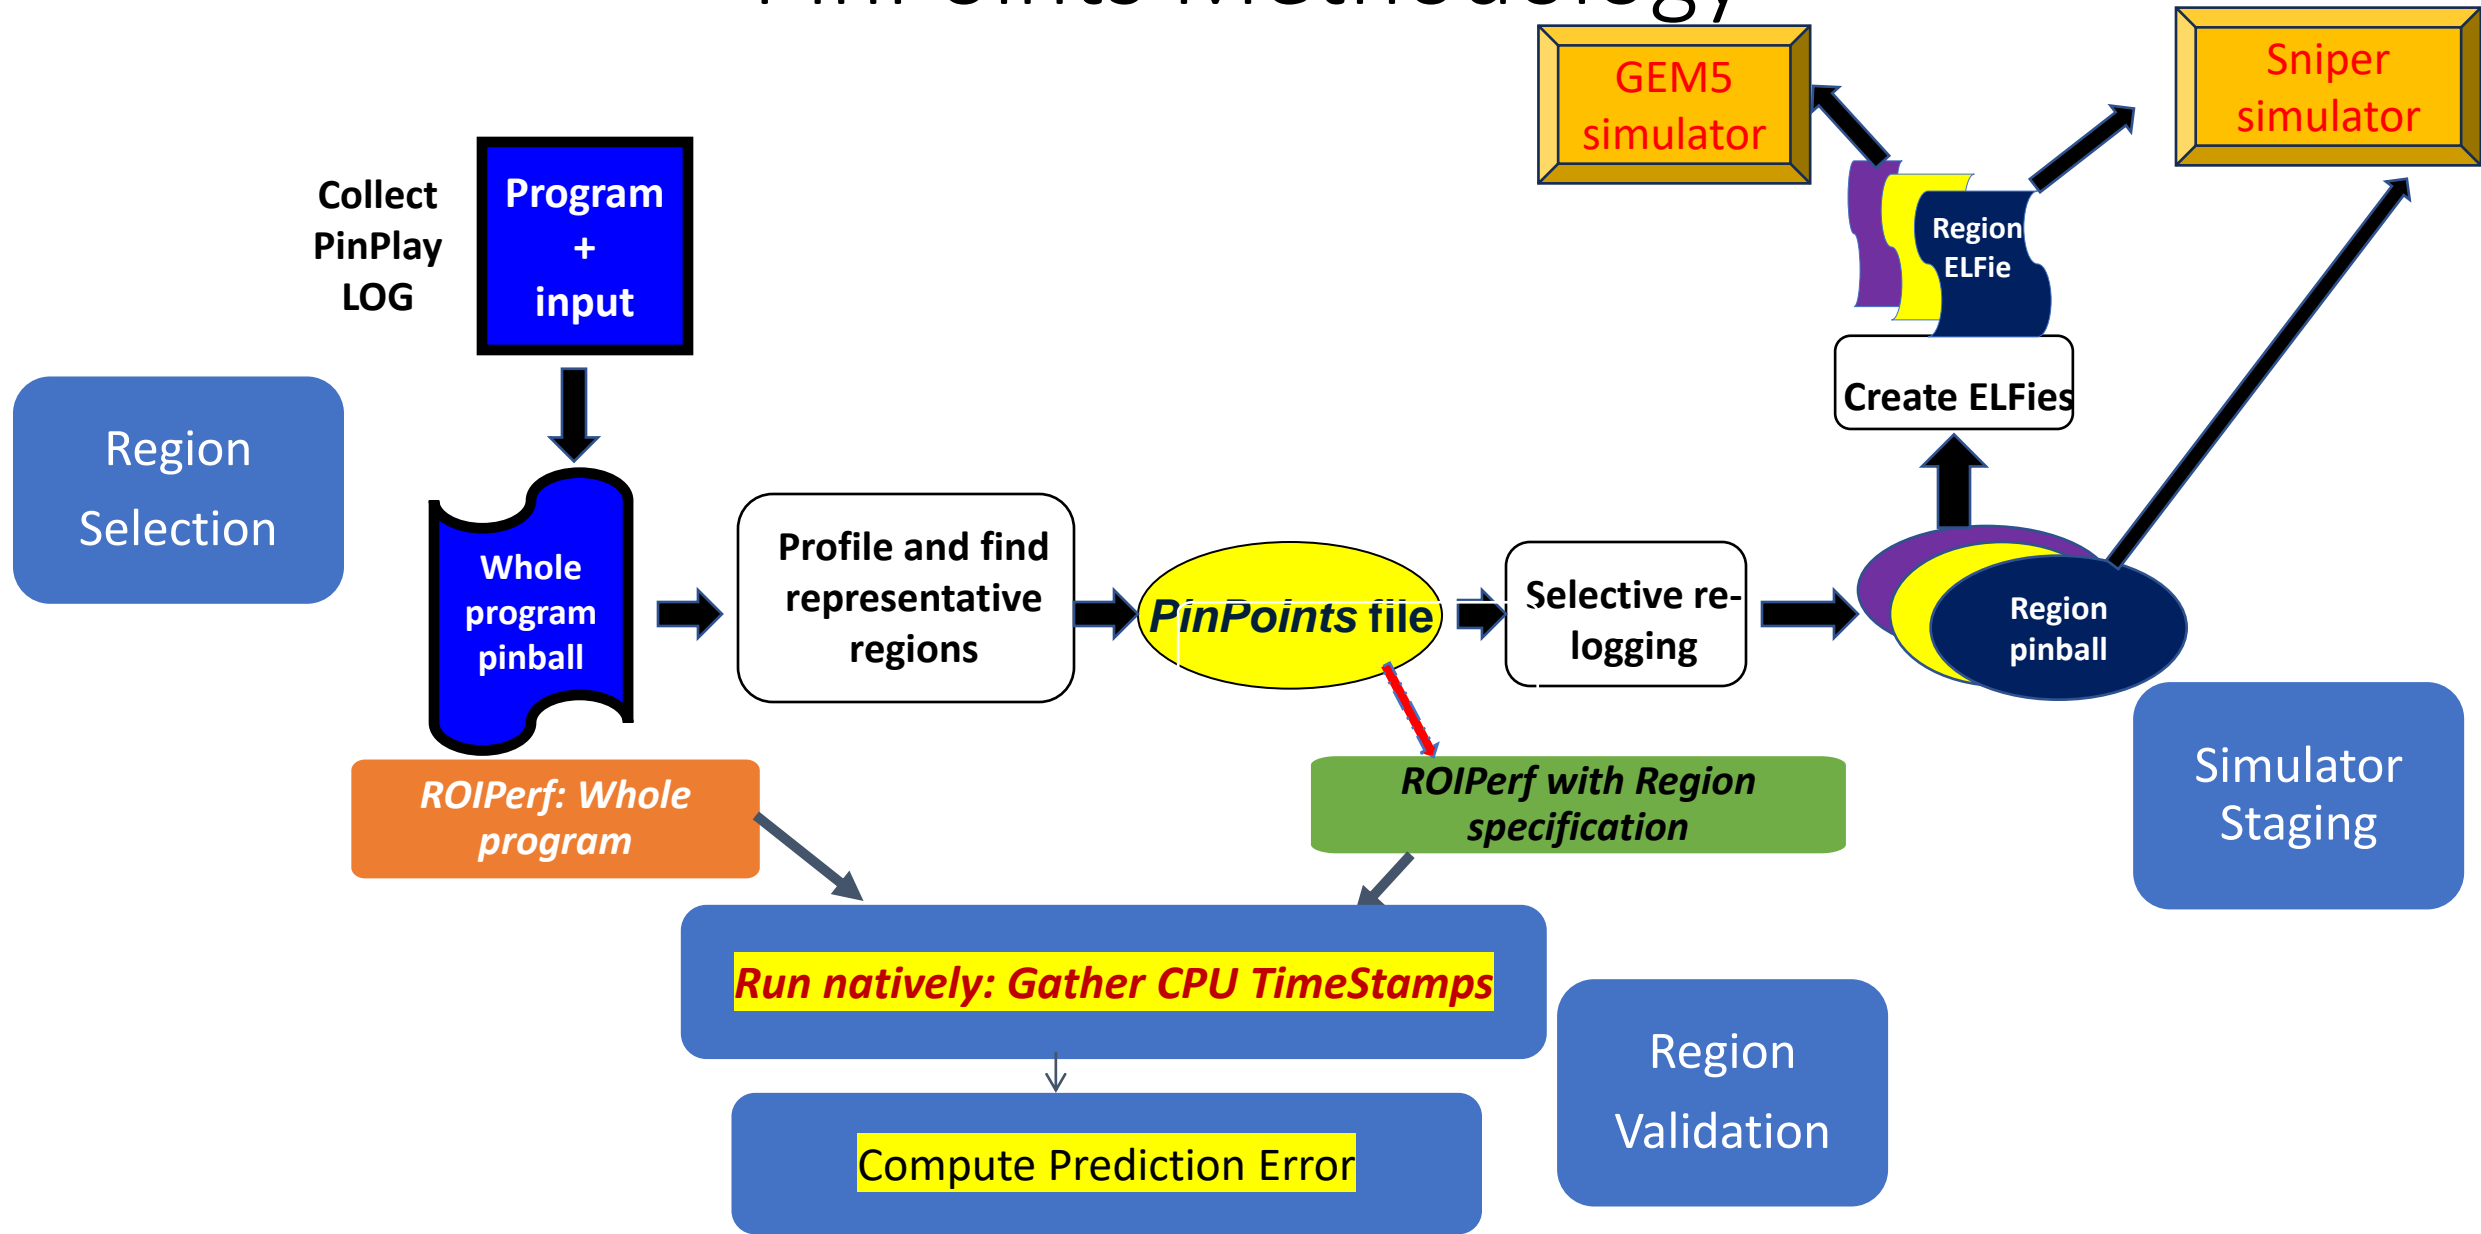
PinPoints Methodology: Validation



PinPoints Methodology: Simulator Staging



PinPoints Methodology



Demo:Part 1: Region Selection + region pinball generation

Pre-requisites

- Pin kit 3.31 : <http://pintool.intel.com> (or [this link](#))
 - export **PIN_ROOT**=<path to the local copy of Pin kit 3.31>
- Intel SDE 9.44 : Use [this link](#)
 - export **SDE_BUILD_KIT**=<path to the local copy of SDE Kit>
- Pinplay-tools repo: <https://github.com/intel/pinplay-tools>
 - export **PINPLAY_TOOLS**=<path to the local clone of the pinplay-tools repo>
- Pinball2elf repo: <https://github.com/intel/pinball2elf>
 - export **PINBALL2ELF**=<path to the local clone of the pinball2elf repo>

< Assumption: using "bash" >

< put "." in PATH >

Build all required tools

- Make sure the pre-requisites are satisfied and the following environment variables are set
 - SDE_BUILD_KIT
 - PIN_ROOT
 - PINBALL2ELF

```
% cd pinplay-tools/PinPoints
```

```
% sde-pin-build-PinPoints.sh
```

Build SDE profiler: 'pcregions_control.so'
(already built by 'sde-pin-build-PinPoints.sh')

```
% cd $SDE_BUILD_KIT/pinkit/sde-example/example
```

```
% make TARGET=intel64 clean; make TARGET=intel64
```

```
% cp obj-intel64/pcregions_control.so $SDE_BUILD_KIT/intel64
```

Test case: \$PINPLAY_TOOLS/PinPoints/Test


- cd \$PINPLAY_TOOLS/PinPoints/Test

- **< Edit** *sde-run.pinpoints.single-threaded.sh* **>**

% make clean; make

% sde-run.pinpoints.single-threaded.sh

```
SLICESIZE=75000
WARMUP_FACTOR=3
MAXK=20
PROGRAM=dotproduct-st
INPUT=1
COMMAND="./dotproduct-st"
```



WARMUP_FACTOR=3



**Include 3 slices before
simulation region for
warmup**

sde-run.pinpoints.single-thread.sh


```
#Whole Program Logging and replay using the default sde tool
# We are recording starting at 'main'
$SDE_BUILD_KIT/pinplay-scripts/sde_pinpoints.py --pin_options "$SDE_ARCH" $GLOBAL $PCCOUNT --program_name=$PROG
RAM --input_name=$INPUT --command="$COMMAND" --delete --mode st --log_options="-start_address main -log:fat -lo
g:mp_mode 0 -log:mp_atomic 0" --replay_options="-replay:strace" -l -r

#Profiling using regular profiler from the default sde tool
$SDE_BUILD_KIT/pinplay-scripts/sde_pinpoints.py --pin_options "$SDE_ARCH" $GLOBAL $PCCOUNT --program_name=$PROG
RAM --input_name=$INPUT --command="$COMMAND" --mode st -S $SLICESIZE -b

#Simpoint
$SDE_BUILD_KIT/pinplay-scripts/sde_pinpoints.py --pin_options "$SDE_ARCH" $GLOBAL $PCCOUNT --program_name=$PROG
GRAM --input_name=$INPUT --command="$COMMAND" $PCCOUNT -S $SLICESIZE $WARMUP --maxk=$MAXK --append_status -s

# Create per-region CSV files
#Create Makefile.regions with commands for relogging all regions
# and use 'make -j Makefile.regions' create all region pinballs in parallel

# Replay each region pinball serially
```



BYOP: Bring Your Own (whole-program) Pinball

- Generate pinball any other way, say at another site
- Say the pinball is input.pinball/log_0

```
% $SDE_BUILD_KIT/pinplay-scripts/sde_pinpoints.py ..  
--whole_pgm_dir input.pinball --pin_options  
"$SDE_ARCH" $GLOBAL $PCCOUNT --program_name=$PROGRAM  
--input_name=$INPUT --command="$COMMAND" --mode st -S  
$SLICESIZE -b
```

Demo:Part 2: Region Selection Validation

Build sde-global-event-icounter tool

(used for PCCount region processing)

(already built by 'sde-pin-build-PinPoints.sh')

```
% cd $PINPLAY_TOOLS/GlobalLoopPoint/EventCounter/
```

```
% make clean TARGET=intel64; make build TARGET=intel64
```

Build ROIPerf tool

(used for time-stamp[rdtsc] generation)
(already built by 'sde-pin-build-PinPoints.sh')

```
% cd $PINBALL2ELF/pintools/ROIProbe
```

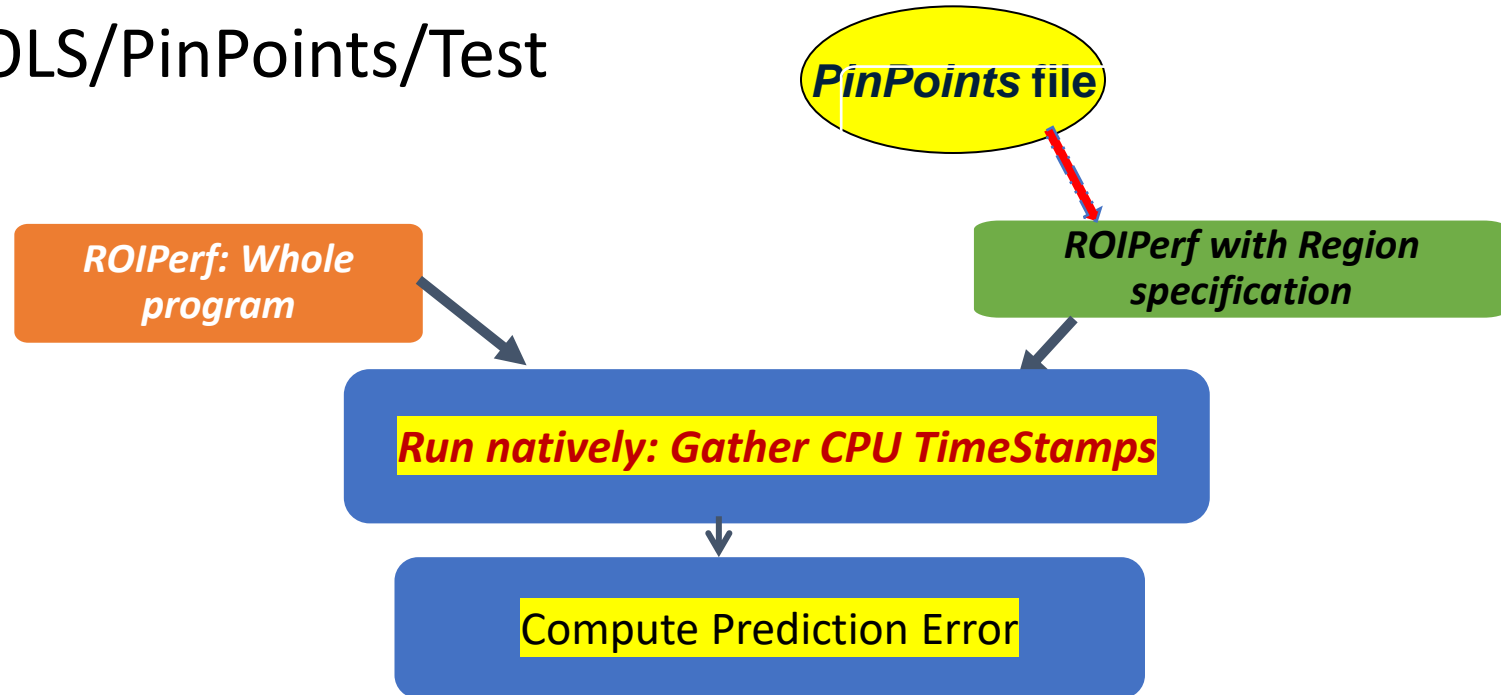
```
% make clean; make
```

```
% cp obj-intel64/pcregions_control.so $SDE_BUILD_KIT/intel64
```

Run ROIPerf-based validation

```
% cd $PINPLAY_TOOLS/PinPoints/Test
```

```
% run.ROIPerf.sh
```



```
wp_rdtsc,   region_rdtsc,   err%  
702607571.25,   504807628.48800004,   28.15%
```

- Prediction error can be tuned by changing SLICESIZE and MAXK
- Error high for short-running programs due to measurement overhead

Validation of PinPoints

SPEC2017 Rate/Train input: Simulation vs ROIperf

	% CPI Prediction Error		
	Intel-Sniper: Skylake Server	ROIperf: BroadWell Server	ROIperf: Skylake Server
perlbench_r.train.1	0.3%	0.5%	0.7%
perlbench_r.train.2	1.3%	0.0%	2.7%
perlbench_r.train.3	0.8%	0.0%	-0.1%
perlbench_r.train.4	0.2%	3.6%	3.0%
perlbench_r.train.5	8.9%	-6.6%	-0.7%
gcc_r.train.1	10.4%	-3.7%	-6.1%
gcc_r.train.2	1.0%	3.7%	0.4%
gcc_r.train.3	0.8%	3.9%	4.3%
mcf_r.train.1	8.8%	2.8%	0.2%
omnetpp_r.train.1	1.2%	2.8%	1.4%
x264_r.train.1	3.3%	0.0%	0.0%
leela_r.train.1	2.1%	-0.2%	-0.4%
exchange2_r.train.1	5.4%	0.0%	0.9%
xz_r.train.1	4.0%	-8.2%	-7.3%
xz_r.train.2	3.5%	-3.7%	-3.3%

Sniper Simulation time: up to 5 weeks
ROIperf validation time (3 trials each): few hours

Demo:Part 3: Creating ELFies and script templates for Sniper and GEM-5

ELFie generation

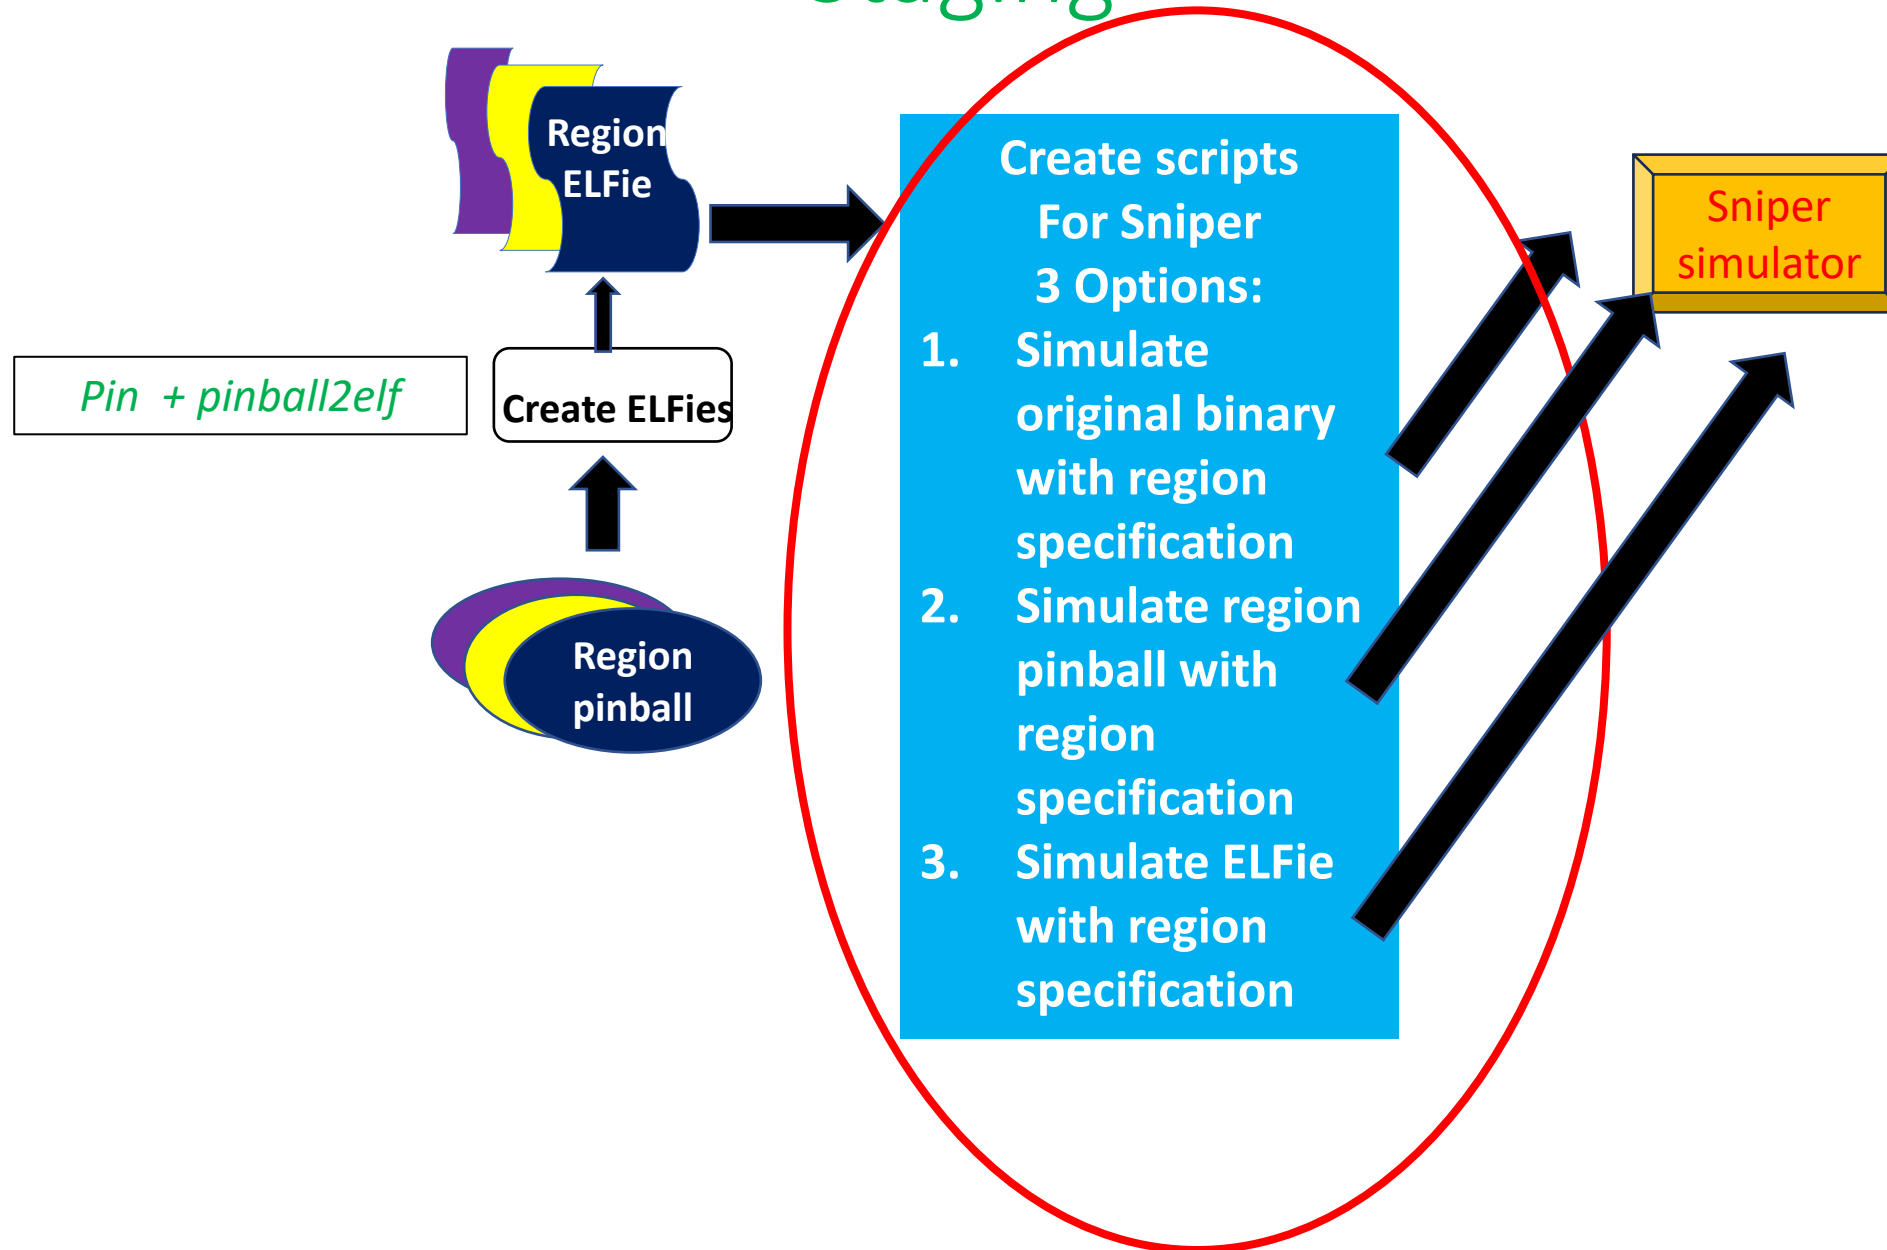
% run.pinball2elf.sh

% run.elfies.sh

<GEM5/Sniper template script generation to be added soon>

PinPoints Methodology: Sniper Simulator

Staging



Staging for Sniper

< Edit create.sniper_region_simulation_scripts.sh >

- %create.sniper_region_simulation_scripts.sh
creates one bash script per region

```
SNIPER_ROOT="ADD ME"  
SNIPER_ARGS="ADD ME"
```

```
#CHANGE  
export SNIPER_ROOT="ADD ME"  
#CHANGE  
export SNIPER_ARGS="ADD ME"
```

Change in the create.snipe*.sh script

```
#Uncomment for Option 1: Simulate the original application with current region specification  
#time $SNIPER_ROOT/run_sniper $SNIPER_ARGS --trace_args="-control start:address:dotproduct-st+0x14e0:  
count4528 -control stop:address:dotproduct-st+0x14e0:count10779" -- ./dotproduct-st
```

```
#Uncomment for Option 2: Simulate the region pinball with current region specification  
#time $SNIPER_ROOT/run_sniper $SNIPER_ARGS --trace_args="-control start:address:0x557e1a0414e0:count6  
52 -control stop:address:0x557e1a0414e0:count6251" --pinballs dotproduct-st.1_1362997.pp/dotproduct-s  
t.1_1362997_1_t0r1_warmupendPC0x557e1a0414e0_warmupendPCCount652_warmuplength225027_endPC0x557e1a0414  
e0_endPCCount6252_length75012_multiplier1951-044_001_0-45235
```

```
#Uncomment for Option 3: Simulate the region ELFie with current region specification  
#time $SNIPER_ROOT/run_sniper $SNIPER_ARGS --trace_args="-control start:address:0x557e1a0414e0:count6  
52 -control stop:address:0x557e1a0414e0:count6251" -- ./dotproduct-st.1_1362997.pp/dotproduct-st.1_13  
62997_1_t0r1_warmupendPC0x557e1a0414e0_warmupendPCCount652_warmuplength225027_endPC0x557e1a0414e0_end  
PCCount6252_length75012_multiplier1951-044_001_0-45235.sim.elfie
```


PinPoints Methodology: GEM5 Simulator

Staging



Resources

- [This tutorial](#)
- Past PinPoints tutorials (commands/kits outdated)
 - [ISCA2014-PinPoints-Tutorial](#)
 - [HPCA2013-PinPoints-Tutorial](#)
- Past PinPlay tutorials (commands/kits outdated)
 - [PLDI2015-PinPlay-Tutorial](#)
 - [PLDI2016-PinPlay-Tutorial](#)
- LoopPoint Methodology (multi-threaded programs): See <https://looppoint.github.io/>
- [Public Release and Validation of SPEC CPU2017 PinPoints](#) (Northwestern University)
- [SPEC2017 Pinballs from UT Austin](#)
- [SPEC highlights PinPoints research](#)