

PinPoints: A methodology for *Simulation Region Selection, Validation, and Simulator Staging* with Intel Pin and SDE

Harish Patil

Principal Engineer, Development Tools Software, Intel Corporation

With input from Alen Sabu (National University of Singapore)

November 7th, 2024

<https://github.com/intel/pinplay-tools/tree/main/PinPoints>

Architecture Simulation : Key Questions

Where to simulate?

Regions of Interest (ROIs)
($\ll 0.1$ % of whole-program)
NOT whole-program

How to simulate?

Drive simulations for ROIs

Are the regions representative?

Region Selection Validation

PinPoints Methodology

Where to simulate?

SDE-profiler + SimPoint

Region
Selection

How to simulate?

1. Checkpoint (pinball) driven
2. Binary-driven
 - Pinballs → ELFies

Simulator
Staging

Are the regions representative?

**ROIPerf:
CPU time-stamps with a Pin-probe tools**

Region
Validation

What is a “Region”?



Region specification: *Icount* vs *Pccount*

- *ICount*: instruction count based

Region Start: *icount1*
Region End: *icount2* (or *length* after Start)

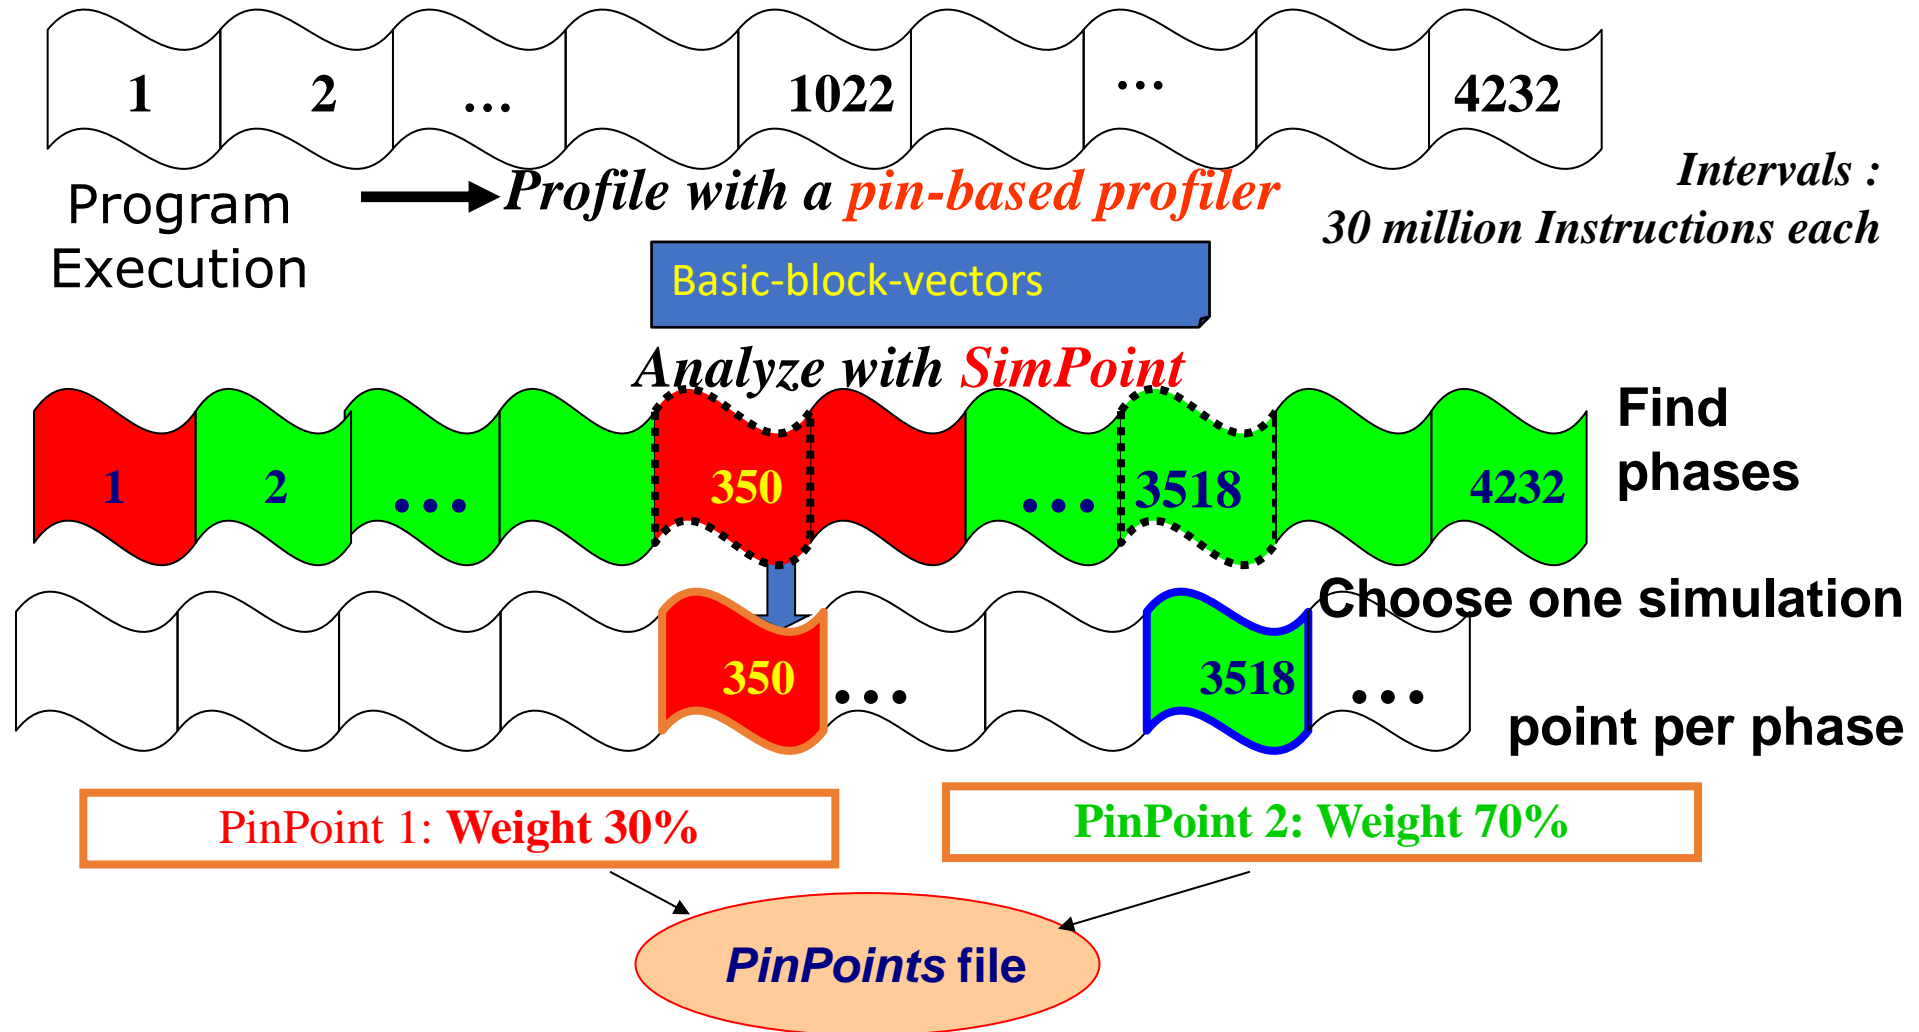
Drawbacks:

1. Different tools count instructions differently (e.g. REP-prefixed instructions as 1 or N)
2. (Multi-threaded programs) Instruction counts can change widely from run to run

- *PCcount*: program-counter (pc) based (select marker PCs with invariant counts across runs)

Region Start: *PC1+count1*
Region End: *PC2+count2* (relative to Start)

PinPoints = Pin + SimPoint (UC San Diego)



Two Phases => Two PinPoints

Projection Formulas

1. With weights (**ICount regions**) : ST only

$$\text{CPI}_{\text{predicted}} = \sum \text{Weight}_{\text{region}} * \text{CPI}_{\text{region}}$$

$$\text{Cycles}_{\text{predicted}} = \text{CPI}_{\text{predicted}} \times \text{Total_Instruction_Count (pathlength)}$$

2. With multipliers (**PCCount regions**) :ST and MT (no instruction counts used directly)

$$\text{Cycles}_{\text{predicted}} = \sum \text{Cycles}_{\text{region}} * \text{Multiplier}_{\text{region}}$$

where

$$\text{Multiplier}_j = \text{inscount}_j / \sum_{i=0..m} \text{inscount}_i$$

for m, the number of regions that are represented by the j^{th} representative region

Agenda

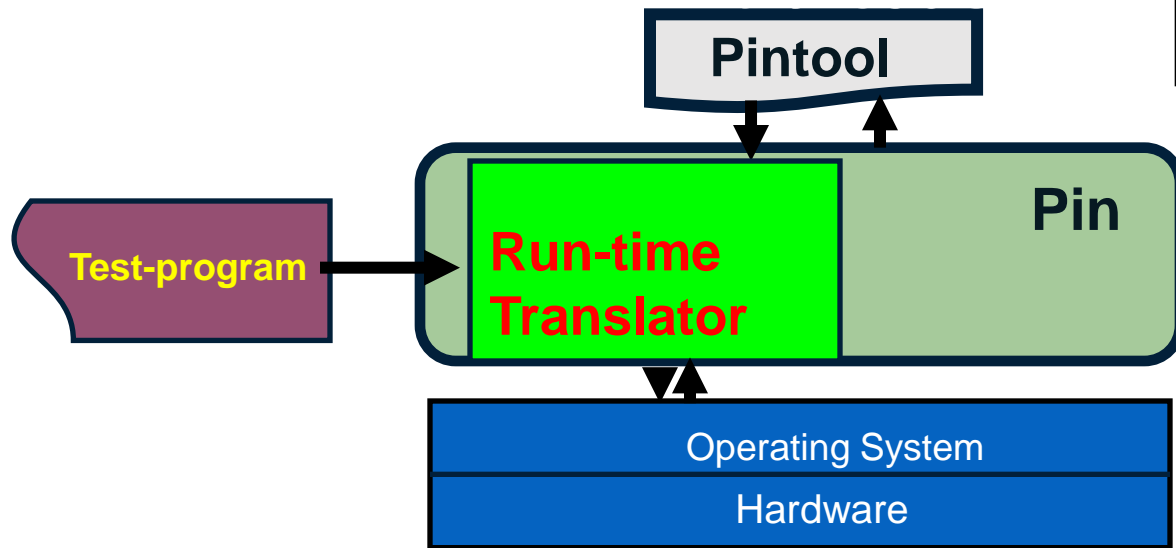
1. Tools Overview
2. Example/Demo
3. Discussion

Tools Overview

Pin: A Tool for Writing Program Analysis Tools

```
sub    $0xff, %edx  
movl   0x8(%ebp), %eax  
jle    <L1>
```

```
counter++; print(IP)  
sub    $0xff, %edx  
counter++; print(EA)  
movl   0x8(%ebp), %eax  
counter++; print(br taken)  
jle    <L1>
```



```
$ pin -t pintool -- test-program
```

Normal output
+ *Analysis*
output

Pin: A Dynamic Instrumentation Framework from Intel
<http://www.pintool.org> (or search "Intel Pintool"

JIT Mode vs Probe Mode

- JIT Mode

- Pin creates a modified copy of the application on-the-fly
- Original code never executes

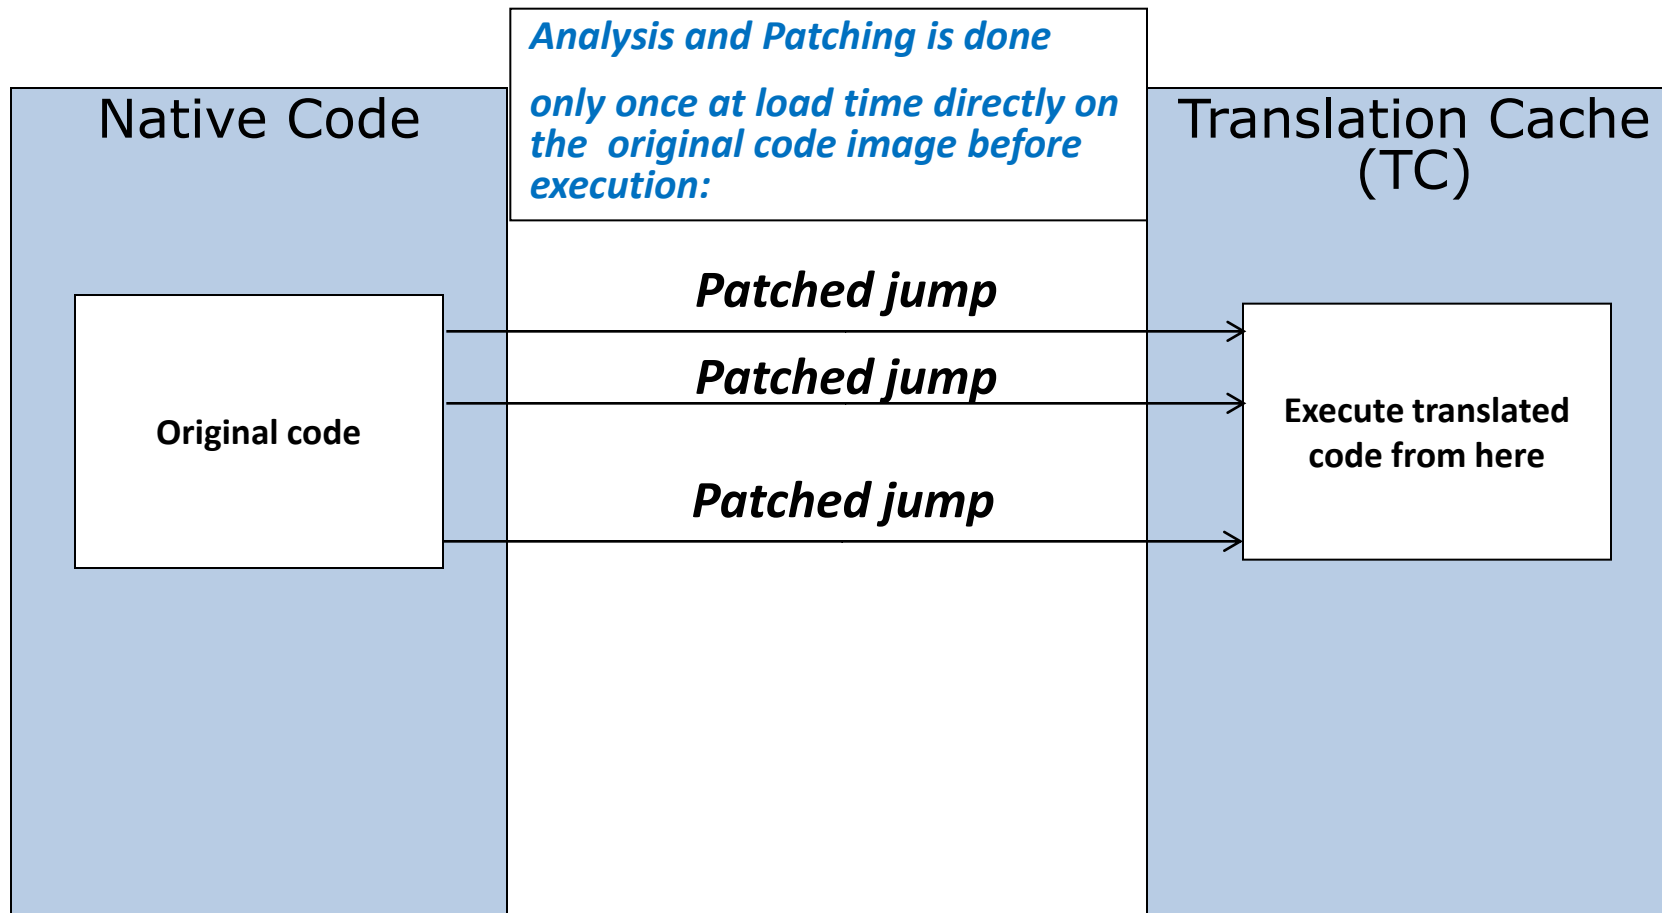
➤ More flexible, more common approach

- Probe Mode

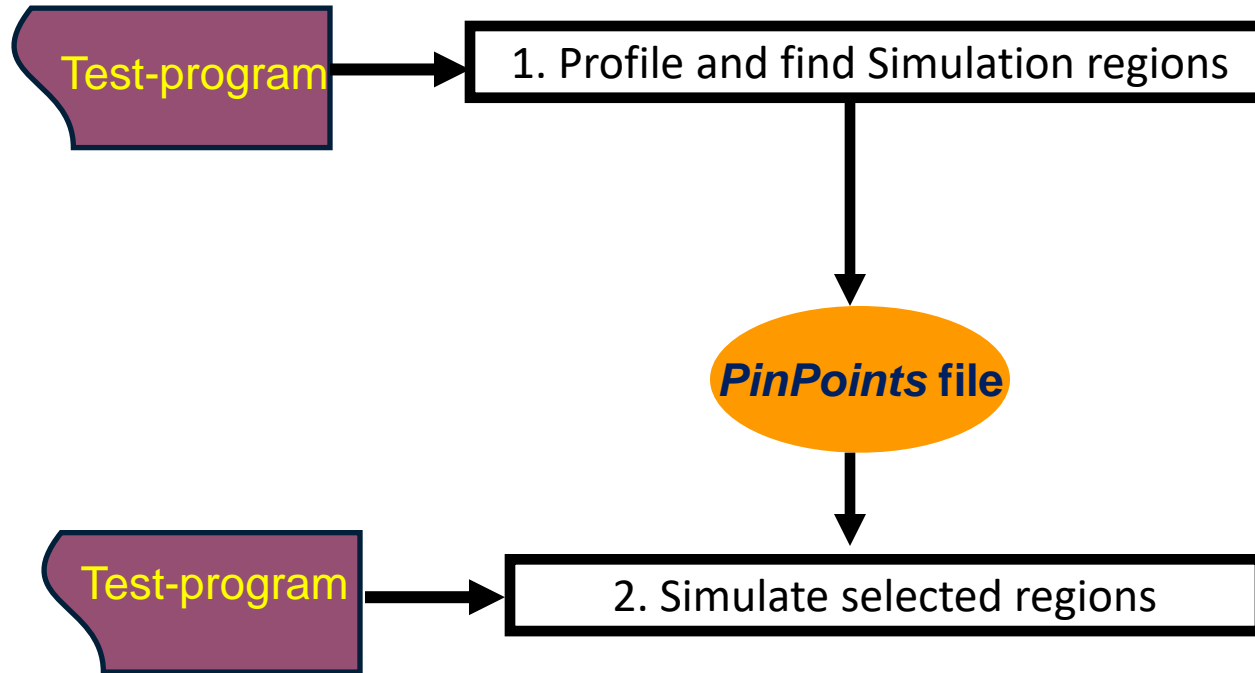
- Pin modifies the original application instructions
- Inserts jumps to instrumentation code (trampolines)

➤ Lower overhead (less flexible) approach

Pin Probe Mode – General Flow



PinPoints : The Repeatability Challenge

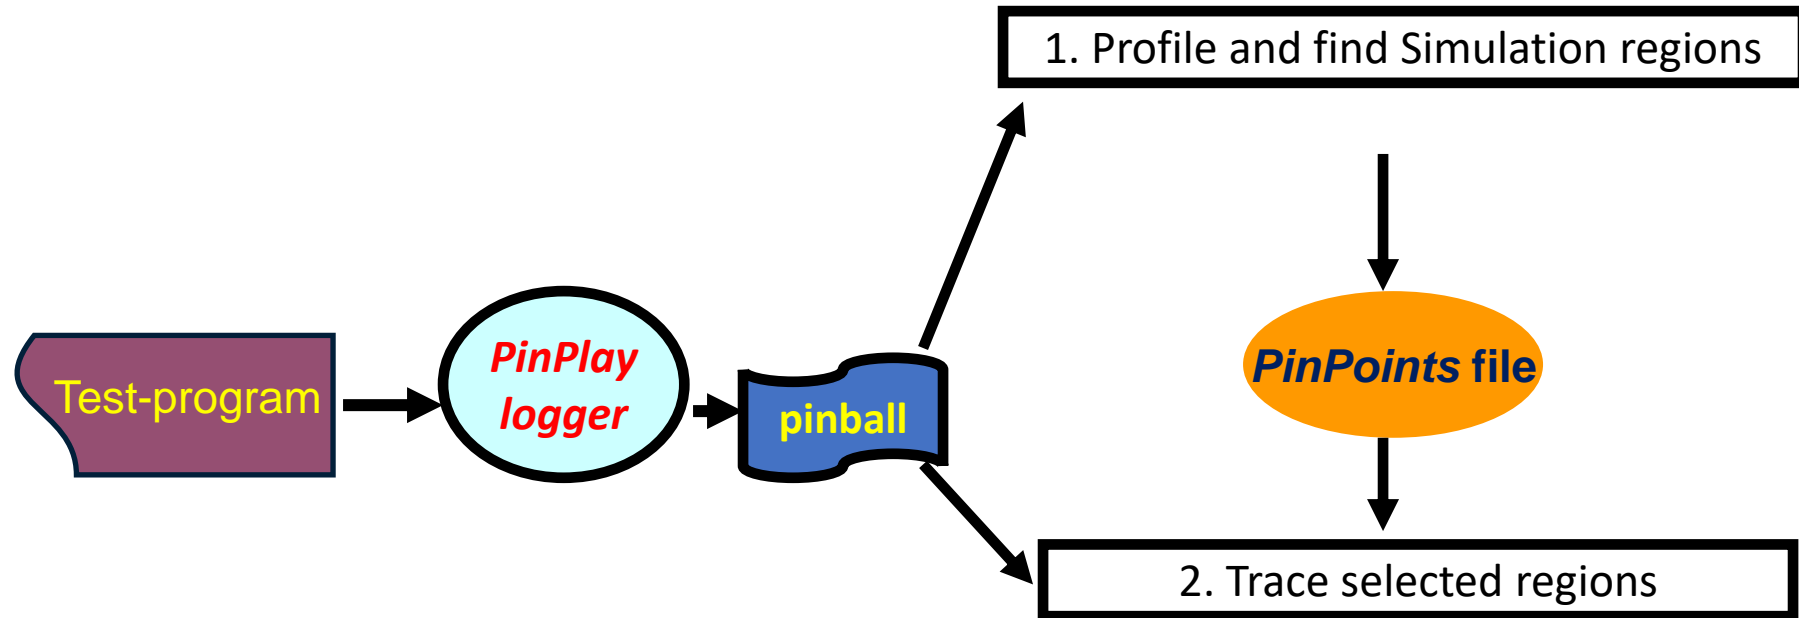


Problem: Two runs are not exactly same → PinPoints missed

Found this for 25/54 SPEC2006 runs!

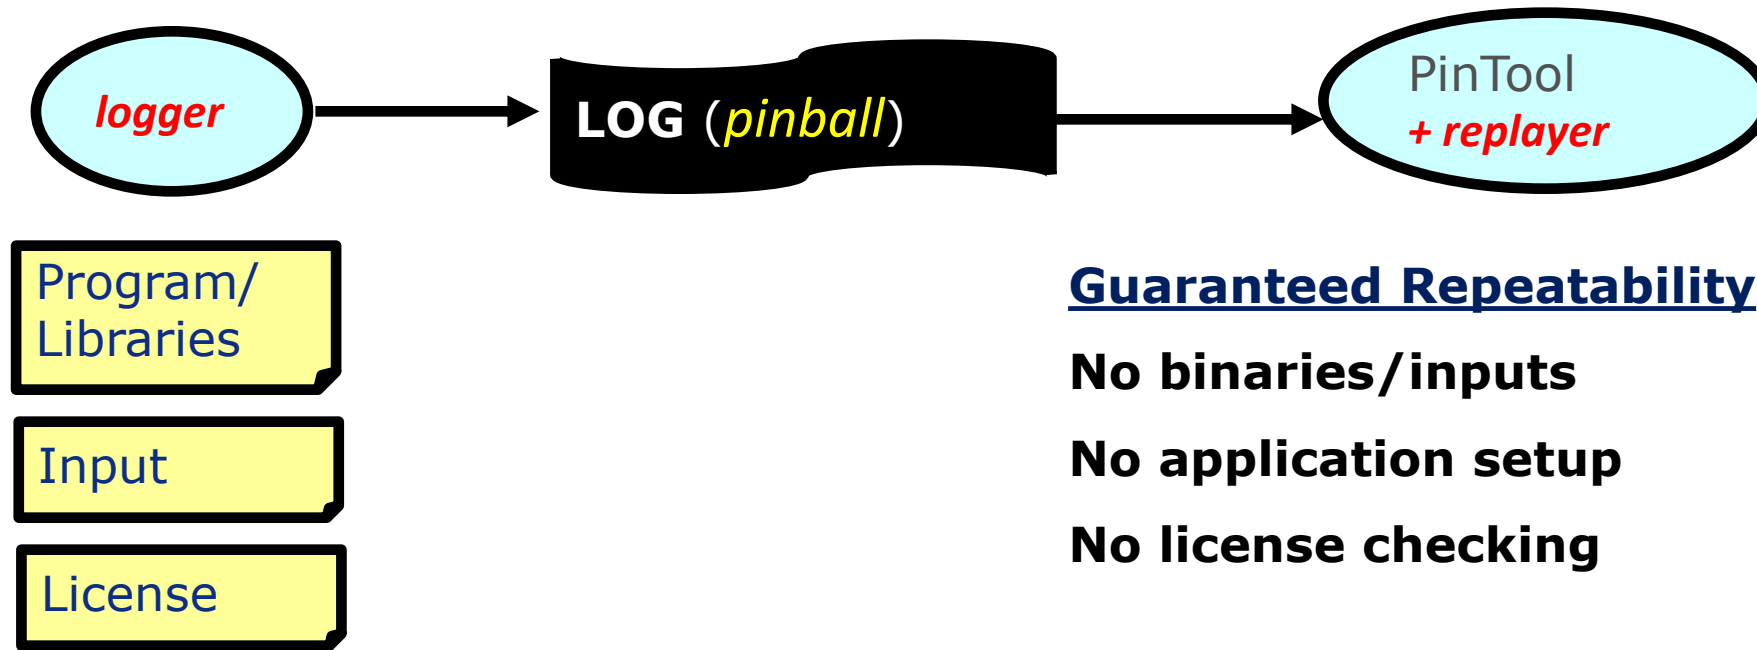
[*"PinPoints out of order" "PinPoint End seen before Start"*]

Enters PinPlay To Provide Repeatability



Two runs are same → PinPoints guaranteed to be reached

PinPlay*: Workload Capture and Deterministic Replay Framework



Record Once Replay + Analyze Multiple Times Anywhere!

PinPlay: Included in SDE

SDE : Pin + Emulation + PinPlay

- **SDE driver** runs with a default tool (sde-mix.so)
 - Emulation, log/replay, loop generation (DCFG json files), checkers.....
- **SDE Kit** also allows people to write tools
 - Regular Pin tools + use of additional SDE API
 - Can use emulation, log/replay, DCFG....

SDE : Intel Software Development Emulator
(search "Intel SDE")

Direct pinball generation with SDE

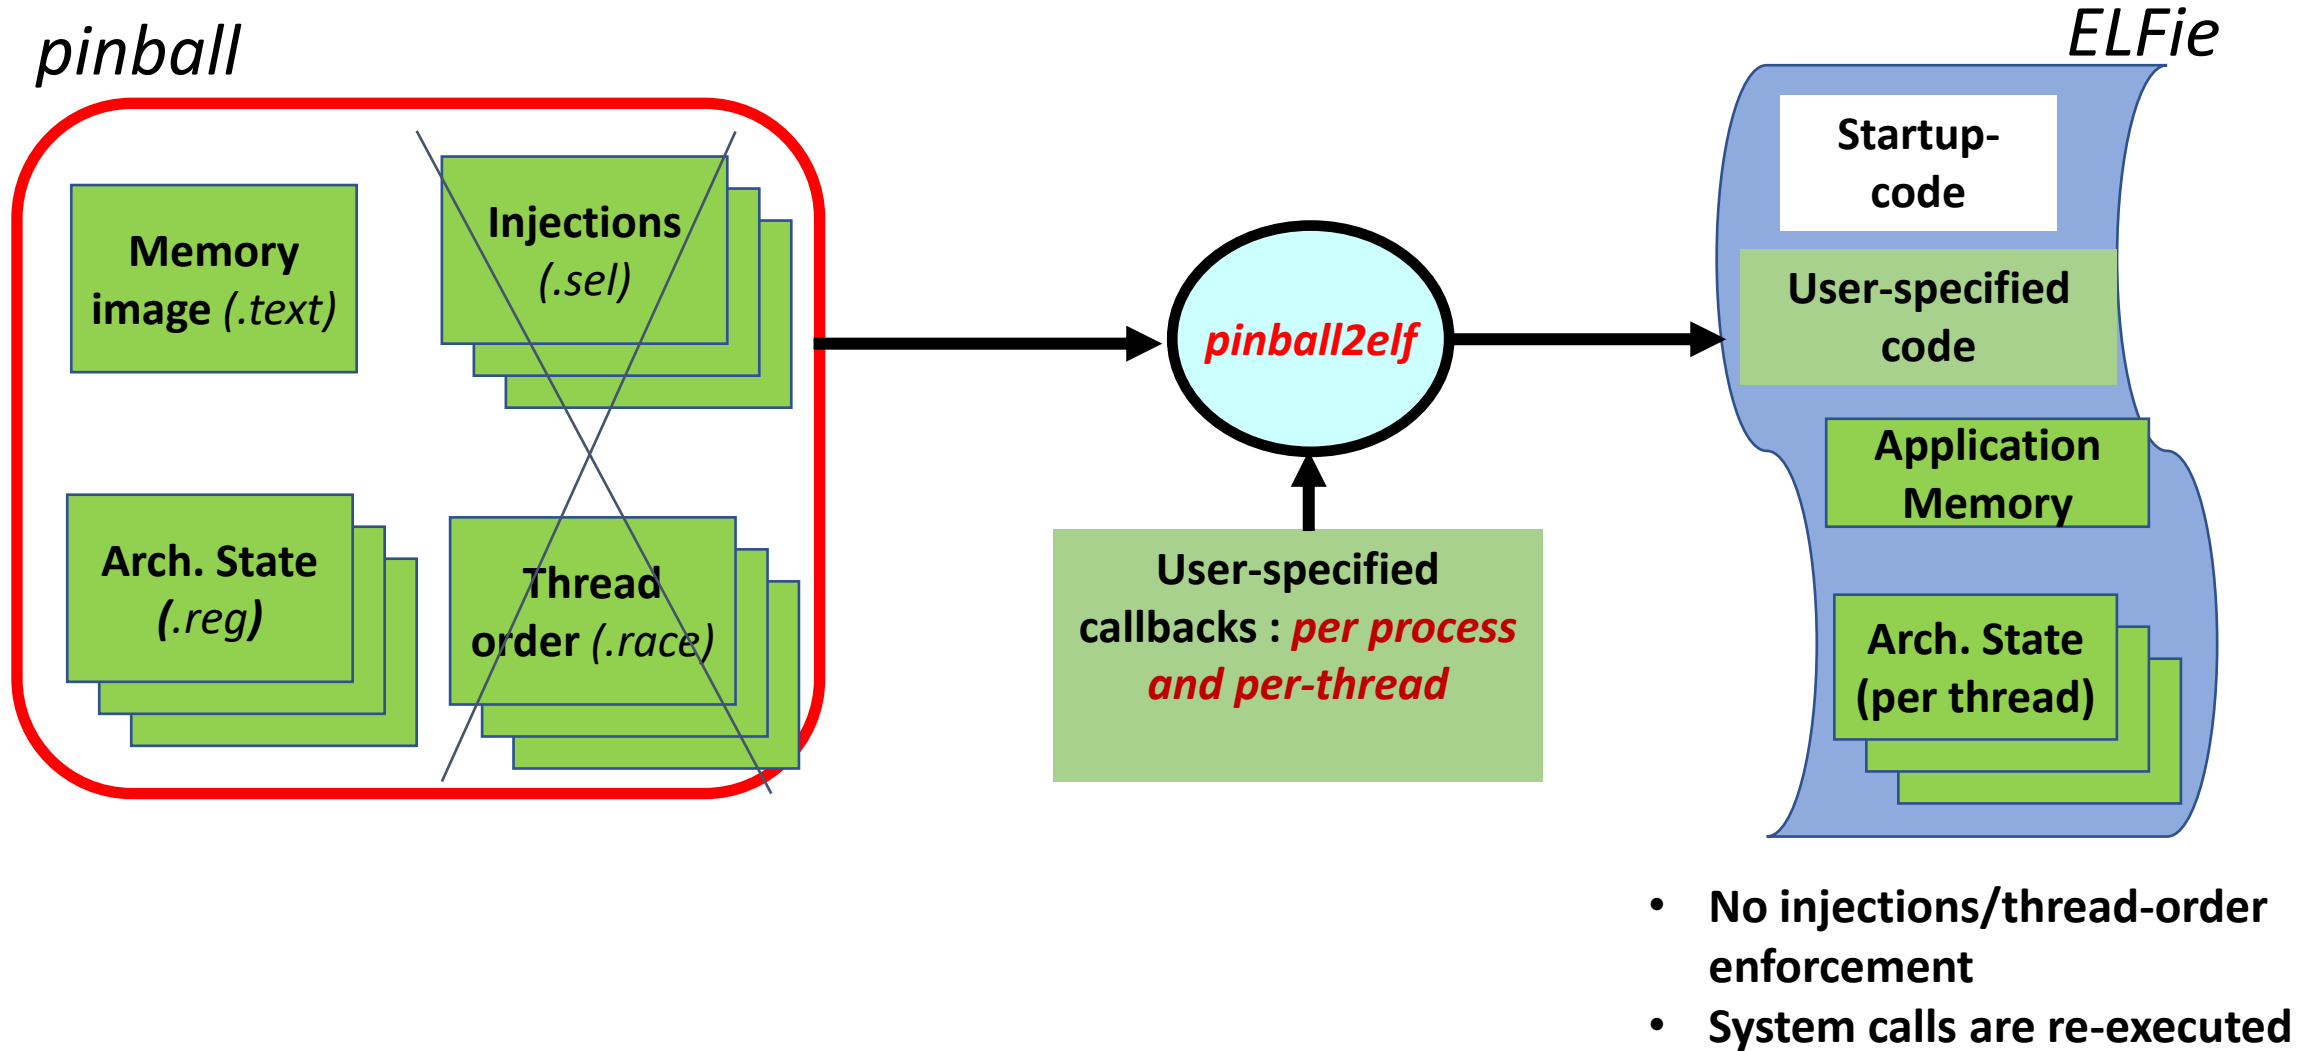
If you know your “Regions of Interest”, you may use SDE directly to generate pinballs for them

- SDE supports a “Controller” mechanism which has multiple ways of specifying ROI
- See [this article](#) or search for “Pintool regions”
- Example: (See run.sde.binary.log.sh in the PinPoints/Test directory)

```
% $SDE_ROOT/sde64 -controller_log -controller_olog region.binary -control  
start:address:dotproduct-st+0x12b0:count2692 -control stop:icount:300000 -log -log:basename  
pinball.region.binary/rpb -- dotproduct-st
```

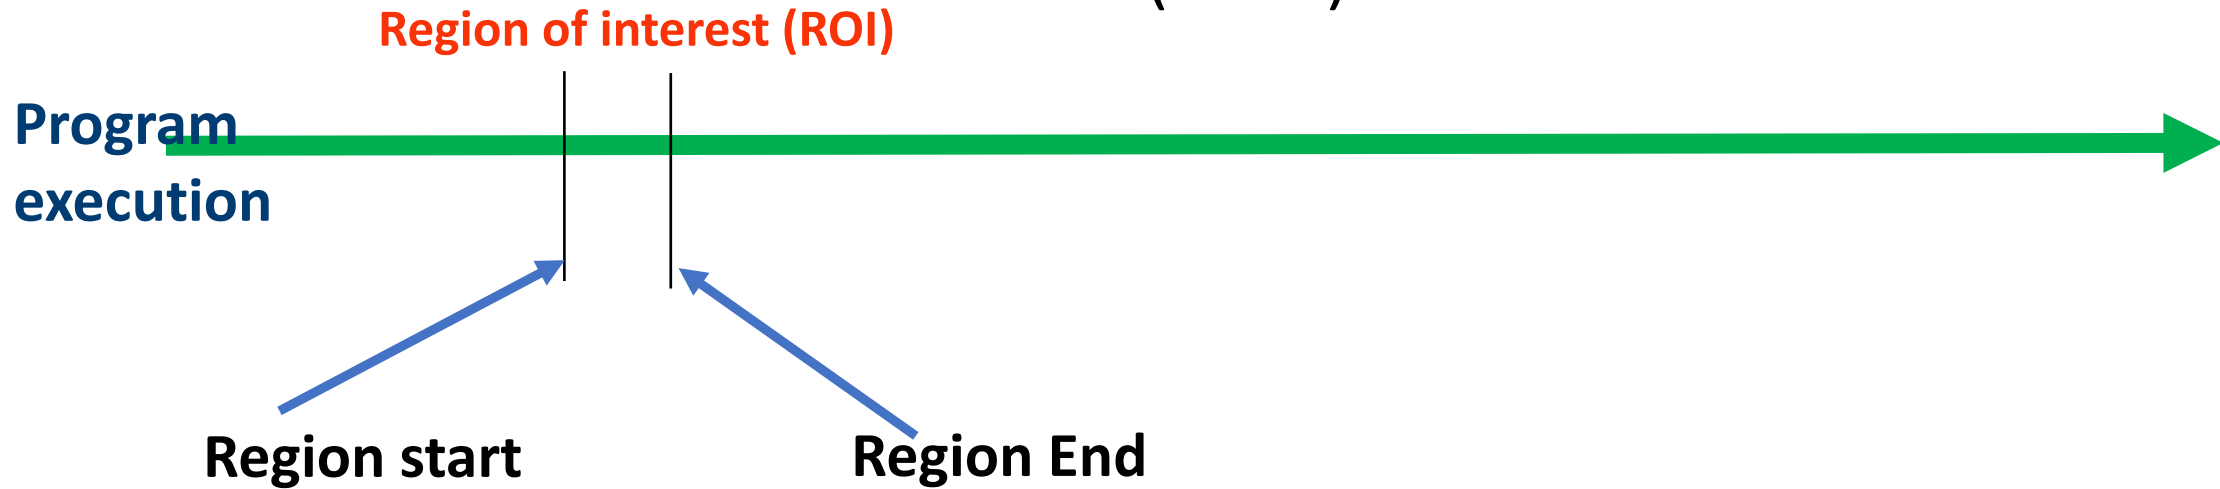
Will create a pinball region starting at the PC=“image+offset”=“dotproduct-st+0x12b0” for a length of 300,000 instructions (as counted by SDE)

pinball2elf: Pinball converter to ELF



Open-sourced on *GitHub* : <http://www.pinelfie.org>
(or search “pinball2elf”)

ROIperf: A Pin-probe tool for Targeted Performance monitoring of a region-of-interest (ROI)



Output: '*rdtsc*' and HW counter values at region boundaries

Hardware performance counter specification:

Comma separated pairs '*perftype:counter*'

ROIperf_LIST="0:0,0:1..."

Counter: 0:0 hw_cpu_cycles, 0:1: hw_instructions

based on /usr/include/linux/perf_event.h

perftype: 0 --> HW 1 --> SW

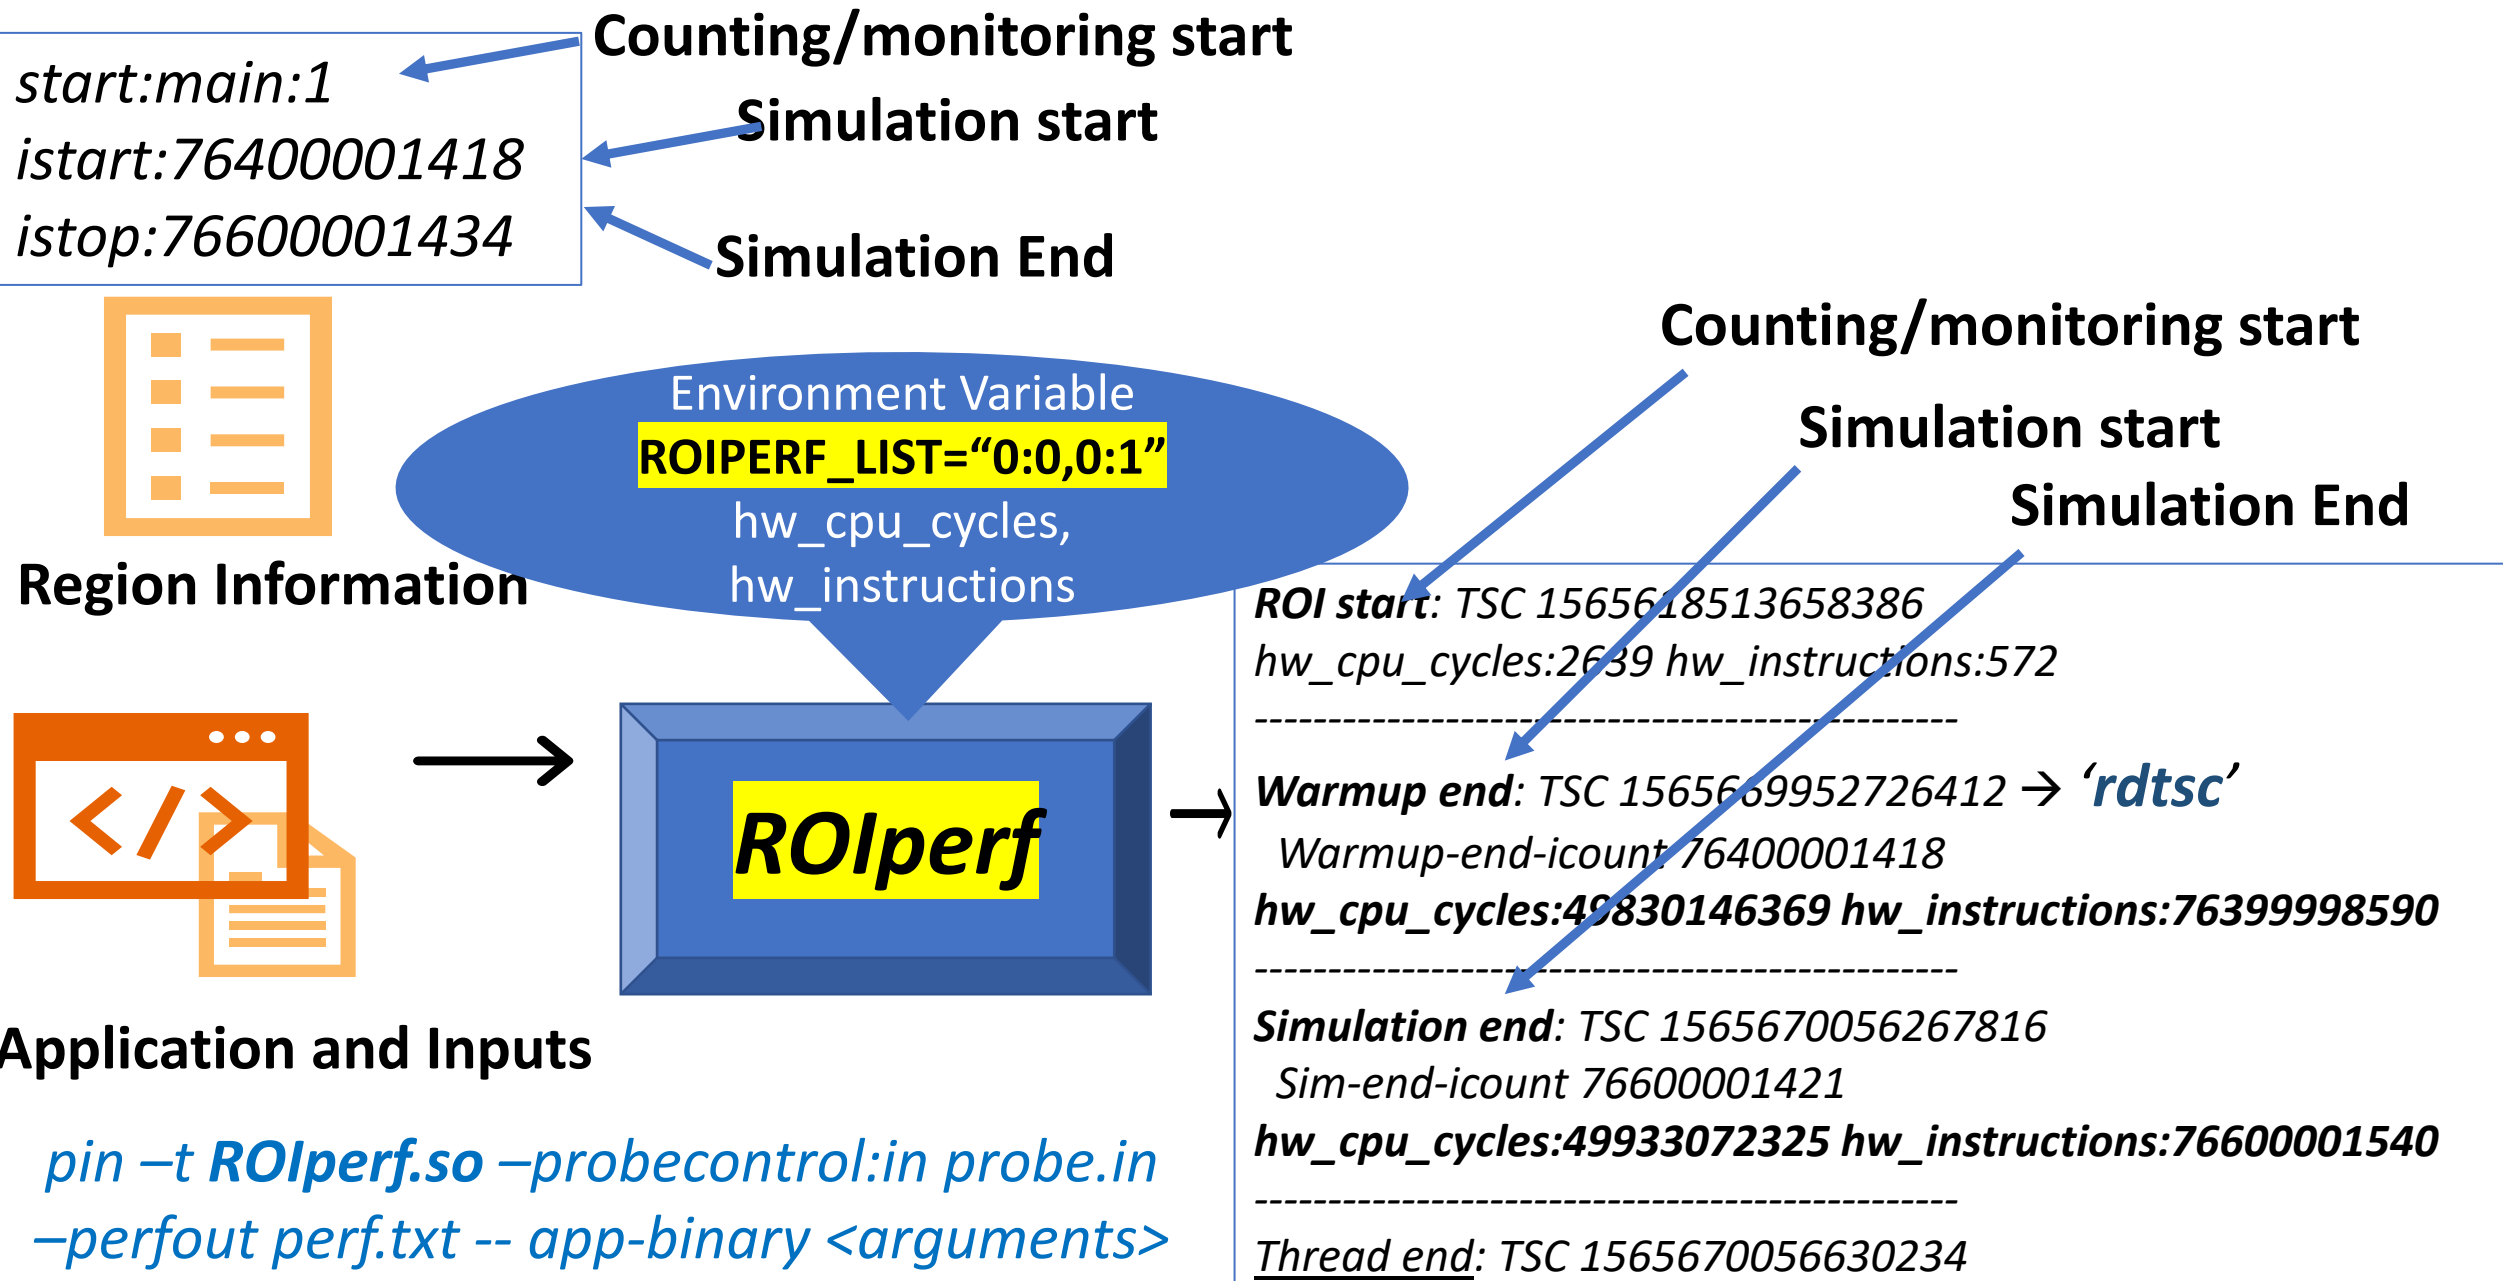
HW counter: 0 --> PERF_COUNT_HW_CPU_CYCLES

HW counter: 1 --> PERF_COUNT_HW_CPU_INSTRUCTIONS

SW counter: 0 --> PERF_COUNT_SW_CPU_CLOCK

... <see 'enum perf_hw_id' and 'enum perf_sw_id'>

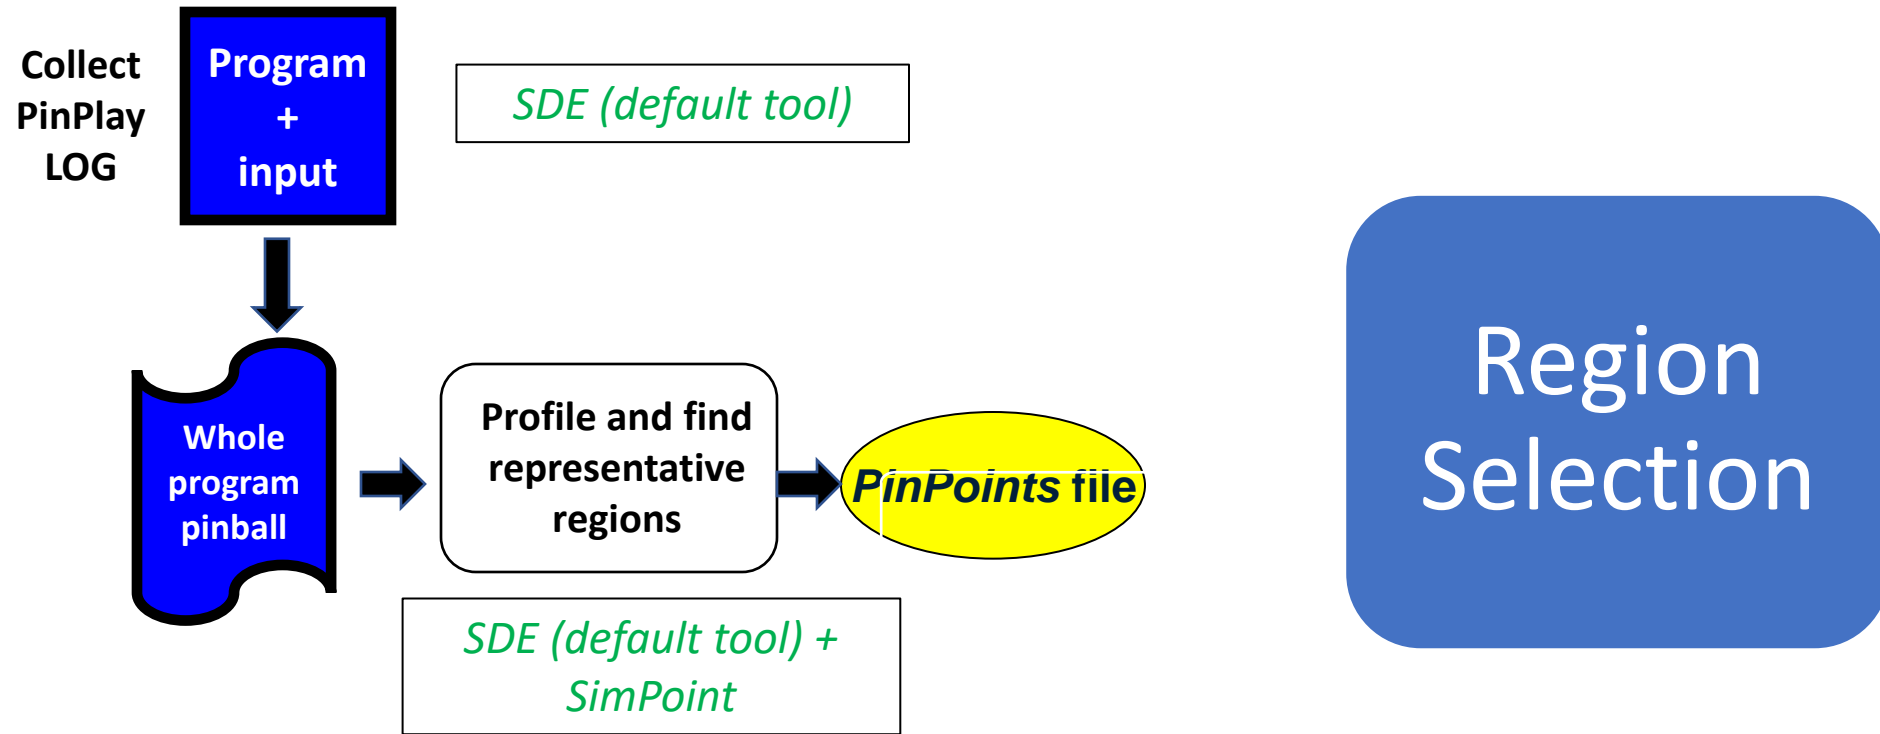
Usage example



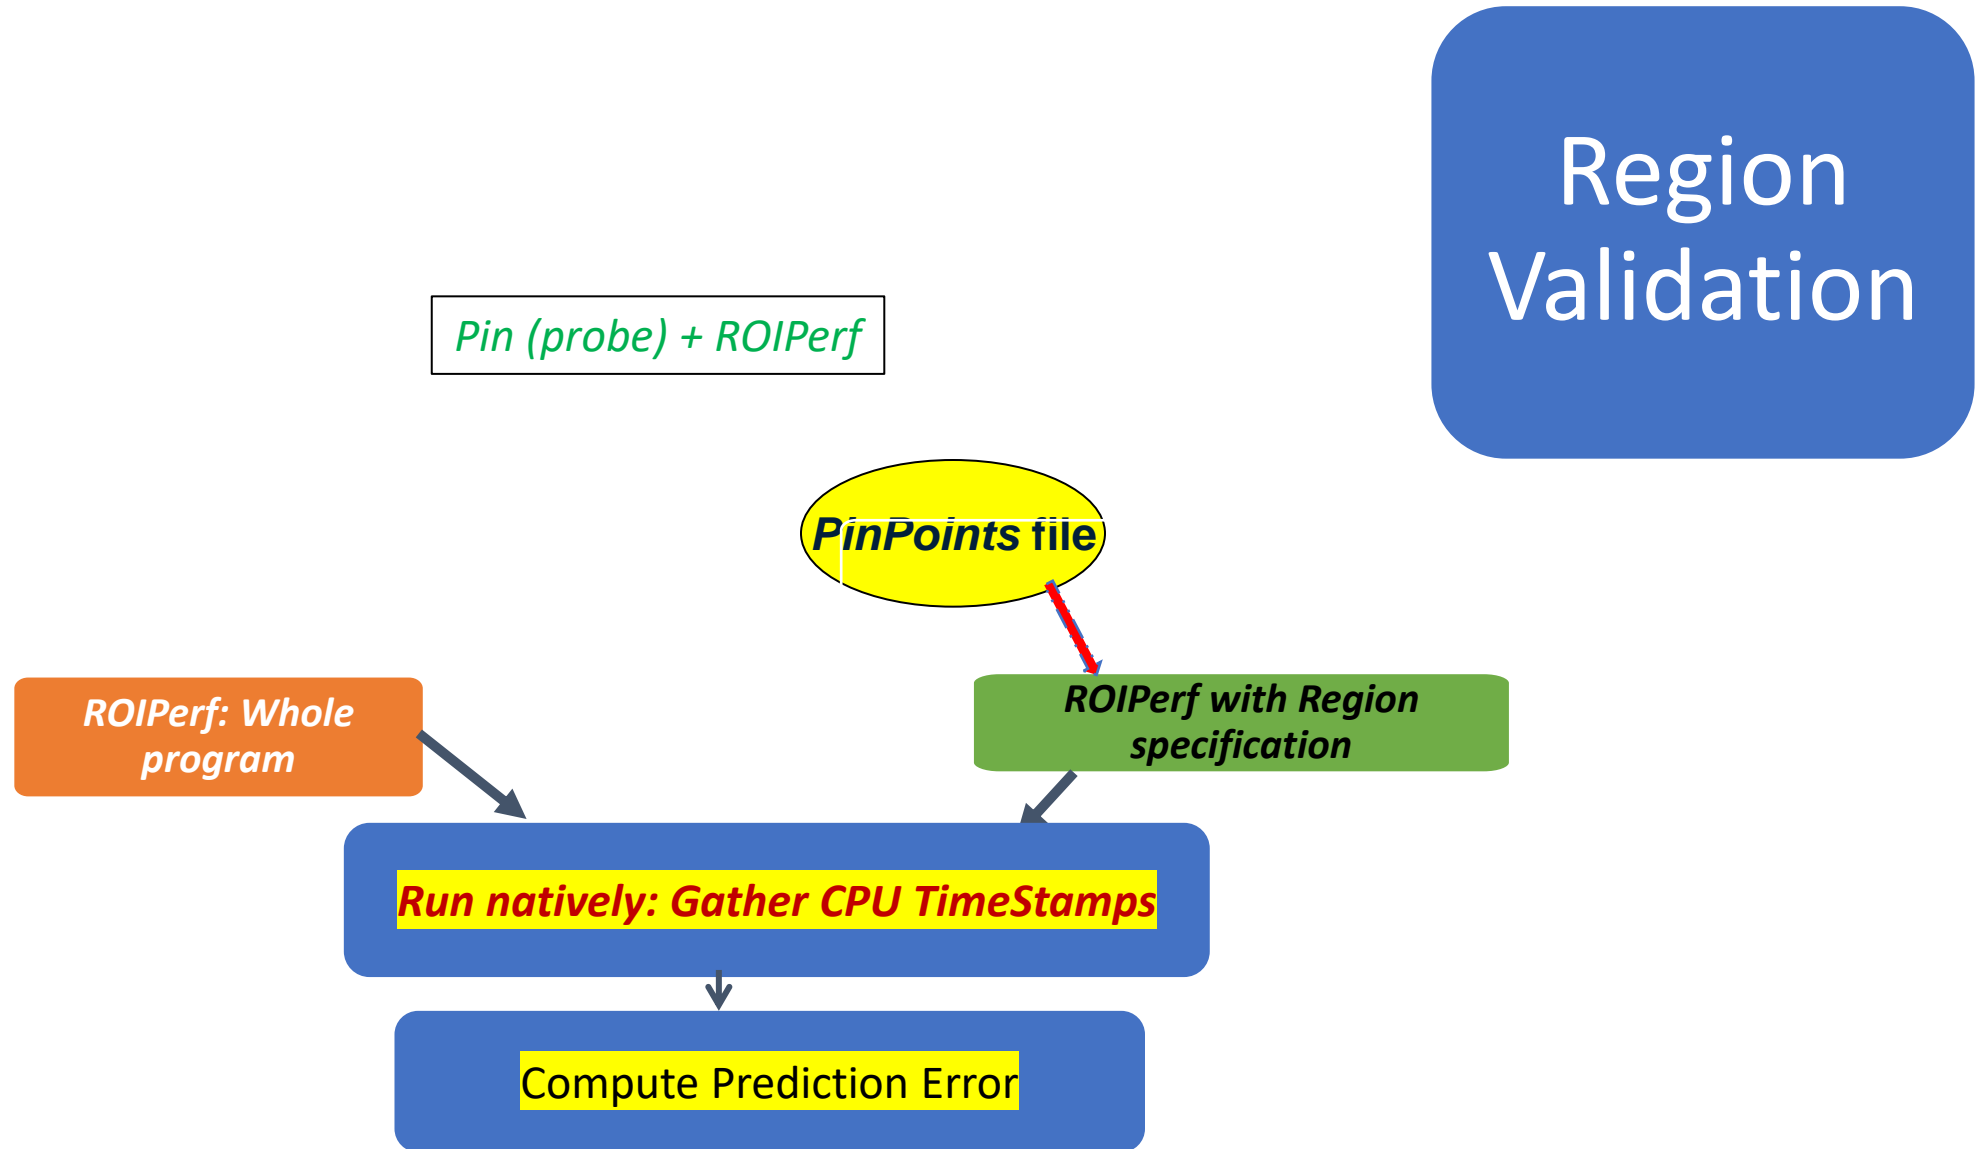
PinPoints Methodology



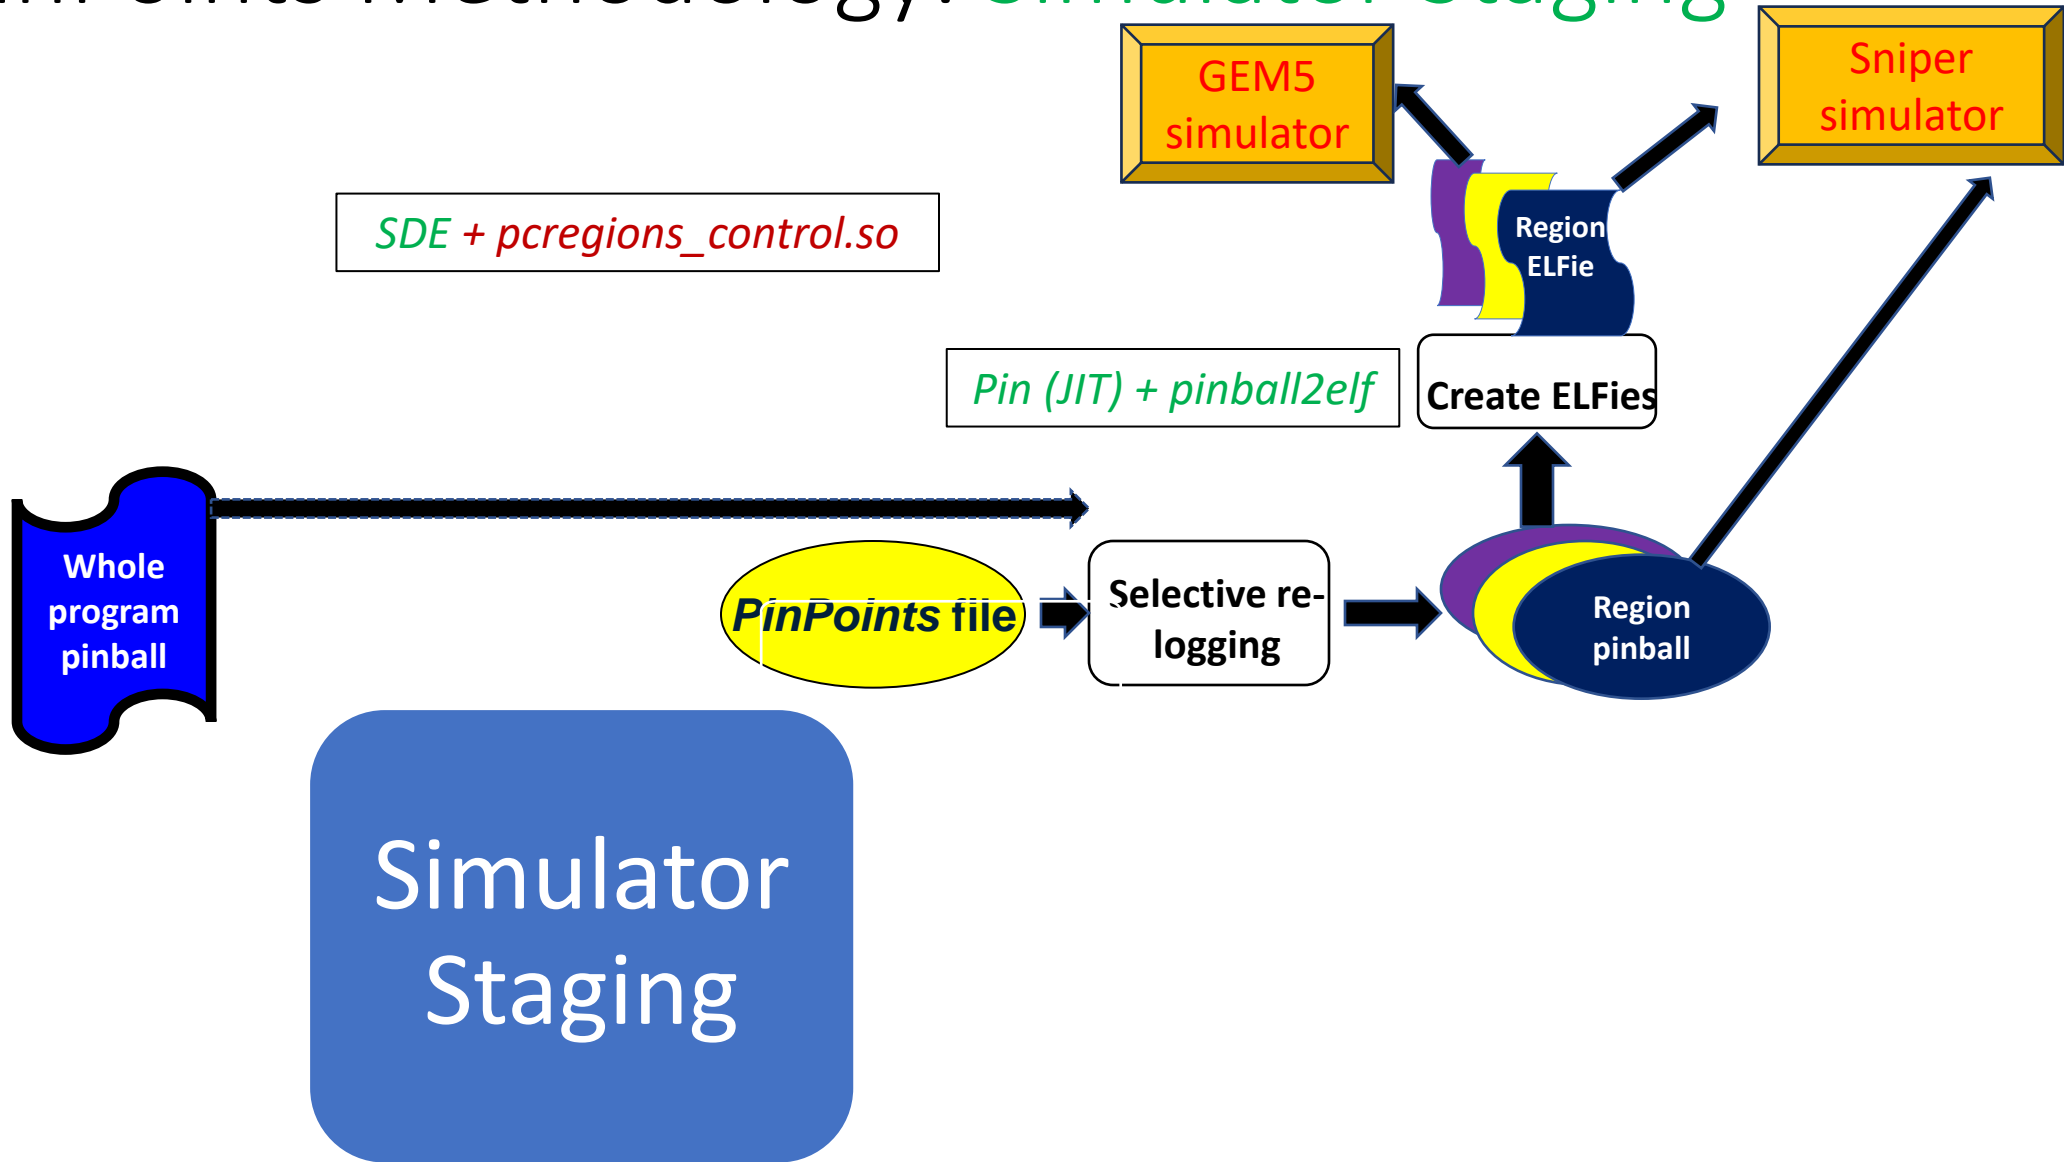
PinPoints Methodology: Simulation Region Selection



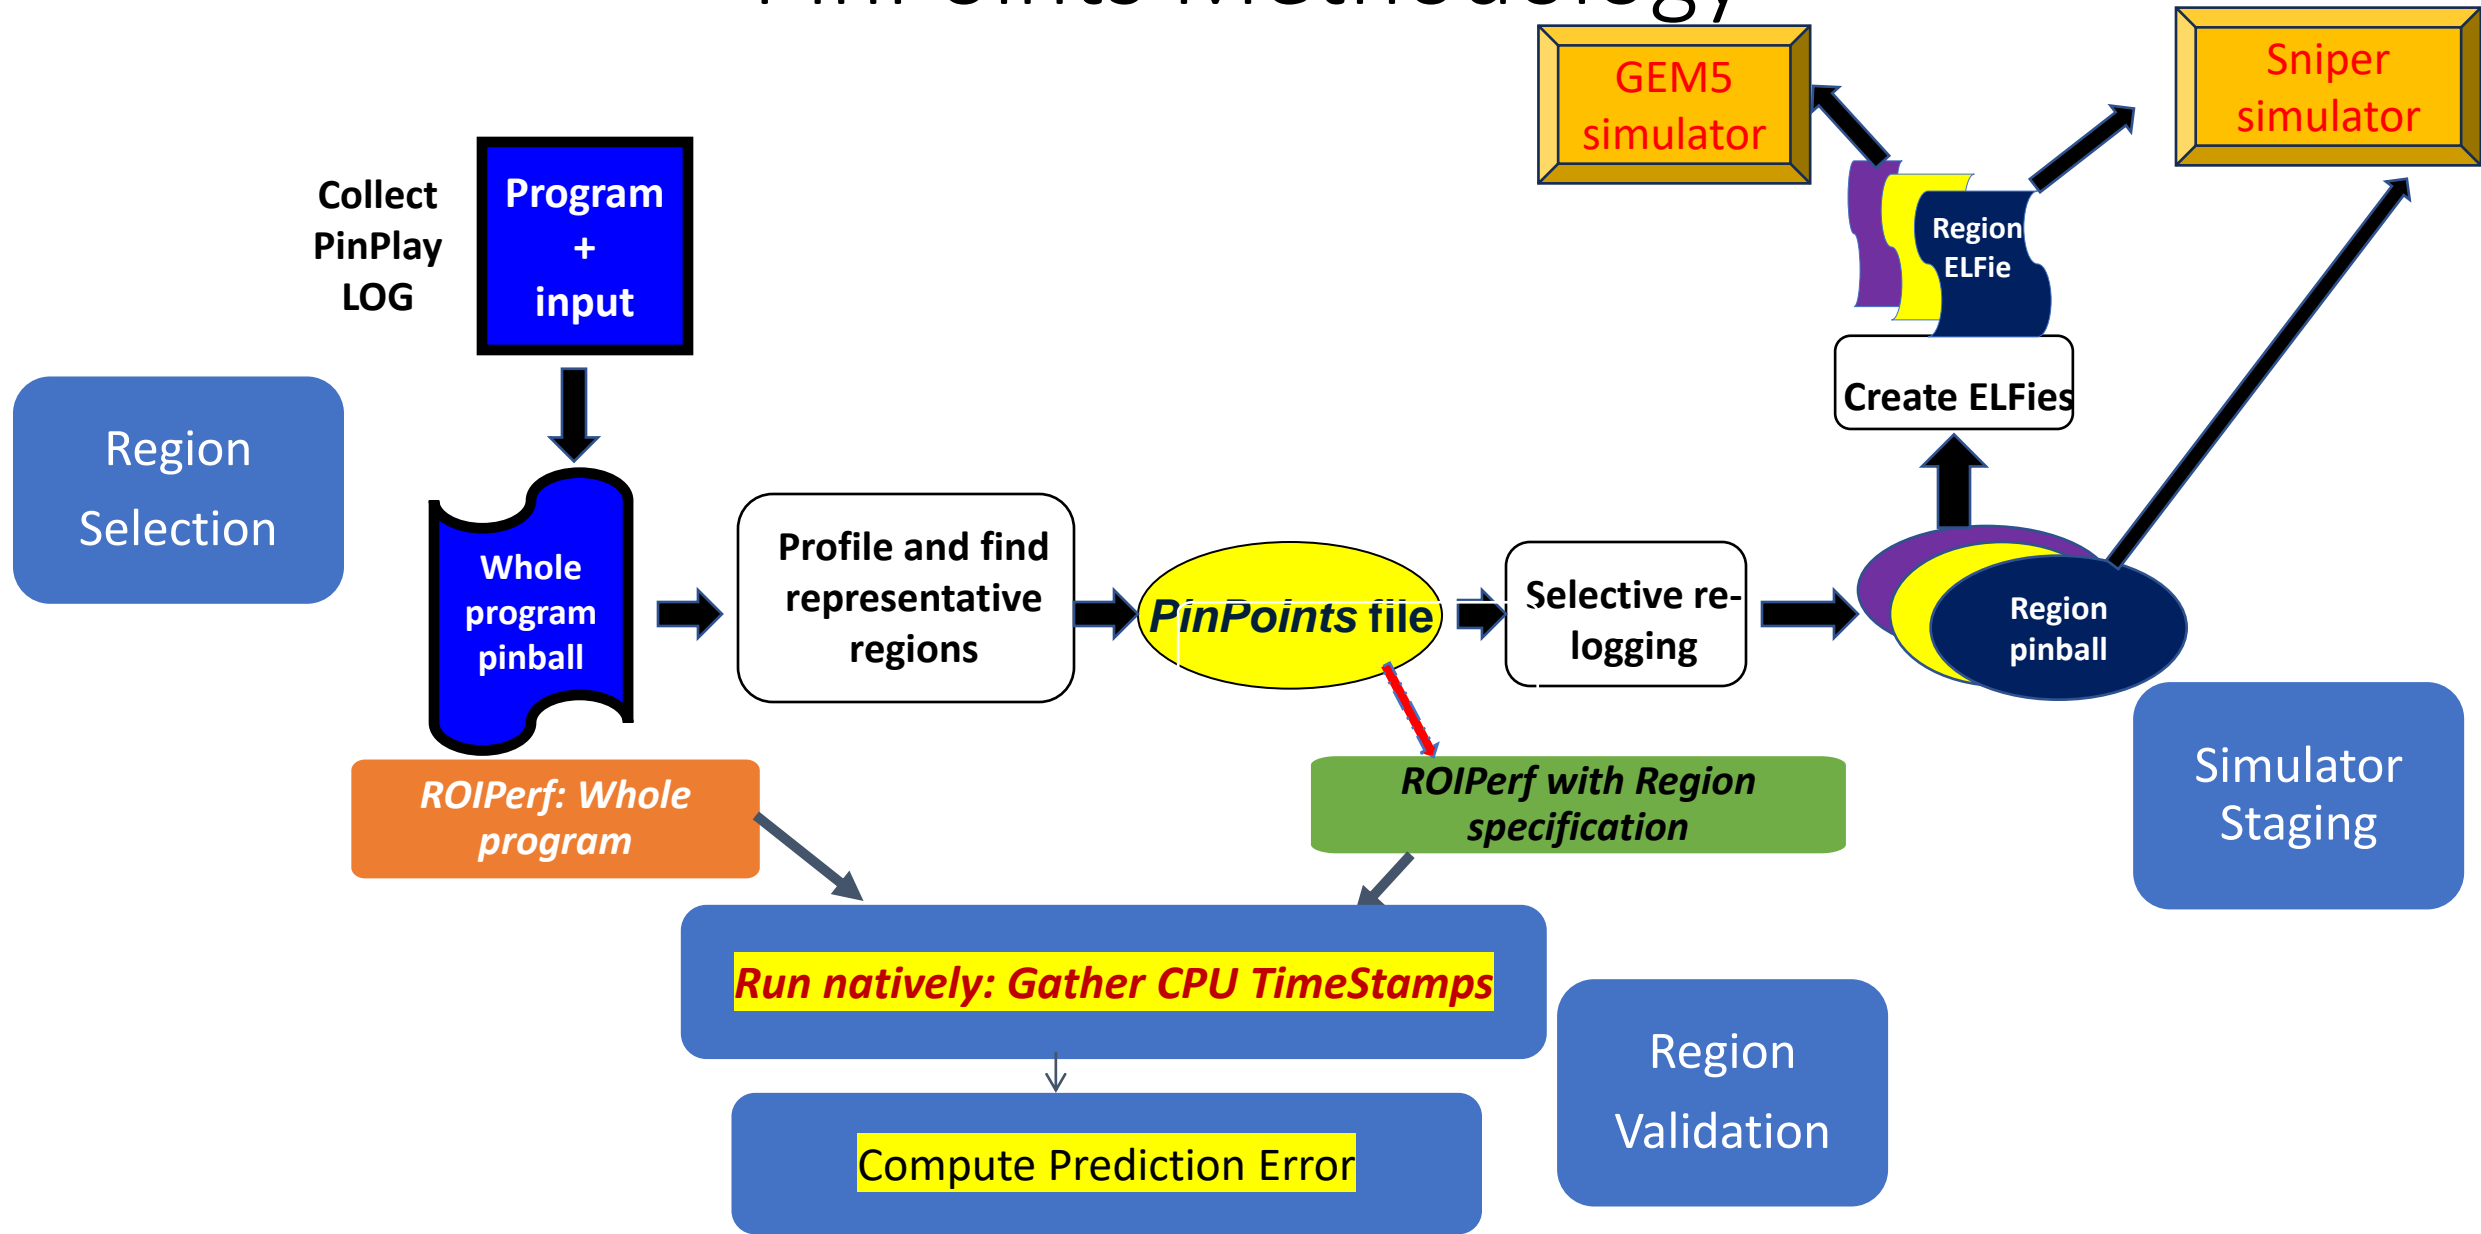
PinPoints Methodology: Validation



PinPoints Methodology: Simulator Staging



PinPoints Methodology



Demo:Part 1: Region Selection + region pinball generation

Pre-requisites

- Pin kit 3.31 : <http://pintool.intel.com> (or [this link](#))
 - export **PIN_ROOT**=<path to the local copy of Pin kit 3.31>
- Intel SDE 9.44 : Use [this link](#)
 - export **SDE_BUILD_KIT**=<path to the local copy of SDE Kit>
- Pinplay-tools repo: <https://github.com/intel/pinplay-tools>
 - export **PINPLAY_TOOLS**=<path to the local clone of the pinplay-tools repo>
- Pinball2elf repo: <https://github.com/intel/pinball2elf>
 - export **PINBALL2ELF**=<path to the local clone of the pinball2elf repo>

< Assumption: using "bash" >

< put "." in PATH >

Build SDE profiler: 'pcregions_control.so'

```
% cd $SDE_BUILD_KIT/pinkit/sde-example/example
```

```
% make TARGET=intel64 clean; make TARGET=intel64
```

```
% cp obj-intel64/pcregions_control.so $SDE_BUILD_KIT/intel64
```

Test case: \$PINPLAY_TOOLS/PinPoints/Test


- cd \$PINPLAY_TOOLS/PinPoints/Test

- **< Edit** *sde-run.pinpoints.single-threaded.sh* **>**

```
SLICESIZE=75000  
WARMUP_FACTOR=3  
MAXK=20  
PROGRAM=dotproduct-st  
INPUT=1  
COMMAND="./dotproduct-st"
```

% make clean; make

% sde-run.pinpoints.single-threaded.sh



WARMUP_FACTOR=3



**Include 3 slices before
simulation region for
warmup**

sde-run.pinpoints.single-thread.sh

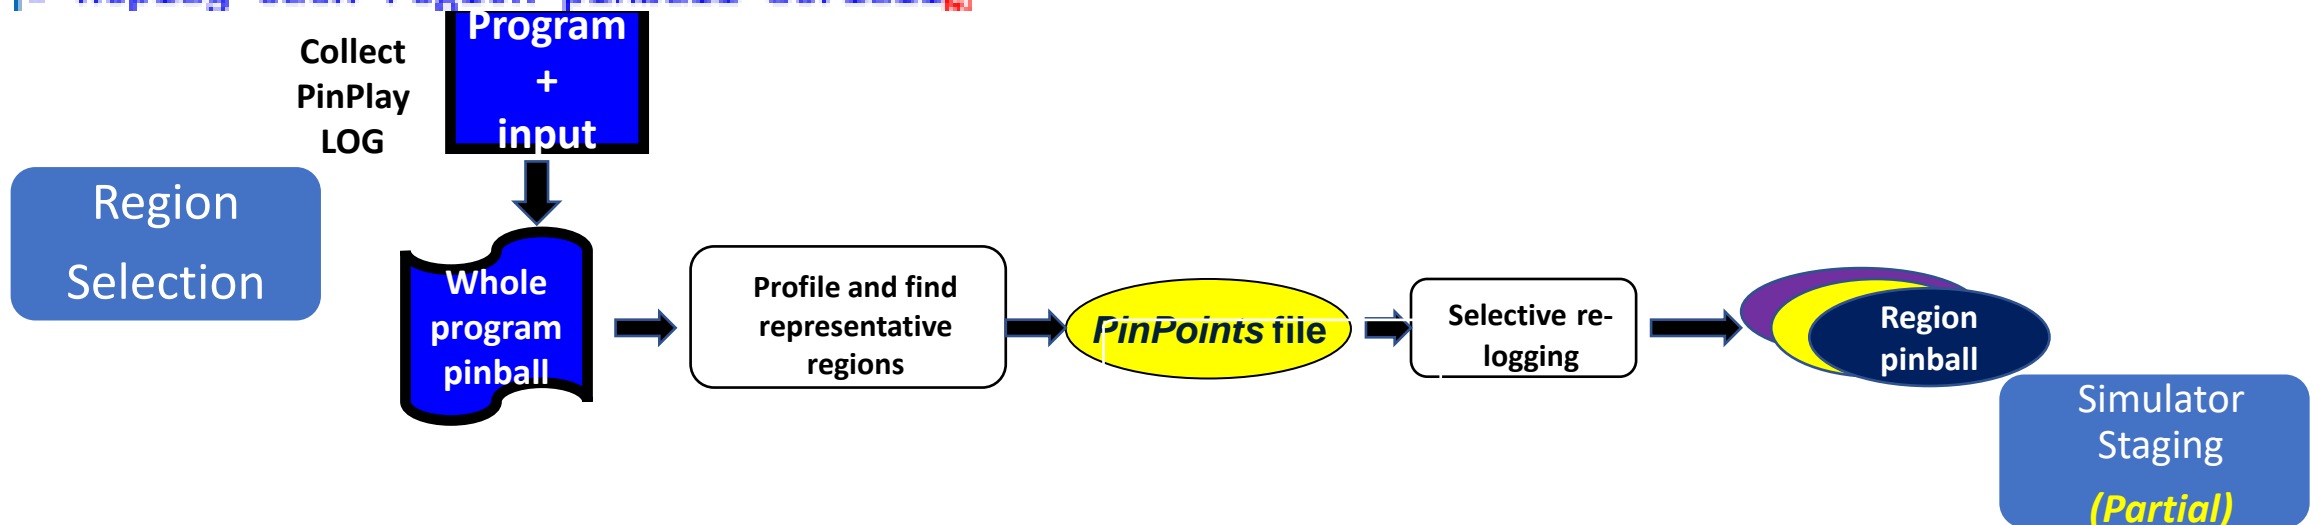
```
#Whole Program Logging and replay using the default sde tool
# We are recording starting at 'main'
$SDE_BUILD_KIT/pinplay-scripts/sde_pinpoints.py --pin_options "$SDE_ARCH" $GLOBAL $PCCOUNT --program_name=$PROG
RAM --input_name=$INPUT --command="$COMMAND" --delete --mode st --log_options="-start_address main -log:fat -lo
g:mp_mode 0 -log:mp_atomic 0" --replay_options="-replay:strace" -l -r

#Profiling using regular profiler from the default sde tool
$SDE_BUILD_KIT/pinplay-scripts/sde_pinpoints.py --pin_options "$SDE_ARCH" $GLOBAL $PCCOUNT --program_name=$PROG
RAM --input_name=$INPUT --command="$COMMAND" --mode st -S $SLICESIZE -b

#Simpoint
$SDE_BUILD_KIT/pinplay-scripts/sde_pinpoints.py --pin_options "$SDE_ARCH" $GLOBAL $PCCOUNT --program_name=$PROG
GRAM --input_name=$INPUT --command="$COMMAND" $PCCOUNT -S $SLICESIZE $WARMUP --maxk=$MAXK --append_status -s

# Create per-region CSV files
#Create Makefile.regions with commands for relogging all regions
# and use 'make -j Makefile.regions' create all region pinballs in parallel

# Replay each region pinball serially
```



BYOP: Bring Your Own (whole-program) Pinball

- Generate pinball any other way, say at another site
- Say the pinball is input.pinball/log_0

```
% $SDE_BUILD_KIT/pinplay-scripts/sde_pinpoints.py ..  
--whole_pgm_dir input.pinball --pin_options  
"$SDE_ARCH" $GLOBAL $PCCOUNT --program_name=$PROGRAM  
--input_name=$INPUT --command="$COMMAND" --mode st -S  
$SLICESIZE -b
```

Demo:Part 2: Region Selection Validation

Build sde-global-event-icounter tool

(used for PCCount region processing)

```
% cd $PINPLAY_TOOLS/GlobalLoopPoint/EventCounter/
```

```
% make clean TARGET=intel64; make build TARGET=intel64
```


Build ROIPerf tool

(used for time-stamp[rdtsc] generation)

```
% cd $PINBALL2ELF/pintools/ROIProbe
```


```
% make clean; make
```

```
% cp obj-intel64/pcregions_control.so $SDE_BUILD_KIT/intel64
```

Run ROIPerf-based validation

```
% cd $PINPLAY_TOOLS/PinPoints/Test
```

```
% run.ROIPerf.sh
```



```
wp_rdtsc,   region_rdtsc,   err%  
702607571.25,   504807628.48800004,   28.15%
```

- Prediction error can be tuned by changing SLICESIZE and MAXK
- Error high for short-running programs due to measurement overhead

Validation of PinPoints

SPEC2017 Rate/Train input: Simulation vs ROIperf

	% CPI Prediction Error		
	Intel-Sniper: Skylake Server	ROIperf: BroadWell Server	ROIperf: Skylake Server
perlbench_r.train.1	0.3%	0.5%	0.7%
perlbench_r.train.2	1.3%	0.0%	2.7%
perlbench_r.train.3	0.8%	0.0%	-0.1%
perlbench_r.train.4	0.2%	3.6%	3.0%
perlbench_r.train.5	8.9%	-6.6%	-0.7%
gcc_r.train.1	10.4%	-3.7%	-6.1%
gcc_r.train.2	1.0%	3.7%	0.4%
gcc_r.train.3	0.8%	3.9%	4.3%
mcf_r.train.1	8.8%	2.8%	0.2%
omnetpp_r.train.1	1.2%	2.8%	1.4%
x264_r.train.1	3.3%	0.0%	0.0%
leela_r.train.1	2.1%	-0.2%	-0.4%
exchange2_r.train.1	5.4%	0.0%	0.9%
xz_r.train.1	4.0%	-8.2%	-7.3%
xz_r.train.2	3.5%	-3.7%	-3.3%

Sniper Simulation time: up to 5 weeks
ROIperf validation time (3 trials each): few hours

Demo:Part 3: Creating ELFies and script templates for Sniper and GEM-5

ELFie generation

% run.pinball2elf.sh

% run.elfies.sh

<GEM5/Sniper template script generation to be added soon>

Resources

- [This tutorial](#)
- Past PinPoints tutorials (commands/kits outdated)
 - [ISCA2014-PinPoints-Tutorial](#)
 - [HPCA2013-PinPoints-Tutorial](#)
- Past PinPlay tutorials (commands/kits outdated)
 - [PLDI2015-PinPlay-Tutorial](#)
 - [PLDI2016-PinPlay-Tutorial](#)
- LoopPoint Methodology (multi-threaded programs): See <https://looppoint.github.io/>